

Calling Conventions

Radare2 uses calling conventions to help in identifying function formal arguments and return types. It is used also as a guide for basic function prototype and type propagation.

```
[0x00000000]> afc?
|Usage: afc[agl?]
| afc convention  Manually set calling convention for current function
| afc             Show Calling convention for the Current function
| afc=[cctype])  Select or show default calling convention
| afcr[j]        Show register usage for the current function
| afca           Analyse function for finding the current calling convention
| afcf[j] [name] Prints return type function(arg1, arg2...), see afij
| afck           List SDB details of call loaded calling conventions
| afcl           List all available calling conventions
| afco path      Open Calling Convention sdb profile from given path
| afcR           Register telescoping using the calling conventions order
[0x00000000]>
```

- To list all available calling conventions for current architecture using `afcl` command

```
[0x00000000]> afcl
amd64
ms
```

- To display function prototype of standard library functions you have `afcf` command

```
[0x00000000]> afcf printf
int printf(const char *format)
[0x00000000]> afcf fgets
char *fgets(char *s, int size, FILE *stream)
```

All this information is loaded via sdb under `/libr/anal/d/cc-[arch]-[bits].sdb`

```
default.cc=amd64
```

```
ms=cc
cc.ms.name=ms
cc.ms.arg1=rcx
cc.ms.arg2=rdx
cc.ms.arg3=r8
cc.ms.arg3=r9
cc.ms.argn=stack
cc.ms.ret=rax
```

`cc.x.argi=rax` is used to set the `ith` argument of this calling convention to register name `rax`

`cc.x.argn=stack` means that all the arguments (or the rest of them in case there was argi for any i as counting number) will be stored in stack from left to right

`cc.x.argn=stack_rev` same as `cc.x.argn=stack` except for it means argument are passed right to left

Code Analysis

Code analysis is a common technique used to extract information from assembly code.

Radare2 has different code analysis techniques implemented in the core and available in different commands.

As long as the whole functionalities of r2 are available with the API as well as using commands. This gives you the ability to implement your own analysis loops using any programming language, even with r2 oneliners, shellscripts, or analysis or core native plugins.

The analysis will show up the internal data structures to identify basic blocks, function trees and to extract opcode-level information.

The most common radare2 analysis command sequence is `aa`, which stands for “analyze all”. That all is referring to all symbols and entry-points. If your binary is stripped you will need to use other commands like `aaa`, `aab`, `aar`, `aac` or so.

Take some time to understand what each command does and the results after running them to find the best one for your needs.

```
[0x08048440]> aa
[0x08048440]> pdf @ main
                ; DATA XREF from 0x08048457 (entry0)
/ (fcn) fcn.08048648 141
|      ;-- main:
|      0x08048648      8d4c2404      lea ecx, [esp+0x4]
|      0x0804864c      83e4f0      and esp, 0xffffffff0
|      0x0804864f      ff71fc      push dword [ecx-0x4]
|      0x08048652      55          push ebp
|      ; CODE (CALL) XREF from 0x08048734 (fcn.080486e5)
|      0x08048653      89e5          mov ebp, esp
|      0x08048655      83ec28      sub esp, 0x28
|      0x08048658      894df4      mov [ebp-0xc], ecx
|      0x0804865b      895df8      mov [ebp-0x8], ebx
|      0x0804865e      8975fc      mov [ebp-0x4], esi
|      0x08048661      8b19          mov ebx, [ecx]
|      0x08048663      8b7104      mov esi, [ecx+0x4]
|      0x08048666      c744240c00. mov dword [esp+0xc], 0x0
```

```

|      0x0804866e    c7442408010. mov dword [esp+0x8], 0x1 ; 0x00000001
|      0x08048676    c7442404000. mov dword [esp+0x4], 0x0
|      0x0804867e    c7042400000. mov dword [esp], 0x0
|      0x08048685    e852fdffff   call sym..imp.ptrace
|      sym..imp.ptrace(unk, unk)
|      0x0804868a    85c0          test eax, eax
| ,=< 0x0804868c    7911          jns 0x0804869f
| |      0x0804868e    c70424cf870. mov dword [esp], str.Don_tuseadebuguer_ ; 0x080487cf
| |      0x08048695    e882fdffff   call sym..imp.puts
| |      sym..imp.puts()
| |      0x0804869a    e80dfdffff   call sym..imp.abort
| |      sym..imp.abort()
| `-> 0x0804869f    83fb02          cmp ebx, 0x2
| ,==< 0x080486a2    7411          je 0x080486b5
| |      0x080486a4    c704240c880. mov dword [esp], str.Youmustgiveapasswordforusethisprogram.
| |      0x080486ab    e86cfdffff   call sym..imp.puts
| |      sym..imp.puts()
| |      0x080486b0    e8f7fcffff   call sym..imp.abort
| |      sym..imp.abort()
| `--> 0x080486b5    8b4604          mov eax, [esi+0x4]
|      0x080486b8    890424          mov [esp], eax
|      0x080486bb    e8e5feffff   call fcn.080485a5
|      fcn.080485a5() ; fcn.080484c6+223
|      0x080486c0    b800000000      mov eax, 0x0
|      0x080486c5    8b4df4          mov ecx, [ebp-0xc]
|      0x080486c8    8b5df8          mov ebx, [ebp-0x8]
|      0x080486cb    8b75fc          mov esi, [ebp-0x4]
|      0x080486ce    89ec          mov esp, ebp
|      0x080486d0    5d          pop ebp
|      0x080486d1    8d61fc          lea esp, [ecx-0x4]
\      0x080486d4    c3          ret

```

In this example, we analyze the whole file (**aa**) and then print disassembly of the **main()** function (**pdf**). The **aa** command belongs to the family of auto analysis commands and performs only the most basic auto analysis steps. In radare2 there are many different types of the auto analysis commands with a different analysis depth, including partial emulation: **aa**, **aaa**, **aab**, **aaaa**, ... There is also a mapping of those commands to the r2 CLI options: **r2 -A**, **r2 -AA**, and so on.

It is a common sense that completely automated analysis can produce non-sequitur results, thus radare2 provides separate commands for the particular stages of the analysis allowing fine-grained control of the analysis process. Moreover, there is a treasure trove of configuration variables for controlling the analysis outcomes. You can find them in **anal.*** and **emu.*** cfg variables' namespaces.

Analyze functions

One of the most important “basic” analysis commands is the set of **af** subcommands. **af** means “analyze function”. Using this command you can either allow automatic analysis of the particular function or perform completely manual one.

```
[0x00000000]> af?
```

```
Usage: af
```

af ([name]) ([addr])	analyze functions (start at addr or \$\$)
afr ([name]) ([addr])	analyze functions recursively
af+ addr name [type] [diff]	hand craft a function (requires afb+)
af- [addr]	clean all function analysis data (or function at addr)
afa	analyze function arguments in a call (afal honors db)
afb+ fcnA bbA sz [j] [f] ([t]([d]))	add bb to function @ fcnaddr
afb[?] [addr]	List basic blocks of given function
afbF([0 1])	Toggle the basic-block 'folded' attribute
afB 16	set current function as thumb (change asm.bits)
afC[lc] ([addr])@[addr]	calculate the Cycles (afC) or Cyclomatic Complexity
afc[?] type @[addr]	set calling convention for function
afd[addr]	show function + delta for given offset
afF[1 0]	fold/unfold/toggle
afi [addr fcn.name]	show function(s) information (verbose afl)
afj [tableaddr] [count]	analyze function jumptable
afl[?] [ls*] [fcn name]	list functions (addr, size, bbs, name) (see afl)
afm name	merge two functions
afM name	print functions map
afn[?] name [addr]	rename name for function at address (change flag too)
afna	suggest automatic name for current offset
afo[?j] [fcn.name]	show address for the function name or current offset
afs[!] ([fcnsign])	get/set function signature at current address (afs!)
afS[stack_size]	set stack frame size for function at current address
afsr [function_name] [new_type]	change type for given function
aft[?]	type matching, type propagation
afu addr	resize and analyze function from current address until
afv[absrx]?	manipulate args, registers and variables in function
afx	list function references

You can use **afl** to list the functions found by the analysis.

There are a lot of useful commands under **afl** such as **aflj**, which lists the function in JSON format and **aflm**, which lists the functions in the syntax found in makefiles.

There's also **afl=**, which displays ASCII-art bars with function ranges.

You can find the rest of them under **afl?**.

Some of the most challenging tasks while performing a function analysis are merge, crop and resize. As with other analysis commands you have two modes: semi-

automatic and manual. For the semi-automatic, you can use `afm <function name>` to merge the current function with the one specified by name as an argument, `aff` to readjust the function after analysis changes or function edits, `afu <address>` to do the resize and analysis of the current function until the specified address.

Apart from those semi-automatic ways to edit/analyze the function, you can hand craft it in the manual mode with `af+` command and edit basic blocks of it using `afb` commands. Before changing the basic blocks of the function it is recommended to check the already presented ones:

```
[0x00003ac0]> afb
0x00003ac0 0x00003b7f 01:001A 191 f 0x00003b7f
0x00003b7f 0x00003b84 00:0000 5 j 0x00003b92 f 0x00003b84
0x00003b84 0x00003b8d 00:0000 9 f 0x00003b8d
0x00003b8d 0x00003b92 00:0000 5
0x00003b92 0x00003ba8 01:0030 22 j 0x00003ba8
0x00003ba8 0x00003bf9 00:0000 81
```

Hand craft function

before start, let's prepare a binary file first, for example:

```
int code_block()
{
    int result = 0;

    for(int i = 0; i < 10; ++i)
        result += 1;

    return result;
}
```

then compile it with `gcc -c example.c -m32 -O0 -fno-pie`, we will get the object file `example.o`. open it with `radare2`.

since we haven't analyzed it yet, the `pdf` command will not print out the disassembly here:

```
$ r2 example.o
[0x08000034]> pdf
p: Cannot find function at 0x08000034
[0x08000034]> pd
      ;-- section..text:
      ;-- .text:
      ;-- code_block:
      ;-- eip:
0x08000034      55      push ebp
0x08000035      89e5      mov ebp, esp
; [01] -r-x section s
```

0x08000037	83ec10	sub esp, 0x10
0x0800003a	c745f8000000.	mov dword [ebp - 8], 0
0x08000041	c745fc000000.	mov dword [ebp - 4], 0
,=< 0x08000048	eb08	jmp 0x8000052
.--> 0x0800004a	8345f801	add dword [ebp - 8], 1
: 0x0800004e	8345fc01	add dword [ebp - 4], 1
:`-> 0x08000052	837dfc09	cmp dword [ebp - 4], 9
`==< 0x08000056	7ef2	jle 0x800004a
0x08000058	8b45f8	mov eax, dword [ebp - 8]
0x0800005b	c9	leave
0x0800005c	c3	ret

our goal is to hand craft a function with the following structure

analyze_one

create a function at 0x8000034 named code_block:

```
[0x8000034]> af+ 0x8000034 code_block
```

In most cases, we use jump or call instructions as code block boundaries. so the range of first block is from 0x08000034 push ebp to 0x08000048 jmp 0x8000052. use afb+ command to add it.

```
[0x8000034]> afb+ code_block 0x8000034 0x800004a-0x8000034 0x8000052
```

note that the basic syntax of afb+ is afb+ function_address block_address block_size [jump] [fail]. the final instruction of this block points to a new address(jmp 0x8000052), thus we add the address of jump target (0x8000052) to reflect the jump info.

the next block (0x08000052 ~ 0x08000056) is more likely an if conditional statement which has two branches. It will jump to 0x800004a if jle-less or equal, otherwise (the fail condition) jump to next instruction - 0x08000058.:

```
[0x8000034]> afb+ code_block 0x8000052 0x8000058-0x8000052 0x800004a 0x8000058
```

follow the control flow and create the remaining two blocks (two branches) :

```
[0x8000034]> afb+ code_block 0x800004a 0x8000052-0x800004a 0x8000052
```

```
[0x8000034]> afb+ code_block 0x8000058 0x800005d-0x8000058
```

check our work:

```
[0x8000034]> afb
```

```
0x08000034 0x800004a 00:0000 22 j 0x8000052
0x0800004a 0x8000052 00:0000 8 j 0x8000052
0x08000052 0x8000058 00:0000 6 j 0x800004a f 0x8000058
0x08000058 0x800005d 00:0000 5
```

```
[0x8000034]> VV
```

handcraft_one

There are two very important commands for this: `afc` and `afB`. The latter is a must-know command for some platforms like ARM. It provides a way to change the “bitness” of the particular function. Basically, allowing to select between ARM and Thumb modes.

`afc` on the other side, allows to manually specify function calling convention. You can find more information on its usage in `calling_conventions`.

Recursive analysis

There are 5 important program wide half-automated analysis commands:

- `aab` - perform basic-block analysis (“Nucleus” algorithm)
- `aac` - analyze function calls from one (selected or current function)
- `aaf` - analyze all function calls
- `aar` - analyze data references
- `aad` - analyze pointers to pointers references

Those are only generic semi-automated reference searching algorithms. Radare2 provides a wide choice of manual references’ creation of any kind. For this fine-grained control you can use `ax` commands.

```
Usage: ax[?d-l*]    # see also 'afx?'
| ax                list refs
| ax*               output radare commands
| ax addr [at]      add code ref pointing to addr (from curseek)
| ax- [at]          clean all refs/refs from addr
| ax-*             clean all refs/refs
| axc addr [at]     add generic code ref
| axC addr [at]     add code call ref
| axg [addr]        show xrefs graph to reach current function
| axg* [addr]       show xrefs graph to given address, use .axg*;aggv
| axgj [addr]       show xrefs graph to reach current function in json format
| axd addr [at]     add data ref
| axq              list refs in quiet/human-readable format
| axj              list refs in json format
| axF [flg-glob]    find data/code references of flags
| axm addr [at]     copy data/code references pointing to addr to also point to curseek (or at)
| axt [addr]        find data/code references to this address
| axf [addr]        find data/code references from this address
| axv [addr]        list local variables read-write-exec references
| ax. [addr]        find data/code references from and to this address
| axff[j] [addr]    find data/code references from this function
| axs addr [at]     add string ref
```

The most commonly used `ax` commands are `axt` and `axf`, especially as a part of various `r2pipe` scripts. Lets say we see the string in the data or a code section and want to find all places it was referenced from, we should use `axt`:

```

[0x0001783a]> pd 2
;-- str.02x:
; STRING XREF from 0x00005de0 (sub.strlen_d50)
; CODE XREF from 0x00017838 (str._s_s_s + 7)
0x0001783a      .string "%%02x" ; len=7
;-- str.src_ls.c:
; STRING XREF from 0x0000541b (sub.free_b04)
; STRING XREF from 0x0000543a (sub.__assert_fail_41f + 27)
; STRING XREF from 0x00005459 (sub.__assert_fail_41f + 58)
; STRING XREF from 0x00005f9e (sub._setjmp_e30)
; CODE XREF from 0x0001783f (str.02x + 5)
0x00017841 .string "src/ls.c" ; len=9
[0x0001783a]> axt
sub.strlen_d50 0x5de0 [STRING] lea rcx, str.02x
(nofunc) 0x17838 [CODE] jae str.02x

```

There are also some useful commands under `axt`. Use `axtg` to generate radare2 commands which will help you to create graphs according to the XREFs.

```

[0x08048320]> s main
[0x080483e0]> axtg
agn 0x8048337 "entry0 + 23"
agn 0x80483e0 "main"
age 0x8048337 0x80483e0

```

Use `axt*` to split the radare2 commands and set flags on those corresponding XREFs.

Also under `ax` is `axg`, which finds the path between two points in the file by showing an XREFs graph to reach the location or function. For example:

```

:> axg sym.imp.printf
- 0x08048a5c fcn 0x08048a5c sym.imp.printf
  - 0x080483e5 fcn 0x080483e0 main
  - 0x080483e0 fcn 0x080483e0 main
    - 0x08048337 fcn 0x08048320 entry0
    - 0x08048425 fcn 0x080483e0 main

```

Use `axg*` to generate radare2 commands which will help you to create graphs using `agn` and `age` commands, according to the XREFs.

Apart from predefined algorithms to identify functions there is a way to specify a function prelude with a configuration option `anal.prelude`. For example, like `e anal.prelude = 0x554889e5` which means

```

push rbp
mov rbp, rsp

```

on `x86_64` platform. It should be specified *before* any analysis commands.

Configuration

Radare2 allows to change the behavior of almost any analysis stages or commands. There are different kinds of the configuration options:

- Flow control
- Basic blocks control
- References control
- IO/Ranges
- Jump tables analysis control
- Platform/target specific options

Control flow configuration

Two most commonly used options for changing the behavior of control flow analysis in radare2 are **anal.hasnext** and **anal.jump.after**. The first one allows forcing radare2 to continue the analysis after the end of the function, even if the next chunk of the code wasn't called anywhere, thus analyzing all of the available functions. The latter one allows forcing radare2 to continue the analysis even after unconditional jumps.

In addition to those we can also set **anal.jump.indir** to follow the indirect jumps, continuing analysis; **anal.pushret** to analyze **push ...; ret** sequence as a jump; **anal.nopskip** to skip the NOP sequences at a function beginning.

For now, radare2 also allows you to change the maximum basic block size with **anal.bb.maxsize** option. The default value just works in most use cases, but it's useful to increase that for example when dealing with obfuscated code. Beware that some of basic blocks control options may disappear in the future in favor of more automated ways to set those.

For some unusual binaries or targets, there is an option **anal.noncode**. Radare2 doesn't try to analyze data sections as a code by default. But in some cases - malware, packed binaries, binaries for embedded systems, it is often a case. Thus - this option.

Reference control

The most crucial options that change the analysis results drastically. Sometimes some can be disabled to save the time and memory when analyzing big binaries.

- **anal.jump.ref** - to allow references creation for unconditional jumps
- **anal.jump.cref** - same, but for conditional jumps
- **anal.datarefs** - to follow the data references in code
- **anal.refstr** - search for strings in data references
- **anal.strings** - search for strings and creating references

Note that strings references control is disabled by default because it increases the analysis time.

Analysis ranges

There are a few options for this:

- `anal.limits` - enables the range limits for analysis operations
- `anal.from` - starting address of the limit range
- `anal.to` - the corresponding end of the limit range
- `anal.in` - specify search boundaries for analysis. You can set it to `io.maps`, `io.sections.exec`, `dbg.maps` and many more. For example:
 - To analyze a specific memory map with `anal.from` and `anal.to`, set `anal.in = dbg.maps`.
 - To analyze in the boundaries set by `anal.from` and `anal.to`, set `anal.in=range`.
 - To analyze in the current mapped segment or section, you can put `anal.in=bin.segment` or `anal.in=bin.section`, respectively.
 - To analyze in the current memory map, specify `anal.in=dbg.map`.
 - To analyze in the stack or heap, you can set `anal.in=dbg.stack` or `anal.in=dbg.heap`.
 - To analyze in the current function or basic block, you can specify `anal.in=anal.fcn` or `anal.in=anal.bb`.

Please see `e anal.in=??` for the complete list.

Jump tables

Jump tables are one of the trickiest targets in binary reverse engineering. There are hundreds of different types, the end result depending on the compiler/linker and LTO stages of optimization. Thus radare2 allows enabling some experimental jump tables detection algorithms using `anal.jump.tbl` option. Eventually, algorithms moved into the default analysis loops once they start to work on every supported platform/target/testcase. Two more options can affect the jump tables analysis results too:

- `anal.jump.indir` - follow the indirect jumps, some jump tables rely on them
- `anal.datarefs` - follow the data references, some jump tables use those

Platform specific controls

There are two common problems when analyzing embedded targets: ARM/Thumb detection and MIPS GP value. In case of ARM binaries radare2 supports some auto-detection of ARM/Thumb mode switches, but beware that it uses partial ESIL emulation, thus slowing the analysis process. If you will not like the results, particular functions' mode can be overridden with `afB` command.

The MIPS GP problem is even trickier. It is a basic knowledge that GP value can be different not only for the whole program, but also for some functions. To partially solve that there are options `anal.gp` and `anal.gpfixed`. The first

one sets the GP value for the whole program or particular function. The latter allows to “constantify” the GP value if some code is willing to change its value, always resetting it if the case. Those are heavily experimental and might be changed in the future in favor of more automated analysis.

Visuals

One of the easiest way to see and check the changes of the analysis commands and variables is to perform a scrolling in a Vv special visual mode, allowing functions preview:

```
vv
```

When we want to check how analysis changes affect the result in the case of big functions, we can use minimap instead, allowing to see a bigger flow graph on the same screen size. To get into the minimap mode type VV then press p twice:

```
vv2
```

This mode allows you to see the disassembly of each node separately, just navigate between them using Tab key.

Analysis hints

It is not an uncommon case that analysis results are not perfect even after you tried every single configuration option. This is where the “analysis hints” radare2 mechanism comes in. It allows to override some basic opcode or meta-information properties, or even to rewrite the whole opcode string. These commands are located under ah namespace:

```
Usage: ah[lba-]  Analysis Hints
| ah?            show this help
| ah? offset     show hint of given offset
| ah            list hints in human-readable format
| ah.           list hints in human-readable format from current offset
| ah-           remove all hints
| ah- offset [size] remove hints at given offset
| ah* offset     list hints in radare commands format
| aha ppc @ 0x42 force arch ppc for all addrs >= 0x42 or until the next hint
| aha 0 @ 0x84   disable the effect of arch hints for all addrs >= 0x84 or until the next hint
| ahb 16 @ 0x42 force 16bit for all addrs >= 0x42 or until the next hint
| ahb 0 @ 0x84   disable the effect of bits hints for all addrs >= 0x84 or until the next hint
| ahc 0x804804   override call/jump address
| ahd foo a0,33  replace opcode string
| ahe 3,eax,+=   set vm analysis string
| ahf 0x804840   override fallback address for call
| ahF 0x10       set stackframe size at current offset
| ahh 0x804840   highlight this address offset in disasm
```

```

| ahi[?] 10          define numeric base for immediates (2, 8, 10, 10u, 16, i, p, S, s)
| ahj               list hints in JSON
| aho call          change opcode type (see aho?) (deprecated, moved to "ahd")
| ahp addr          set pointer hint
| ahr val           set hint for return value of a function
| ahs 4             set opcode size=4
| ahS jz            set asm.syntax=jz for this opcode
| aht [?] <type>    Mark immediate as a type offset (deprecated, moved to "aho")
| ahv val           change opcode's val field (useful to set jmptbl sizes in jmp rax)

```

One of the most common cases is to set a particular numeric base for immediates:

```

[0x00003d54]> ahi?
Usage: ahi [2|8|10|10u|16|bodhipSs] [@ offset]  Define numeric base
| ahi <base> set numeric base (2, 8, 10, 16)
| ahi 10|d   set base to signed decimal (10), sign bit should depend on receiver size
| ahi 10u|du set base to unsigned decimal (11)
| ahi b      set base to binary (2)
| ahi o      set base to octal (8)
| ahi h      set base to hexadecimal (16)
| ahi i      set base to IP address (32)
| ahi p      set base to htons(port) (3)
| ahi S      set base to syscall (80)
| ahi s      set base to string (1)

```

```

[0x00003d54]> pd 2
0x00003d54      0583000000      add eax, 0x83
0x00003d59      3d13010000      cmp eax, 0x113
[0x00003d54]> ahi d
[0x00003d54]> pd 2
0x00003d54      0583000000      add eax, 131
0x00003d59      3d13010000      cmp eax, 0x113
[0x00003d54]> ahi b
[0x00003d54]> pd 2
0x00003d54      0583000000      add eax, 10000011b
0x00003d59      3d13010000      cmp eax, 0x113

```

It is notable that some analysis stages or commands add the internal analysis hints, which can be checked with **ah** command:

```

[0x00003d54]> ah
0x00003d54 - 0x00003d54 => immbase=2
[0x00003d54]> ah*
ahi 2 @ 0x3d54

```

Sometimes we need to override jump or call address, for example in case of tricky relocation, which is unknown for radare2, thus we can change the value manually. The current analysis information about a particular opcode can be checked with

ao command. We can use `ahc` command for performing such a change:

```
[0x00003cee]> pd 2
0x00003cee      e83d080100      call sub.__errno_location_530
0x00003cf3      85c0          test eax, eax
[0x00003cee]> ao
address: 0x3cee
opcode: call 0x14530
mnemonic: call
prefix: 0
id: 56
bytes: e83d080100
refptr: 0
size: 5
sign: false
type: call
cycles: 3
esil: 83248,rip,8,rsp,-=,rsp,=[],rip,=
jump: 0x00014530
direction: exec
fail: 0x00003cf3
stack: null
family: cpu
stackop: null
[0x00003cee]> ahc 0x5382
[0x00003cee]> pd 2
0x00003cee      e83d080100      call sub.__errno_location_530
0x00003cf3      85c0          test eax, eax
[0x00003cee]> ao
address: 0x3cee
opcode: call 0x14530
mnemonic: call
prefix: 0
id: 56
bytes: e83d080100
refptr: 0
size: 5
sign: false
type: call
cycles: 3
esil: 83248,rip,8,rsp,-=,rsp,=[],rip,=
jump: 0x00005382
direction: exec
fail: 0x00003cf3
stack: null
family: cpu
```

```
stackop: null
[0x00003cee]> ah
0x00003cee - 0x00003cee => jump: 0x5382
```

As you can see, despite the unchanged disassembly view the jump address in opcode was changed (`jump` option).

If anything of the previously described didn't help, you can simply override shown disassembly with anything you like:

```
[0x00003d54]> pd 2
0x00003d54      0583000000      add eax, 10000011b
0x00003d59      3d13010000      cmp eax, 0x113
[0x00003d54]> "ahd myopcode bla, foo"
[0x00003d54]> pd 2
0x00003d54      myopcode bla, foo
0x00003d55      830000      add dword [rax], 0
```

Emulation

One of the most important things to remember in reverse engineering is a core difference between static analysis and dynamic analysis. As many already know, static analysis suffers from the path explosion problem, which is impossible to solve even in the most basic way without at least a partial emulation.

Thus many professional reverse engineering tools use code emulation while performing an analysis of binary code, and radare2 is no difference here.

For partial emulation (or imprecise full emulation) radare2 uses its own ESIL intermediate language and virtual machine.

Radare2 supports this kind of partial emulation for all platforms that implement ESIL uplifting (x86/x86_64, ARM, arm64, MIPS, powerpc, sparc, AVR, 8051, Gameboy, ...).

One of the most common usages of such emulation is to calculate indirect jumps and conditional jumps.

To see the ESIL representation of the program one can use the `ao` command or enable the `asm.esil` configuration variable, to check if the program uplifted correctly, and to grasp how ESIL works:

```
[0x00001660]> pdf
. (fcn) fcn.00001660 40
|   fcn.00001660 ();
|       ; CALL XREF from 0x00001713 (entry2.fini)
|   0x00001660  lea rdi, obj.__progrname      ; 0x207220
|   0x00001667  push rbp
|   0x00001668  lea rax, obj.__progrname      ; 0x207220
```

```

|      0x0000166f  cmp rax, rdi
|      0x00001672  mov rbp, rsp
|  .-< 0x00001675  je 0x1690
| |      0x00001677  mov rax, qword [reloc._ITM_deregisterTMCloneTable] ; [0x206fd8:8]=0
| |      0x0000167e  test rax, rax
| .--< 0x00001681  je 0x1690
| ||     0x00001683  pop rbp
| ||     0x00001684  jmp rax
| ``-> 0x00001690  pop rbp
|      0x00001691  ret
[0x00001660]> e asm.esil=true
[0x00001660]> pdf
. (fcn) fcn.00001660 40
|      fcn.00001660 ();
|      ; CALL XREF from 0x00001713 (entry2.fini)
|      0x00001660  0x205bb9,rip,+,rdi,=
|      0x00001667  rbp,8,rsp,-,rsp,=[8]
|      0x00001668  0x205bb1,rip,+,rax,=
|      0x0000166f  rdi,rax,==,$z,zf,=,$b64,cf,=,$p,pf,=,$s,sf,=,$o,of,=
|      0x00001672  rsp,rbp,=
|  .-< 0x00001675  zf,{,5776,rip,=,}
| |      0x00001677  0x20595a,rip,+, [8],rax,=
| |      0x0000167e  0,rax,rax,&,==,$z,zf,=,$p,pf,=,$s,sf,=,$o,cf,=,$o,of,=
| .--< 0x00001681  zf,{,5776,rip,=,}
| ||     0x00001683  rsp,[8],rbp,=,8,rsp,+=
| ||     0x00001684  rax,rip,=
| ``-> 0x00001690  rsp,[8],rbp,=,8,rsp,+=
|      0x00001691  rsp,[8],rip,=,8,rsp,+=

```

To manually setup the ESIL imprecise emulation you need to run this command sequence:

- **aei** to initialize ESIL VM
- **aeim** to initialize ESIL VM memory (stack)
- **aeip** to set the initial ESIL VM IP (instruction pointer)
- a sequence of **aer** commands to set the initial register values.

While performing emulation, please remember, that ESIL VM cannot emulate external calls or system calls, along with SIMD instructions. Thus the most common scenario is to emulate only a small chunk of the code, like encryption/decryption, unpacking or calculating something.

After we successfully set up the ESIL VM we can interact with it like with a usual debugging mode. Commands interface for ESIL VM is almost identical to the debugging one:

- **aes** to step (or **s** key in visual mode)
- **aesi** to step over the function calls

- **aesu** <address> to step until some specified address
- **aesue** <ESIL expression> to step until some specified ESIL expression met
- **aec** to continue until break (Ctrl-C), this one is rarely used though, due to the omnipresence of external calls

In visual mode, all of the debugging hotkeys will work also in ESIL emulation mode.

Along with usual emulation, there is a possibility to record and replay mode:

- **aets** to list all current ESIL R&R sessions
- **aets+** to create a new one
- **aesb** to step back in the current ESIL R&R session

More about this operation mode you can read in Reverse Debugging chapter.

Emulation in analysis loop

Apart from the manual emulation mode, it can be used automatically in the analysis loop. For example, the **aaaa** command performs the ESIL emulation stage along with others. To disable or enable its usage you can use **anal.esil** configuration variable. There is one more important option, though setting it might be quite dangerous, especially in the case of malware - **emu.write** which allows ESIL VM to modify memory. Sometimes it is required though, especially in the process of deobfuscating or unpacking code.

To show the process of emulation you can set **asm.emu** variable, which will show calculated register and memory values in disassembly comments:

```
[0x00001660]> e asm.emu=true
[0x00001660]> pdf
. (fcn) fcn.00001660 40
|   fcn.00001660 ();
|       ; CALL XREF from 0x00001713 (entry2.fini)
|   0x00001660 lea rdi, obj.__prognam ; 0x207220 ; rdi=0x207220 -> 0x464c457f
|   0x00001667 push rbp                ; rsp=0xfffffffffffffffff8
|   0x00001668 lea rax, obj.__prognam ; 0x207220 ; rax=0x207220 -> 0x464c457f
|   0x0000166f cmp rax, rdi            ; zf=0x1 -> 0x2464c45 ; cf=0x0 ; pf=0x1 -> 0x2464c45
|   0x00001672 mov rbp, rsp           ; rbp=0xfffffffffffffffff8
|   .-< 0x00001675 je 0x1690            ; rip=0x1690 -> 0x1f0fc35d ; likely
| |   0x00001677 mov rax, qword [reloc._ITM_deregisterTMCloneTable] ; [0x206fd8:8]=0 ; rax=0
| |   0x0000167e test rax, rax         ; zf=0x1 -> 0x2464c45 ; pf=0x1 -> 0x2464c45 ; sf=0
| .--< 0x00001681 je 0x1690            ; rip=0x1690 -> 0x1f0fc35d ; likely
| ||  0x00001683 pop rbp              ; rbp=0xfffffffffffffffff8 -> 0x4c457fff ; rsp=0x0
| ||  0x00001684 jmp rax              ; rip=0x0 ..
| ``-> 0x00001690 pop rbp             ; rbp=0x10102464c457f ; rsp=0x8 -> 0x464c457f
| `   0x00001691 ret                 ; rip=0x0 ; rsp=0x10 -> 0x3e0003
```


Note here **likely** comments, which indicates that ESIL emulation predicted for particular conditional jump to happen.

Apart from the basic ESIL VM setup, you can change the behavior with other options located in `emu.` and `esil.` configuration namespaces.

For manipulating ESIL working with memory and stack you can use the following options:

- `esil.stack` to enable or disable temporary stack for `asm.emu` mode
- `esil.stack.addr` to set stack address in ESIL VM (like `aeim` command)
- `esil.stack.size` to set stack size in ESIL VM (like `aeim` command)
- `esil.stack.depth` limits the number of PUSH operations into the stack
- `esil.romem` specifies read-only access to the ESIL memory
- `esil.fillstack` and `esil.stack.pattern` allows you to use a various pattern for filling ESIL VM stack upon initialization
- `esil.nonull` when set stops ESIL execution upon NULL pointer read or write.

Graph commands

When analyzing data it is usually handy to have different ways to represent it in order to get new perspectives to allow the analyst to understand how different parts of the program interact.

Representing basic block edges, function calls, string references as graphs show a very clear view of this information.

Radare2 supports various types of graph available through commands starting with `ag`:

```
[0x00005000]> ag?
|Usage: ag<graphtype><format> [addr]
| Graph commands:
| aga[format]           Data references graph
| agA[format]           Global data references graph
| agc[format]           Function callgraph
| agC[format]           Global callgraph
| agd[format] [fcn addr] Diff graph
| agf[format]           Basic blocks function graph
| agi[format]           Imports graph
| agr[format]           References graph
| agR[format]           Global references graph
| agx[format]           Cross references graph
| agg[format]           Custom graph
| ag-                   Clear the custom graph
| agn[?] title body     Add a node to the custom graph
| age[?] title1 title2  Add an edge to the custom graph
```

Output formats:

<blank>	Ascii art
*	r2 commands
d	Graphviz dot
g	Graph Modelling Language (gml)
j	json ('J' for formatted disassembly)
k	SDB key-value
t	Tiny ascii art
v	Interactive ascii art
w [path]	Write to path or display graph image (see graph.gv.format and graph.gv.view)

The structure of the commands is as follows: `ag <graph type> <output format>`.

For example, `agid` displays the imports graph in dot format, while `aggj` outputs the custom graph in JSON format.

Here's a short description for every output format available:

Ascii Art ** (e.g. agf)

Displays the graph directly to stdout using ASCII art to represent blocks and edges.

Warning: displaying large graphs directly to stdout might prove to be computationally expensive and will make r2 not responsive for some time. In case of a doubt, prefer using the interactive view (explained below).

Interactive Ascii Art (e.g. agfv)

Displays the ASCII graph in an interactive view similar to `VV` which allows to move the screen, zoom in / zoom out, ...

Tiny Ascii Art (e.g. agft)

Displays the ASCII graph directly to stdout in tiny mode (which is the same as reaching the maximum zoom out level in the interactive view).

Graphviz dot (e.g. agfd)

Prints the dot source code representing the graph, which can be interpreted by programs such as graphviz or online viewers like this

JSON (e.g. agfj)

Prints a JSON string representing the graph.

- In case of the **f** format (basic blocks of function), it will have detailed information about the function and will also contain the disassembly of the function (use **J** format for the formatted disassembly).
- In all other cases, it will only have basic information about the nodes of the graph (id, title, body, and edges).

Graph Modelling Language (e.g. **agfg**)

Prints the GML source code representing the graph, which can be interpreted by programs such as yEd

SDB key-value (e.g. **agfk**)

Prints key-value strings representing the graph that was stored by sdb (radare2's string database).

R2 custom graph commands (e.g. **agf***)

Prints r2 commands that would recreate the desired graph. The commands to construct the graph are **agn** [**title**] [**body**] to add a node and **age** [**title1**] [**title2**] to add an edge. The [**body**] field can be expressed in base64 to include special formatting (such as newlines).

To easily execute the printed commands, it is possible to prepend a dot to the command (**.agf***).

Web / image (e.g. **agfw**)

Radare2 will convert the graph to dot format, use the **dot** program to convert it to a **.gif** image and then try to find an already installed viewer on your system (**xdg-open**, **open**, ...) and display the graph there.

The extension of the output image can be set with the **graph.extension** config variable. Available extensions are **png**, **jpg**, **gif**, **pdf**, **ps**.

*Note: for particularly large graphs, the most recommended extension is **svg** as it will produce images of much smaller size*

If **graph.web** config variable is enabled, radare2 will try to display the graph using the browser (*this feature is experimental and unfinished, and disabled by default.*)

Data and Code Analysis

Radare2 has a very rich set of commands and configuration options to perform data and code analysis, to extract useful information from a binary, like pointers, string references, basic blocks, opcode data, jump targets, cross references and much more. These operations are handled by the **a** (analyze) command family:

```

|Usage: a[abdefFghoprxtc] [...]
| aa[?]          analyze all (fcns + bbs) (aa0 to avoid sub renaming)
| a8 [hexpairs]  analyze bytes
| ab[b] [addr]   analyze block at given address
| abb [len]      analyze N basic blocks in [len] (section.size by default)
| abt [addr]     find paths in the bb function graph from current offset to given address
| ac [cycles]    analyze which op could be executed in [cycles]
| ad[?]          analyze data trampoline (wip)
| ad [from] [to] analyze data pointers to (from-to)
| ae[?] [expr]   analyze opcode eval expression (see ao)
| af[?]          analyze Functions
| aF             same as above, but using anal.depth=1
| ag[?] [options] draw graphs in various formats
| ah[?]          analysis hints (force opcode size, ...)
| ai [addr]      address information (show perms, stack, heap, ...)
| an [name] [@addr] show/rename/create whatever flag/function is used at addr
| ao[?] [len]    analyze Opcodes (or emulate it)
| a0[?] [len]    Analyze N instructions in M bytes
| ap             find prelude for current offset
| ar[?]          like 'dr' but for the esil vm. (registers)
| as[?] [num]    analyze syscall using dbg.reg
| av[?] [.]      show vtables
| ax[?]          manage refs/xrefs (see also afx?)

```

In fact, a namespace is one of the biggest in radare2 tool and allows to control very different parts of the analysis:

- Code flow analysis
- Data references analysis
- Using loaded symbols
- Managing different type of graphs, like CFG and call graph
- Manage variables
- Manage types
- Emulation using ESIL VM
- Opcode introspection
- Objects information, like virtual tables

Symbols

Radare2 automatically parses available imports and exports sections in the binary, moreover, it can load additional debugging information if present. Two main formats are supported: DWARF and PDB (for Windows binaries). Note that, unlike many tools radare2 doesn't rely on Windows API to parse PDB files, thus they can be loaded on any other supported platform - e.g. Linux or OS X.

DWARF debug info loads automatically by default because usually it's stored

right in the executable file. PDB is a bit of a different beast - it is always stored as a separate binary, thus the different logic of handling it.

At first, one of the common scenarios is to analyze the file from Windows distribution. In this case, all PDB files are available on the Microsoft server, which is by default is in options. See all pdb options in radare2:

```
pdb.autoload = 0
pdb.extract = 1
pdb.server = https://msdl.microsoft.com/download/symbols
pdb.useragent = Microsoft-Symbol-Server/6.11.0001.402
```

Using the variable `pdb.server` you can change the address where radare2 will try to download the PDB file by the GUID stored in the executable header. You can make use of multiple symbol servers by separating each URL with a semi-colon:

```
e pdb.server = https://msdl.microsoft.com/download/symbols;https://symbols.mozilla.org/
```

On Windows, you can also use local network share paths (UNC paths) as symbol servers.

Usually, there is no reason to change default `pdb.useragent`, but who knows where could it be handy?

Because those PDB files are stored as “cab” archives on the server, `pdb.extract=1` says to automatically extract them.

Note that for the automatic downloading to work you need “cabextract” tool, and wget/curl installed.

Sometimes you don’t need to do that from the radare2 itself, thus - two handy rabin2 options:

```
-P          show debug/pdb information
-PP         download pdb file for binary
```

where `-PP` automatically downloads the pdb for the selected binary, using those `pdb.*` config options. `-P` will dump the contents of the PDB file, which is useful sometimes for a quick understanding of the symbols stored in it.

Apart from the basic scenario of just opening a file, PDB information can be additionally manipulated by the `id` commands:

```
[0x000051c0]> id?
|Usage: id Debug information
| Output mode:
| '*'          Output in radare commands
| id           Source lines
| idp [file.pdb] Load pdb file information
| idpi [file.pdb] Show pdb file information
| idpd         Download pdb file on remote server
```

Where `idpi` is basically the same as `rabin2 -P`. Note that `idp` can be also used not only in the static analysis mode, but also in the debugging mode, even if connected via WinDbg.

For simplifying the loading PDBs, especially for the processes with many linked DLLs, radare2 can autoloading all required PDBs automatically - you need just set the `e pdb.autoload=true` option. Then if you load some file in debugging mode in Windows, using `r2 -d file.exe` or `r2 -d 2345` (attach to pid 2345), all related PDB files will be loaded automatically.

DWARF information loading, on the other hand, is completely automated. You don't need to run any commands/change any options:

```
r2 `which rabin2`
[0x00002437 8% 300 /usr/local/bin/rabin2]> pd $r
0x00002437 jne 0x2468 ;[1]
0x00002439 cmp qword reloc.__cxa_finalize_224, 0
0x00002441 push rbp
0x00002442 mov rbp, rsp
0x00002445 je 0x2453 ;[2]
0x00002447 lea rdi, obj.__dso_handle ; 0x207c40 ; "@| "
0x0000244e call 0x2360 ;[3]
0x00002453 call sym.deregister_tm_clones ;[4]
0x00002458 mov byte [obj.completed.6991], 1 ; obj.__TMC_END__ ; [0x2082f0:1]=0
0x0000245f pop rbp
0x00002460 ret
0x00002461 nop dword [rax]
0x00002468 ret
0x0000246a nop word [rax + rax]
;-- entry1.init:
;-- frame_dummy:
0x00002470 push rbp
0x00002471 mov rbp, rsp
0x00002474 pop rbp
0x00002475 jmp sym.register_tm_clones ;[5]
;-- blob_version:
0x0000247a push rbp ; ../blob/version.c:18
0x0000247b mov rbp, rsp
0x0000247e sub rsp, 0x10
0x00002482 mov qword [rbp - 8], rdi
0x00002486 mov eax, 0x32 ; ../blob/version.c:24 ; '2'
0x0000248b test al, al ; ../blob/version.c:19
0x0000248d je 0x2498 ;[6]
0x0000248f lea rax, str.2.0.1_182_gf1aa3aa4d ; 0x60b8 ; "2.0.1-182-gf1aa3aa4d"
0x00002496 jmp 0x249f ;[7]
0x00002498 lea rax, 0x000060cd
0x0000249f mov rsi, qword [rbp - 8]
```

```

0x000024a3  mov r8, rax
0x000024a6  mov ecx, 0x40                ; section_end.ehdr
0x000024ab  mov edx, 0x40c0
0x000024b0  lea rdi, str._s_2.1.0_git__d__linux_x86__d_git._s_n ; 0x60d0 ; "%s 2.1.0-git %
0x000024b7  mov eax, 0
0x000024bc  call 0x2350                  ; [8]
0x000024c1  mov eax, 0x66                ; ../blob/version.c:25 ; 'f'
0x000024c6  test al, al
0x000024c8  je 0x24d6                    ; [9]
0x000024ca  lea rdi, str.commit:_f1aa3aa4d2599c1ad60e3ecbe5f4d8261b282385_build:_2017_11_06
0x000024d1  call sym.imp.puts            ; [?]
0x000024d6  mov eax, 0                    ; ../blob/version.c:28
0x000024db  leave                         ; ../blob/version.c:29
0x000024dc  ret
;-- rabin_show_help:
0x000024dd  push rbp                      ; ../rabin2.c:27

```

As you can see, it loads function names and source line information.

Syscalls

Radare2 allows manual search for assembly code looking like a syscall operation. For example on ARM platform usually they are represented by the `svc` instruction, on the others can be a different instructions, e.g. `syscall` on x86 PC.

```

[0x0001ece0]> /ad/ svc
...
0x000187c2  # 2: svc 0x76
0x000189ea  # 2: svc 0xa9
0x00018a0e  # 2: svc 0x82
...

```

Syscalls detection is driven by `asm.os`, `asm.bits`, and `asm.arch`. Be sure to setup those configuration options accordingly. You can use `asl` command to check if syscalls' support is set up properly and as you expect. The command lists syscalls supported for your platform.

```

[0x0001ece0]> asl
...
sd_softdevice_enable = 0x80.16
sd_softdevice_disable = 0x80.17
sd_softdevice_is_enabled = 0x80.18
...

```

If you setup ESIL stack with `aei` or `aeim`, you can use `/as` command to search the addresses where particular syscalls were found and list them.

```
[0x0001ece0]> aei
[0x0001ece0]> /as
0x000187c2 sd_ble_gap_disconnect
0x000189ea sd_ble_gatts_sys_attr_set
0x00018a0e sd_ble_gap_sec_info_reply
...
```

To reduce searching time it is possible to restrict the searching range for only executable segments or sections with `/as @e:search.in=io.maps.x`

Using the ESIL emulation radare2 can print syscall arguments in the disassembly output. To enable the linear (but very rough) emulation use `asm.emu` configuration variable:

```
[0x0001ece0]> e asm.emu=true
[0x0001ece0]> s 0x000187c2
[0x000187c2]> pdf~svc
    0x000187c2    svc 0x76    ; 118 = sd_ble_gap_disconnect
[0x000187c2]>
```

In case of executing `aae` (or `aaaa` which calls `aae`) command radare2 will push found syscalls to a special `syscall. flagspace`, which can be useful for automation purpose:

```
[0x000187c2]> fs
0    0 * imports
1    0 * symbols
2 1523 * functions
3  420 * strings
4  183 * syscalls
[0x000187c2]> f~syscall
...
0x000187c2 1 syscall.sd_ble_gap_disconnect.0
0x000189ea 1 syscall.sd_ble_gatts_sys_attr_set
0x00018a0e 1 syscall.sd_ble_gap_sec_info_reply
...
```

It also can be interactively navigated through within HUD mode (`V_`)

```
0> syscall.sd_ble_gap_disconnect
- 0x000187b2  syscall.sd_ble_gap_disconnect
    0x000187c2  syscall.sd_ble_gap_disconnect.0
    0x00018a16  syscall.sd_ble_gap_disconnect.1
    0x00018b32  syscall.sd_ble_gap_disconnect.2
    0x0002ac36  syscall.sd_ble_gap_disconnect.3
```

When debugging in radare2, you can use `dcs` to continue execution until the next syscall. You can also run `dcs*` to trace all syscalls.

```
[0xf7fb9120]> dcs*
```



```

Running child until syscalls:-1
child stopped with signal 133
--> SN 0xf7fd3d5b syscall 45 brk (0xffffffffda)
child stopped with signal 133
--> SN 0xf7fd28f3 syscall 384 arch_prctl (0xffffffffda 0x3001)
child stopped with signal 133
--> SN 0xf7fc81b2 syscall 33 access (0xffffffffda 0xf7fd8bf1)
child stopped with signal 133

```

radare2 also has a syscall name to syscall number utility. You can return the syscall name of a given syscall number or vice versa, without leaving the shell.

```

[0x08048436]> asl 1
exit
[0x08048436]> asl write
4
[0x08048436]> ask write
0x80,4,3,iZi

```

See `as?` for more information about the utility.

Types

Radare2 supports the C-syntax data types description. Those types are parsed by a C11-compatible parser and stored in the internal SDB, thus are introspectable with `k` command.

Most of the related commands are located in `t` namespace:

```

[0x00000000]> t?
| Usage: t      # cparse types commands
| t              List all loaded types
| tj             List all loaded types as json
| t <type>       Show type in 'pf' syntax
| t*            List types info in r2 commands
| t- <name>      Delete types by its name
| t-*           Remove all types
| tail [filename] Output the last part of files
| tc [type.name] List all/given types in C output format
| te[?]         List all loaded enums
| td[?] <string> Load types from string
| tf            List all loaded functions signatures
| tk <sdb-query> Perform sdb query
| tl[?]         Show/Link type to an address
| tn[?] [-][addr] manage noreturn function attributes and marks
| to -          Open cfg.editor to load types
| to <path>     Load types from C header file

```

toe [type.name]	Open cfg.editor to edit types
tos <path>	Load types from parsed Sdb database
tp <type> [addr varname]	cast data at <address> to <type> and print it (XXX: type can co
tpv <type> @ [value]	Show offset formatted for given type
tpx <type> <hexpairs>	Show value for type with specified byte sequence (XXX: type can
ts[?]	Print loaded struct types
tu[?]	Print loaded union types
tx[f?]	Type xrefs
tt[?]	List all loaded typedefs

Note that the basic (atomic) types are not those from C standard - not `char`, `_Bool`, or `short`. Because those types can be different from one platform to another, radare2 uses **definite** types like as `int8_t` or `uint64_t` and will convert `int` to `int32_t` or `int64_t` depending on the binary or debuggee platform/compiler.

Basic types can be listed using `t` command. For the structured types you need to use `ts`, for unions use `tu` and for enums — `te`.

```
[0x00000000]> t
char
char *
double
float
gid_t
int
int16_t
int32_t
int64_t
int8_t
long
long long
pid_t
short
size_t
uid_t
uint16_t
uint32_t
uint64_t
uint8_t
unsigned char
unsigned int
unsigned short
void *
```

Loading types

There are three easy ways to define a new type: * Directly from the string using `td` command * From the file using `to <filename>` command * Open an `$EDITOR` to type the definitions in place using `to -`

```
[0x00000000]> "td struct foo {char* a; int b;}"
[0x00000000]> cat ~/radare2-regressions/bins/headers/s3.h
struct S1 {
    int x[3];
    int y[4];
    int z;
};
[0x00000000]> to ~/radare2-regressions/bins/headers/s3.h
[0x00000000]> ts
foo
S1
```

Also note there is a `config` option to specify include directories for types parsing

```
[0x00000000]> e? dir.types
dir.types: Default path to look for cparse type files
[0x00000000]> e dir.types
/usr/include
```

Printing types

Notice below we have used `ts` command, which basically converts the C type description (or to be precise it's SDB representation) into the sequence of `pf` commands. See more about print format.

The `tp` command uses the `pf` string to print all the members of type at the current offset/given address:

```
[0x00000000]> "td struct foo {char* a; int b;}"
[0x00000000]> wx 68656c6c6f00c00000
[0x00000000]> wz world @ 0x00000010 ; wx 17 @ 0x00000016
[0x00000000]> px
[0x00000000]> ts foo
pf zd a b
[0x00000000]> tp foo
a : 0x00000000 = "hello"
b : 0x00000006 = 12
[0x00000000]> tp foo @ 0x00000010
a : 0x00000010 = "world"
b : 0x00000016 = 23
```

Also, you could fill your own data into the struct and print it using `tpx` command

```
[0x00000000]> tpx foo 414243440010000000
```

```
a : 0x00000000 = "ABCD"
b : 0x00000005 = 16
```

Linking Types

The `tp` command just performs a temporary cast. But if we want to link some address or variable with the chosen type, we can use `tl` command to store the relationship in SDB.

```
[0x000051c0]> tl S1 = 0x51cf
[0x000051c0]> tll
(S1)
x : 0x000051cf = [ 2315619660, 1207959810, 34803085 ]
y : 0x000051db = [ 2370306049, 4293315645, 3860201471, 4093649307 ]
z : 0x000051eb = 4464399
```

Moreover, the link will be shown in the disassembly output or visual mode:

```
[0x000051c0 15% 300 /bin/ls]> pd $r @ entry0
;-- entry0:
0x000051c0      xor ebp, ebp
0x000051c2      mov r9, rdx
0x000051c5      pop rsi
0x000051c6      mov rdx, rsp
0x000051c9      and rsp, 0xfffffffffffffff0
0x000051cd      push rax
0x000051ce      push rsp
(S1)
x : 0x000051cf = [ 2315619660, 1207959810, 34803085 ]
y : 0x000051db = [ 2370306049, 4293315645, 3860201471, 4093649307 ]
z : 0x000051eb = 4464399
0x000051f0      lea rdi, loc._edata          ; 0x21f248
0x000051f7      push rbp
0x000051f8      lea rax, loc._edata          ; 0x21f248
0x000051ff      cmp rax, rdi
0x00005202      mov rbp, rsp
```

Once the struct is linked, radare2 tries to propagate structure offset in the function at current offset, to run this analysis on whole program or at any targeted functions after all structs are linked you have `aat` command:

```
[0x00000000]> aa?
| aat [fcn]          Analyze all/given function to convert immediate to linked structure of
```

Note sometimes the emulation may not be accurate, for example as below :

```
|0x000006da  push rbp
|0x000006db  mov rbp, rsp
|0x000006de  sub rsp, 0x10
```

```
|0x000006e2  mov edi, 0x20                ; "@"
|0x000006e7  call sym.imp.malloc          ; void *malloc(size_t size)
|0x000006ec  mov qword [local_8h], rax
|0x000006f0  mov rax, qword [local_8h]
```

The return value of `malloc` may differ between two emulations, so you have to set the hint for return value manually using `ahr` command, so run `t1` or `aat` command after setting up the return value hint.

```
[0x000006da]> ah?
| ahr val          set hint for return value of a function
```

Structure Immediates

There is one more important aspect of using types in radare2 - using `aht` you can change the immediate in the opcode to the structure offset. Lets see a simple example of [R]SI-relative addressing

```
[0x000052f0]> pd 1
0x000052f0      mov rax, qword [rsi + 8]      ; [0x8:8]=0
```

Here 8 - is some offset in the memory, where `rsi` probably holds some structure pointer. Imagine that we have the following structures

```
[0x000052f0]> "td struct ms { char b[8]; int member1; int member2; };"
[0x000052f0]> "td struct ms1 { uint64_t a; int member1; };"
[0x000052f0]> "td struct ms2 { uint16_t a; int64_t b; int member1; };"
```

Now we need to set the proper structure member offset instead of 8 in this instruction. At first, we need to list available types matching this offset:

```
[0x000052f0]> ahts 8
ms.member1
ms1.member1
```

Note, that `ms2` is not listed, because it has no members with offset 8. After listing available options we can link it to the chosen offset at the current address:

```
[0x000052f0]> aht ms1.member1
[0x000052f0]> pd 1
0x000052f0      488b4608      mov rax, qword [rsi + ms1.member1]      ; [0x8:8]=0
```

Managing enums

- Printing all fields in enum using `te` command

```
[0x00000000]> "td enum Foo {COW=1,BAR=2};"
[0x00000000]> te Foo
COW = 0x1
BAR = 0x2
```

- Finding matching enum member for given bitfield and vice-versa

```
[0x00000000]> te Foo 0x1
COW
[0x00000000]> teb Foo COW
0x1
```

Internal representation

To see the internal representation of the types you can use `tk` command:

```
[0x000051c0]> tk~S1
S1=struct
struct.S1=x,y,z
struct.S1.x=int32_t,0,3
struct.S1.x.meta=4
struct.S1.y=int32_t,12,4
struct.S1.y.meta=4
struct.S1.z=int32_t,28,0
struct.S1.z.meta=0
[0x000051c0]>
```

Defining primitive types requires an understanding of basic `pf` formats, you can find the whole list of format specifier in `pf??`:

format	explanation
b	byte (unsigned)
c	char (signed byte)
d	0x%%08x hexadecimal value (4 bytes)
f	float value (4 bytes)
i	%%i integer value (4 bytes)
o	0x%%08o octal value (4 byte)
p	pointer reference (2, 4 or 8 bytes)
q	quadword (8 bytes)
s	32bit pointer to string (4 bytes)
S	64bit pointer to string (8 bytes)
t	UNIX timestamp (4 bytes)
T	show Ten first bytes of buffer
u	uleb128 (variable length)
w	word (2 bytes unsigned short in hex)
x	0x%%08x hex value and flag (fd @ addr)
X	show formatted hexpairs
z	\0 terminated string
Z	\0 terminated wide string

there are basically 3 mandatory keys for defining basic data types: `X=type`
`type.X=format_specifier type.X.size=size_in_bits` For example, let's de-

fine UNIT, according to Microsoft documentation UINT is just equivalent of standard C `unsigned int` (or `uint32_t` in terms of TCC engine). It will be defined as:

```
UINT=type
type.UINT=d
type.UINT.size=32
```

Now there is an optional entry:

```
X.type.pointto=Y
```

This one may only be used in case of pointer `type.X=p`, one good example is `LPFILETIME` definition, it is a pointer to `_FILETIME` which happens to be a structure. Assuming that we are targeting only 32-bit windows machine, it will be defined as the following:

```
LPFILETIME=type
type.LPFILETIME=p
type.LPFILETIME.size=32
type.LPFILETIME.pointto=_FILETIME
```

This last field is not mandatory because sometimes the data structure internals will be proprietary, and we will not have a clean representation for it.

There is also one more optional entry:

```
type.UINT.meta=4
```

This entry is for integration with C parser and carries the type class information: integer size, signed/unsigned, etc.

Structures

Those are the basic keys for structs (with just two elements):

```
X=struct
struct.X=a,b
struct.X.a=a_type,a_offset,a_number_of_elements
struct.X.b=b_type,b_offset,b_number_of_elements
```

The first line is used to define a structure called `X`, the second line defines the elements of `X` as comma separated values. After that, we just define each element info.

For example. we can have a struct like this one:

```
struct _FILETIME {
    DWORD dwLowDateTime;
    DWORD dwHighDateTime;
}
```

assuming we have `DWORD` defined, the struct will look like this

```

_FILETIME=struct
struct._FILETIME=dwLowDateTime,dwHighDateTime
struct._FILETIME.dwLowDateTime=DWORD,0,0
struct._FILETIME.dwHighDateTime=DWORD,4,0

```

Note that the number of elements field is used in case of arrays only to identify how many elements are in arrays, other than that it is zero by default.

Unions

Unions are defined exactly like structs the only difference is that you will replace the word `struct` with the word `union`.

Function prototypes

Function prototypes representation is the most detail oriented and the most important one of them all. Actually, this is the one used directly for type matching

```

X=func
func.X.args=NumberOfArgs
func.x.arg0=Arg_type,arg_name
.
.
.
func.X.ret=Return_type
func.X.cc=calling_convention

```

It should be self-explanatory. Let's do `strncasecmp` as an example for x86 arch for Linux machines. According to man pages, `strncasecmp` is defined as the following:

```
int strncasecmp(const char *s1, const char *s2, size_t n);
```

When converting it into its sdb representation it will look like the following:

```

strncasecmp=func
func.strncasecmp.args=3
func.strncasecmp.arg0=char *,s1
func.strncasecmp.arg1=char *,s2
func.strncasecmp.arg2=size_t,n
func.strncasecmp.ret=int
func.strncasecmp.cc=cdecl

```

Note that the `.cc` part is optional and if it didn't exist the default calling-convention for your target architecture will be used instead. There is one extra optional key

```
func.x.noreturn=true/false
```


This key is used to mark functions that will not return once called, such as `exit` and `_exit`.

Managing variables

Radare2 allows managing local variables, no matter their location, stack or registers. The variables' auto analysis is enabled by default but can be disabled with `anal.vars` configuration option.

The main variables commands are located in `afv` namespace:

```
Usage: afv [rbs]
| afv*                output r2 command to add args/locals to flagspace
| afv-([name])        remove all or given var
| afv=               list function variables and arguments with disasm refs
| afva              analyze function arguments/locals
| afvb[?]           manipulate bp based arguments/locals
| afvd name          output r2 command for displaying the value of args/locals in
| afvf              show BP relative stackframe variables
| afvn [new_name] ([old_name]) rename argument/local
| afvr[?]           manipulate register based arguments
| afvR [varname]     list addresses where vars are accessed (READ)
| afvs[?]           manipulate sp based arguments/locals
| afvt [name] [new_type] change type for given argument/local
| afvW [varname]     list addresses where vars are accessed (WRITE)
| afvx              show function variable xrefs (same as afvR+afvW)
```

`afvr`, `afvb` and `afvs` commands are uniform but allow manipulation of register-based arguments and variables, BP/FP-based arguments and variables, and SP-based arguments and variables respectively. If we check the help for `afvr` we will get the way two others commands works too:

```
|Usage: afvr [reg] [type] [name]
| afvr                list register based arguments
| afvr*              same as afvr but in r2 commands
| afvr [reg] [name] ([type]) define register arguments
| afvrj              return list of register arguments in JSON format
| afvr- [name]        delete register arguments at the given index
| afvrg [reg] [addr]  define argument get reference
| afvrs [reg] [addr]  define argument set reference
```

Like many other things variables detection is performed by radare2 automatically, but results can be changed with those arguments/variables control commands. This kind of analysis relies heavily on preloaded function prototypes and the calling-convention, thus loading symbols can improve it. Moreover, after changing something we can rerun variables analysis with `afva` command. Quite often variables analysis is accompanied with types analysis, see `afta` command.

The most important aspect of reverse engineering - naming things. Of course, you can rename variable too, affecting all places it was referenced. This can be achieved with `afvn` for *any* type of argument or variable. Or you can simply remove the variable or argument with `afv-` command.

As mentioned before the analysis loop relies heavily on types information while performing variables analysis stages. Thus comes next very important command - `afvt`, which allows you to change the type of variable:

```
[0x00003b92]> afvs
var int local_8h @ rsp+0x8
var int local_10h @ rsp+0x10
var int local_28h @ rsp+0x28
var int local_30h @ rsp+0x30
var int local_32h @ rsp+0x32
var int local_38h @ rsp+0x38
var int local_45h @ rsp+0x45
var int local_46h @ rsp+0x46
var int local_47h @ rsp+0x47
var int local_48h @ rsp+0x48
[0x00003b92]> afvt local_10h char*
[0x00003b92]> afvs
var int local_8h @ rsp+0x8
var char* local_10h @ rsp+0x10
var int local_28h @ rsp+0x28
var int local_30h @ rsp+0x30
var int local_32h @ rsp+0x32
var int local_38h @ rsp+0x38
var int local_45h @ rsp+0x45
var int local_46h @ rsp+0x46
var int local_47h @ rsp+0x47
var int local_48h @ rsp+0x48
```

Less commonly used feature, which is still under heavy development - distinction between variables being read and written. You can list those being read with `afvR` command and those being written with `afvW` command. Both commands provide a list of the places those operations are performed:

```
[0x00003b92]> afvR
local_48h  0x48ee
local_30h  0x3c93,0x520b,0x52ea,0x532c,0x5400,0x3cfb
local_10h  0x4b53,0x5225,0x53bd,0x50cc
local_8h   0x4d40,0x4d99,0x5221,0x53b9,0x50c8,0x4620
local_28h  0x503a,0x51d8,0x51fa,0x52d3,0x531b
local_38h
local_45h  0x50a1
local_47h
local_46h
```

```

local_32h  0x3cb1
[0x00003b92]> afvW
local_48h  0x3adf
local_30h  0x3d3e,0x4868,0x5030
local_10h  0x3d0e,0x5035
local_8h   0x3d13,0x4d39,0x5025
local_28h  0x4d00,0x52dc,0x53af,0x5060,0x507a,0x508b
local_38h  0x486d
local_45h  0x5014,0x5068
local_47h  0x501b
local_46h  0x5083
local_32h
[0x00003b92]>

```

Type inference

The type inference for local variables and arguments is well integrated with the command `afta`.

Let's see an example of this with a simple `hello_world` binary

```

[0x000007aa]> pdf
|          ;-- main:
/ (fcn) sym.main 157
| sym.main ();
| ; var int local_20h @ rbp-0x20
| ; var int local_1ch @ rbp-0x1c
| ; var int local_18h @ rbp-0x18
| ; var int local_10h @ rbp-0x10
| ; var int local_8h @ rbp-0x8
| ; DATA XREF from entry0 (0x6bd)
| 0x000007aa  push rbp
| 0x000007ab  mov rbp, rsp
| 0x000007ae  sub rsp, 0x20
| 0x000007b2  lea rax, str.Hello          ; 0x8d4 ; "Hello"
| 0x000007b9  mov qword [local_18h], rax
| 0x000007bd  lea rax, str.r2_folks      ; 0x8da ; " r2-folks"
| 0x000007c4  mov qword [local_10h], rax
| 0x000007c8  mov rax, qword [local_18h]
| 0x000007cc  mov rdi, rax
| 0x000007cf  call sym.imp.strlen        ; size_t strlen(const char *s)

```

- After applying `afta`

```

[0x000007aa]> afta
[0x000007aa]> pdf
| ;-- main:
| ;-- rip:

```

```

/ (fcn) sym.main 157
| sym.main ();
| ; var size_t local_20h @ rbp-0x20
| ; var size_t size @ rbp-0x1c
| ; var char *src @ rbp-0x18
| ; var char *s2 @ rbp-0x10
| ; var char *dest @ rbp-0x8
| ; DATA XREF from entry0 (0x6bd)
| 0x000007aa  push rbp
| 0x000007ab  mov rbp, rsp
| 0x000007ae  sub rsp, 0x20
| 0x000007b2  lea rax, str.Hello          ; 0x8d4 ; "Hello"
| 0x000007b9  mov qword [src], rax
| 0x000007bd  lea rax, str.r2_folks       ; 0x8da ; " r2-folks"
| 0x000007c4  mov qword [s2], rax
| 0x000007c8  mov rax, qword [src]
| 0x000007cc  mov rdi, rax                ; const char *s
| 0x000007cf  call sym.imp.strlen         ; size_t strlen(const char *s)

```

It also extracts type information from format strings like `printf ("fmt : %s , %u , %d", ...)`, the format specifications are extracted from `anal/d/spec.sdb`

You could create a new profile for specifying a set of format chars depending on different libraries/operating systems/programming languages like this :

```

win=spec
spec.win.u32=unsigned int

```

Then change your default specification to newly created one using this config variable `e anal.spec = win`

For more information about primitive and user-defined types support in radare2 refer to types chapter.

Virtual Tables

There is a basic support of virtual tables parsing (RTTI and others). The most important thing before you start to perform such kind of analysis is to check if the `anal.cpp.abi` option is set correctly, and change if needed.

All commands to work with virtual tables are located in the `av` namespace. Currently, the support is very basic, allowing you only to inspect parsed tables.

```

|Usage: av[?jr*] C++ vtables and RTTI
| av          search for vtables in data sections and show results
| avj         like av, but as json
| av*        like av, but as r2 commands
| avr[j@addr] try to parse RTTI at vtable addr (see anal.cpp.abi)

```

```
| avra[j]          search for vtables and try to parse RTTI at each of them
```

The main commands here are **av** and **avr**. **av** lists all virtual tables found when r2 opened the file. If you are not happy with the result you may want to try to parse virtual table at a particular address with **avr** command. **avra** performs the search and parsing of all virtual tables in the binary, like r2 does during the file opening.

Using r2 with 8051

Features

- ☒ Disassembler
- ☒ Assembler
- ☒ Emulation (esil)
- ☒ Basic address space mapping

Untested

- ☐ Debugger

Missing

- ☐ Full emulation of memory-mapped registers
- ☐ Memory banking for address spaces > 64K
- ☐ Advanced analysis like local variables, function parameters ..
- ☐ More predefined CPU models

r2 configuration

Set architecture to 8051:

```
r2 -a 8051
```

Set cpu to desired model:

```
e asm.cpu = ?
```

After changing the cpu model, run 'aei' to initialize/reset the registers and mapped memory. For example:

```
e asm.cpu = 8051-generic
aei
```

Address spaces and memory mapping

Pseudo-registers are used to control how r2 emulates the multiple address spaces of the 8051. The registers hold the base address where the 8051 memory area is located in r2 address space.

register	address space	comment
<code>_code</code>	CODE	Program memory. Typically located at 0.
<code>_idata</code>	IDATA	256 bytes of internal RAM.
<code>_sfr</code>	SFR	128 bytes for special function registers. <code>_sfr</code> is the base address. Registers start at <code>_sfr+0x80</code> .
<code>_xdata</code>	XDATA	64K of external RAM.
<code>_pdata</code>	PDATA, XREG	MSB of address of 256-byte page in XDATA accessed with 'movx @Ri' op codes.

The registers are initialized based on the selected CPU. See command 'e asm.cpu=?'.

The registers can be viewed and modified with the 'ar' command. When modifying the pseudo-registers or updating 'asm.cpu', memory will be (re)allocated automatically when the analyzer is invoked the next time (e.g. during 'aei'). Use the 'om' command to see the list of allocated memory blocks.

```
[0x00000000]> e asm.cpu=8051.generic
[0x00000000]> aei
[0x00000000]> om
 4 fd: 6 +0x00000000 0x00000000 - 0x0000ffff -rwx
 3 fd: 5 +0x00000000 0x20000000 - 0x2000ffff -rw- xdata
 2 fd: 4 +0x00000000 0x10000180 - 0x100001ff -rw- sfr
 1 fd: 3 +0x00000000 0x10000000 - 0x100000ff -rw- idata
```

Analysis and emulation rely on the address mapping. Setup the pseudo-registers before running analysis, or rerun analysis after updating pseudo-registers.

Address spaces can overlap in r2 memory. This allows emulating 8051 variants that mirror IDATA and SFR into XDATA, or have shared XDATA and CODE address spaces.

For example, the CC2430 from Texas Instruments maps SFR to 0xDF80 and IDATA to 0xFF00 in XDATA memory space. In r2 this can be setup with:

```
ar _sfr = _xdata + 0xdf00
ar _idata = _xdata + 0xff00
```

For overlapping areas, r2 will prioritize smaller memory blocks over larger ones. For example, if IDATA is mapped into XDATA, all r2 operations will use IDATA

in the overlapping addresses. If you want to use XDATA instead, you can delete the offending map with the command 'om-'. See 'om?' for more information.

For using emulation with overlapping code and RAM spaces, the r2 memory holding the firmware must allow write access. This is best achieved with the command 'omf 4 rwx', with 4 being the id of the firmware file's IO map entry. See 'om?' for more information.

Some 8051 variants use memory banking to address memory spaces larger than 64K. Currently, memory banking is not supported by r2.

Tips & tricks

Use pseudo-registers in r2 commands to calculate addresses. For example:

Hex dump of all special function registers:

```
px @ _sfr
```

Write a value to a location in external RAM

```
wx deadbeef @ _xdata + 0x1234
```

Set a flag for a variable stored at 0x20 in internal RAM:

```
f sym.secret @ _idata + 0x20
```

Adding support for new 8051 variants

Follow these steps to add support for new 8051 variants to r2.

1. Clone latest version of radare2
2. In '/libr/anal/p/anal_8051.c' add a new entry to array 'cpu_models[]' to define a name and a memory mapping. The name of the last entry in array must be NULL
3. In 'libr/asm/p/asm_8051.c' append entry with the same name to 'cpus' attribute
4. Compile, test your addition, and submit a pull request

Architectures

This chapter covers architecture specific topics.

Even though most examples in the radare2 book are showing Intel x86 code, radare2 supports an extensive list of computer architectures. The concepts of radare2 apply to all architectures, but there are a few differences in the configuration and usage.

Block Size

The block size determines how many bytes radare2 commands will process when not given an explicit size argument. You can temporarily change the block size by specifying a numeric argument to the print commands. For example `px 20`.

```
[0x00000000]> b?
Usage: b[f] [arg] # Get/Set block size
| b 33      set block size to 33
| b eip+4   numeric argument can be an expression
| b        display current block size
| b+3      increase blocksize by 3
| b-16     decrease blocksize by 16
| b*       display current block size in r2 command
| bf foo    set block size to flag size
| bj       display block size information in JSON
| bm 1M     set max block size
```

The `b` command is used to change the block size:

```
[0x00000000]> b 0x100 # block size = 0x100
[0x00000000]> b+16    # ... = 0x110
[0x00000000]> b-32    # ... = 0xf0
```

The `bf` command is used to change the block size to value specified by a flag. For example, in symbols, the block size of the flag represents the size of the function. To make that work, you have to either run function analysis `af` (which is included in `aa`) or manually seek and define some functions e.g. via `Vd`.

```
[0x00000000]> bf sym.main # block size = sizeof(sym.main)
[0x00000000]> pD @ sym.main # disassemble sym.main
```

You can combine two operations in a single `pdf` command. Except that `pdf` neither uses nor affects global block size.

```
[0x00000000]> pdf @ sym.main # disassemble sym.main
```

Another way around is to use special variables `$FB` and `$FS` which denote Function's Beginning and Size at the current seek. Read more about Usable variables.

```
[0x00000000]> s sym.main + 0x04
[0x00001ec9]> pD @ $FB !$FS # disassemble current function
/ 211: int main (int argc, char **argv, char **envp);
|          0x00001ec5      55                push rbp
|          0x00001ec6      4889e5          mov rbp, rsp
|          0x00001ec9      4881ecc0000000  sub rsp, 0xc0
...
\          0x00001f97      c3                ret
```

Note: don't put space after `!` size designator. See also Command Format.

Comparison Watchers

Watchers are used to record memory at 2 different points in time, then report if and how it changed.

```
[0x00000000]> cw?
Usage: cw [args]  Manage compare watchers; See if and how memory changes
| cw??           Show more info about watchers
| cw addr sz cmd  Add a compare watcher
| cw[*qj] [addr]  Show compare watchers (*=r2 commands, q=quiet)
| cwd [addr]      Delete watcher
| cwr [addr]      Revert watcher
| cwu [addr]      Update watcher
```

Basic watcher usage

First, create one with `cw addr sz cmd`. This will record `sz` bytes at `addr`. The command is stored and used to print the memory when shown.

```
# Create a watcher at 0x0 of size 4 using p8 as the command
[0x00000000]> cw 0 4 p8
```

To record the second state, use `cwu`. Now, when you run `cw`, the watcher will report if the bytes changed and run the command given at creation with the size and address. When an address is an optional argument, the command will apply to all watchers if you don't pass one.

```
# Introduce a change to the block of data we're watching
[0x00000000]> wx 11223344
# Update all watchers
[0x00000000]> cwu
# Show changes
[0x00000000]> cw
0x00000000 modified
11223344
```

You may overwrite any watcher by creating another at the same address. This will discard the existing watcher completely.

```
# Overwrite our existing watcher to display a bistream instead of
# hexpairs, and make the watched area larger
[0x00000000]> cw 0 8 pB
# Check that it's no longer "modified" as this is a new watcher
[0x00000000]> cw
0x00000000
000100010010001000110011010001000000000000000000000000000000000
```

Reverting watcher state

When you create a watcher, the data read from memory is marked as “new”. Updating the watcher with `cwu` will mark this data as “old”, and then read the “new” data.

`cwr` will mark the current “old” state as being “new”, letting you reuse it as your new base state when updating with `cwu`. Any existing “new” state from running `cwu` previously is lost in this process. Showing a watcher without updating will still run the command, but it will not report changes.

```
# Create a basic watcher
[0x00000000]> cw 0 4 p8
[0x00000000]> cw
0x00000000
00000000
# Modify the memory and update the watcher
[0x00000000]> wx 11223344
[0x00000000]> cwu
# Watcher reports modification
# The "new" state is 11223344, and the "old" state is 00000000
[0x00000000]> cw
0x00000000 modified
11223344
# Revert the watcher
[0x00000000]> cwr
# The "new" state is 00000000 again, and there is no "old" state
# The watcher reports no change since it is no longer up-to-date
[0x00000000]> cw
0x00000000
11223344
```

Overlapping watcher areas

Watched memory areas may overlap with no ill effects, but may have unexpected results if you update some but not others.

```
# Create a watcher that watches 512 bytes starting at 0
[0x00000000]> cw 0 0x200 p8
# Create a watcher that watches 16 bytes starting at 0x100
[0x00000000]> cw 0x100 0x10 p8
# Modify memory watched by both watchers
[0x00000000]> wx 11223344 @ 0x100
# Watchers aren't updated, so they don't report a change
[0x00000000]> cw*
cw 0x00000000 512 p8
cw 0x00000100 16 p8
# Update only the watcher at 0x100
```

```

[0x00000000]> cwu 0x100
# Since only one watcher was updated, the other can't
# report the change
[0x00000000]> cw*
cw 0x00000000 512 p8
cw 0x00000100 16 p8 # differs

```

Watching for code modification

Here is an example of using a disassembly command to watch code being modified.

```

# Write an initial binary blob for the example
[0x00000000]> wx 5053595a
# Use pD since it counts by bytes
[0x00000000]> cw 0 4 pD
# Watcher prints disassembly
[0x00000000]> cw
0x00000000
           0x00000000      50           push rax
           0x00000001      53           push rbx
           0x00000002      59           pop rcx
           0x00000003      5a           pop rdx

# Modify the code
[0x00000000]> wx 585b5152
[0x00000000]> cwu
# Watcher prints different disassembly and reports a change
[0x00000000]> cw
0x00000000 modified
           0x00000000      58           pop rax
           0x00000001      5b           pop rbx
           0x00000002      51           push rcx
           0x00000003      52           push rdx

```

Comparing Bytes

For most generic reverse engineering tasks like finding the differences between two binary files, which bytes has changed, find differences in the graphs of the code analysis results, and other diffing operations you can just use `radiff2`:

```
$ radiff2 -h
```

Inside `r2`, the functionalities exposed by `radiff2` are available with the `c` command.

`c` (short for “compare”) allows you to compare arrays of bytes from different sources. The command accepts input in a number of formats and then compares it against values found at current seek position.

```
[0x00404888]> c?
```

```

Usage: c[?dfx] [argument]    # Compare
| c [string]                  Compare a plain with escaped chars string
| c* [string]                 Same as above, but printing r2 commands instead
| c1 [addr]                   Compare 8 bits from current offset
| c2 [value]                  Compare a word from a math expression
| c4 [value]                  Compare a doubleword from a math expression
| c8 [value]                  Compare a quadword from a math expression
| cat [file]                  Show contents of file (see pwd, ls)
| cc [at]                     Compares in two hexdump columns of block size
| ccc [at]                    Same as above, but only showing different lines
| ccd [at]                    Compares in two disasm columns of block size
| ccdd [at]                   Compares decompiler output (e cmd.pdc=pdg|pdd)
| cf [file]                   Compare contents of file at current seek
| cg[?] [o] [file]           Graphdiff current file and [file]
| cu[?] [addr] @at           Compare memory hexdumps of $$ and dst in unified diff
| cud [addr] @at             Unified diff disasm from $$ and given address
| cv[1248] [hexpairs] @at    Compare 1,2,4,8-byte (silent return in $?)
| cV[1248] [addr] @at        Compare 1,2,4,8-byte address contents (silent, return in $?)
| cw[?] [us?] [...]          Compare memory watchers
| cx [hexpair]                Compare hexpair string (use '.' as nibble wildcard)
| cx* [hexpair]               Compare hexpair string (output r2 commands)
| cX [addr]                   Like 'cc' but using hexdiff output
| cd [dir]                    chdir
| cl|cls|clear                Clear screen, (clear0 to goto 0, 0 only)

```

To compare memory contents at current seek position against a given string of values, use `cx`:

```

[0x08048000]> p8 4
7f 45 4c 46

[0x08048000]> cx 7f 45 90 46
Compare 3/4 equal bytes
0x00000002 (byte=03)  90 ' ' -> 4c 'L'
[0x08048000]>

```

Another subcommand of the `c` command is `cc` which stands for “compare code”.

To compare a byte sequence with a sequence in memory:

```

[0x4A13B8C0]> cc 0x39e8e089 @ 0x4A13B8C0

```

To compare contents of two functions specified by their names:

```

[0x08049A80]> cc sym.main2 @ sym.main

```

`c8` compares a quadword from the current seek (in the example below, `0x00000000`) against a math expression:

```

[0x00000000]> c8 4

```

```

Compare 1/8 equal bytes (0%)
0x00000000 (byte=01)  7f ' ' -> 04 ' '
0x00000001 (byte=02)  45 'E' -> 00 ' '
0x00000002 (byte=03)  4c 'L' -> 00 ' '

```

The number parameter can, of course, be math expressions which use flag names and anything allowed in an expression:

```
[0x00000000]> cx 7f469046
```

```

Compare 2/4 equal bytes
0x00000001 (byte=02)  45 'E' -> 46 'F'
0x00000002 (byte=03)  4c 'L' -> 90 ' '

```

You can use the compare command to find differences between a current block and a file previously dumped to a disk:

```

r2 /bin/true
[0x08049A80]> s 0
[0x08048000]> cf /bin/true
Compare 512/512 equal bytes

```

Dietline

Radare2 comes with the lean readline-like input capability through the lean library to handle the command edition and history navigation. It allows users to perform cursor movements, search the history, and implements autocompletion. Moreover, due to the radare2 portability, dietline provides the uniform experience among all supported platforms. It is used in all radare2 subshells - main prompt, SDB shell, visual prompt, and offsets prompt. It also implements the most common features and keybindings compatible with the GNU Readline.

Dietline supports two major configuration modes : Emacs-mode and Vi-mode.

It also supports the famous **Ctrl-R** reverse history search. Using **TAB** key it allows to scroll through the autocompletion options.

Autocompletion

In the every shell and radare2 command autocompletion is supported. There are multiple modes of it - files, flags, and SDB keys/namespaces. To provide the easy way to select possible completion options the scrollable popup widget is available. It can be enabled with `scr.prompt.popup`, just set it to the `true`.

Emacs (default) mode

By default dietline mode is compatible with readline Emacs-like mode key bindings. Thus active are:

Moving

- **Ctrl-a** - move to the beginning of the line
- **Ctrl-e** - move to the end of the line
- **Ctrl-b** - move one character backward
- **Ctrl-f** - move one character forward

Deleting

- **Ctrl-w** - delete the previous word
- **Ctrl-u** - delete the whole line
- **Ctrl-h** - delete a character to the left
- **Ctrl-d** - delete a character to the right
- **Alt-d** - cuts the character after the cursor

Killing and Yanking

- **Ctrl-k** - kill the text from point to the end of the line.
- **Ctrl-x** - kill backward from the cursor to the beginning of the current line.
- **Ctrl-t** - kill from point to the end of the current word, or if between words, to the end of the next word. Word boundaries are the same as forward-word.
- **Ctrl-w** - kill the word behind point, using white space as a word boundary. The killed text is saved on the kill-ring.
- **Ctrl-y** - yank the top of the kill ring into the buffer at point.
- **Ctrl-]** - rotate the kill-ring, and yank the new top. You can only do this if the prior command is yank or yank-pop.

History

- **Ctrl-r** - the reverse search in the command history

Vi mode

Radare2 also comes with in vi mode that can be enabled by toggling `scr.prompt.vi`. The various keybindings available in this mode are:

Entering command modes

- **ESC** - enter into the control mode

- i - enter into the insert mode

Moving

- j - acts like up arrow key
- k - acts like down arrow key
- a - move cursor forward and enter into insert mode
- I - move to the beginning of the line and enter into insert mode
- A - move to the end of the line and enter into insert mode
- ^ - move to the beginning of the line
- 0 - move to the beginning of the line
- \$ - move to the end of the line
- h - move one character backward
- l - move one character forward

Deleting and Yanking

- x - cuts the character
- dw - delete the current word
- diw - deletes the current word.
- db - delete the previous word
- D - delete the whole line
- dh - delete a character to the left
- dl - delete a character to the right
- d\$ - kill the text from point to the end of the line.
- d^ - kill backward from the cursor to the beginning of the current line.
- de - kill from point to the end of the current word, or if between words, to the end of the next word. Word boundaries are the same as forward-word.
- p - yank the top of the kill ring into the buffer at point.
- c - acts similar to d based commands, but goes into insert mode in the end by prefixing the commands with numbers, the command is performed multiple times.

If you are finding it hard to keep track of which mode you are in, just set `scr.prompt.mode=true` to update the color of the prompt based on the vi-mode.

Flags

Flags are conceptually similar to bookmarks. They associate a name with a given offset in a file. Flags can be grouped into ‘flag spaces’. A flag space is a namespace for flags, grouping together flags of similar characteristics or type. Examples for flag spaces: sections, registers, symbols.

To create a flag:

```
[0x4A13B8C0]> f flag_name @ offset
```

You can remove a flag by appending the `-` character to command. Most commands accept `-` as argument-prefix as an indication to delete something.

```
[0x4A13B8C0]> f-flag_name
```

To switch between or create new flagspaces use the `fs` command:

```
[0x00005310]> fs?
|Usage: fs [*] [+][flag|space|addr] # Manage flagspaces
| fs                display flagspaces
| fs*               display flagspaces as r2 commands
| fsj               display flagspaces in JSON
| fs *              select all flagspaces
| fs flagspace      select flag space or create if it doesn't exist
| fs-flag space     remove flag space
| fs-*              remove all flagspaces
| fs+foo            push previous flag space and set
| fs-               pop to the previous flag space
| fs-              remove the current flag space
| fsq               list flagspaces in quiet mode
| fsm [addr]        move flags at given address to the current flag space
| fss               display flagspaces stack
| fss*              display flagspaces stack in r2 commands
| fssj              display flagspaces stack in JSON
| fsr newname       rename selected flag space
[0x00005310]> fs
0 439 * strings
1  17 * symbols
2  54 * sections
3  20 * segments
4 115 * relocations
5 109 * imports
[0x00005310]>
```

Here there are some command examples:

```
[0x4A13B8C0]> fs symbols ; select only flags in symbols flag space
[0x4A13B8C0]> f          ; list only flags in symbols flag space
[0x4A13B8C0]> fs *       ; select all flag spaces
[0x4A13B8C0]> f myflag   ; create a new flag called 'myflag'
[0x4A13B8C0]> f-myflag   ; delete the flag called 'myflag'
```

You can rename flags with `fr`.

Local flags

Every flag name should be unique for addressing reasons. But it is quite a common need to have the flags, for example inside the functions, with simple and ubiquitous names like `loop` or `return`. For this purpose you can use so called

“local” flags, which are tied to the function where they reside. It is possible to add them using `f.` command:

```
[0x00003a04]> pd 10
|      0x00003a04      48c705c9cc21.  mov qword [0x002206d8], 0xffffffffffffffff ;
[0x2206d8:8]=0
|      0x00003a0f      c60522cc2100.  mov byte [0x00220638], 0      ; [0x220638:1]=0
|      0x00003a16      83f802      cmp eax, 2
|      .-< 0x00003a19      0f84880d0000  je 0x47a7
|      | 0x00003a1f      83f803      cmp eax, 3
|      .--< 0x00003a22      740e      je 0x3a32
|      || 0x00003a24      83e801      sub eax, 1
|      .---< 0x00003a27      0f84ed080000  je 0x431a
|      |||| 0x00003a2d      e8fef8ffff    call sym.imp.abort      ; void abort(void)
|      |||| ; CODE XREF from main (0x3a22)
|      ||\--> 0x00003a32      be07000000    mov esi, 7
[0x00003a04]> f. localflag @ 0x3a32
[0x00003a04]> f.
0x00003a32 localflag      [main + 210]
[0x00003a04]> pd 10
|      0x00003a04      48c705c9cc21.  mov qword [0x002206d8], 0xffffffffffffffff ;
[0x2206d8:8]=0
|      0x00003a0f      c60522cc2100.  mov byte [0x00220638], 0      ; [0x220638:1]=0
|      0x00003a16      83f802      cmp eax, 2
|      .-< 0x00003a19      0f84880d0000  je 0x47a7
|      | 0x00003a1f      83f803      cmp eax, 3
|      .--< 0x00003a22      740e      je 0x3a32      ; main.localflag
|      || 0x00003a24      83e801      sub eax, 1
|      .---< 0x00003a27      0f84ed080000  je 0x431a
|      |||| 0x00003a2d      e8fef8ffff    call sym.imp.abort      ; void abort(void)
|      |||| ; CODE XREF from main (0x3a22)
|      ||\--> .localflag:
|      |||| ; CODE XREF from main (0x3a22)
|      ||\--> 0x00003a32      be07000000    mov esi, 7
[0x00003a04]>
```

Flag Zones

radare2 offers flag zones, which lets you label different offsets on the scrollbar, for making it easier to navigate through large binaries. You can set a flag zone on the current seek using:

```
[0x00003a04]> fz flag-zone-name
```

Set `scr.scrollbar=1` and go to the Visual mode, to see your flag zone appear on the scrollbar on the right end of the window.

See `fz?` for more information.

Basic Commands

Most command names in radare are derived from action names. They should be easy to remember, as they are short. Actually, all commands are single letters. Subcommands or related commands are specified using the second character of the command name. For example, `/ foo` is a command to search plain string, while `/x 90 90` is used to look for hexadecimal pairs.

The general format for a valid command (as explained in the Command Format chapter) looks like this:

```
[.][times][cmd][~grep][@[@iter]addr!size][!>pipe] ; ...
```

For example,

```
> 3s +1024      ; seeks three times 1024 from the current seek
```

If a command starts with `!=`, the rest of the string is passed to the currently loaded IO plugin (a debugger, for example). Most plugins provide help messages with `!=?` or `!=help`.

```
$ r2 -d /bin/ls
```

```
> !=help        ; handled by the IO plugin
```

If a command starts with `!`, `posix_system()` is called to pass the command to your shell. Check `!?` for more options and usage examples.

```
> !ls           ; run `ls` in the shell
```

The meaning of the arguments (iter, addr, size) depends on the specific command. As a rule of thumb, most commands take a number as an argument to specify the number of bytes to work with, instead of the currently defined block size. Some commands accept math expressions or strings.

```
> px 0x17        ; show 0x17 bytes in hexs at current seek
> s base+0x33     ; seeks to flag 'base' plus 0x33
> / lib          ; search for 'lib' string.
```

The `@` sign is used to specify a temporary offset location or a seek position at which the command is executed, instead of current seek position. This is quite useful as you don't have to seek around all the time.

```
> p8 10 @ 0x4010  ; show 10 bytes at offset 0x4010
> f patata @ 0x10 ; set 'patata' flag at offset 0x10
```

Using `@@` you can execute a single command on a list of flags matching the glob. You can think of this as a foreach operation:

```
> s 0
> / lib                ; search 'lib' string
> p8 20 @@ hit0_*      ; show 20 hexpairs at each search hit
```

The `>` operation is used to redirect the output of a command into a file (overwriting it if it already exists).

```
> pr > dump.bin ; dump 'raw' bytes of current block to file named 'dump.bin'
> f > flags.txt ; dump flag list to 'flags.txt'
```

The `|` operation (pipe) is similar to what you are used to expect from it in a *NIX shell: an output of one command as input to another.

```
[0x4A13B8C0]> f | grep section | grep text
0x0805f3b0 512 section._text
0x080d24b0 512 section._text_end
```

You can pass several commands in a single line by separating them with a semicolon `;`:

```
> px ; dr
```

Using `_`, you can print the result that was obtained by the last command.

```
[0x00001060]> axt 0x00002004
main 0x1181 [DATA] lea rdi, str.argv__2d:__s
[0x00001060]> _
main 0x1181 [DATA] lea rdi, str.argv__2d:__s
```

Mapping Files

Radare's I/O subsystem allows you to map the contents of files into the same I/O space used to contain a loaded binary. New contents can be placed at random offsets.

The `o` command permits the user to open a file, this is mapped at offset 0 unless it has a known binary header and then the maps are created in virtual addresses.

Sometimes, we want to rebase a binary, or maybe we want to load or map the file in a different address.

When launching `r2`, the base address can be changed with the `-B` flag. But you must notice the difference when opening files with unknown headers, like bootloaders, so we need to map them using the `-m` flag (or specifying it as argument to the `o` command).

`radare2` is able to open files and map portions of them at random places in memory specifying attributes like permissions and name. It is the perfect basic tooling to reproduce an environment like a core file, a debug session, by also loading and mapping all the libraries the binary depends on.

Opening files (and mapping them) is done using the `o` (open) command. Let's read the help:

```
[0x00000000]> o?
|Usage: o [com- ] [file] ([offset])
```

o	list opened files
o-1	close file descriptor 1
o-!*	close all opened files
o--	close all files, analysis, binfiles, flags, same as !r2 --
o [file]	open [file] file in read-only
o+ [file]	open file in read-write mode
o [file] 0x4000 rwx	map file at 0x4000
oa[-] [A] [B] [filename]	Specify arch and bits for given file
oq	list all open files
o*	list opened files in r2 commands
o. [len]	open a malloc://[len] copying the bytes from current offset
o=	list opened files (ascii-art bars)
ob[?] [lbdos] [...]	list opened binary files backed by fd
oc [file]	open core file, like relaunching r2
of [file]	open file and map it at addr 0 as read-only
oi[- idx]	alias for o, but using index instead of fd
oj[?]	list opened files in JSON format
oL	list all IO plugins registered
om[?]	create, list, remove IO maps
on [file] 0x4000	map raw file at 0x4000 (no r_bin involved)
oo[?]	reopen current file (kill+fork in debugger)
oo+	reopen current file in read-write
ood[r] [args]	reopen in debugger mode (with args)
oo[bnm] [...]	see oo? for help
op [fd]	prioritize given fd (see also ob)
ox fd fdx	exchange the descs of fd and fdx and keep the mapping

Prepare a simple layout:

```
$ rabin2 -l /bin/ls
[Linked libraries]
libselinux.so.1
librt.so.1
libacl.so.1
libc.so.6
```

4 libraries

Map a file:

```
[0x00001190]> o /bin/zsh 0x499999
```

List mapped files:

```
[0x00000000]> o
- 6 /bin/ls @ 0x0 ; r
- 10 /lib/ld-linux.so.2 @ 0x100000000 ; r
- 14 /bin/zsh @ 0x499999 ; r
```

Print hexadecimal values from /bin/zsh:

```
[0x00000000]> px @ 0x499999
```

Unmap files using the o- command. Pass the required file descriptor to it as an argument:

```
[0x00000000]> o-14
```

You can also view the ascii table showing the list of the opened files:

```
[0x00000000]> ob=
```

Print Modes

One of the key features of radare2 is displaying information in many formats. The goal is to offer a selection of display choices to interpret binary data in the best possible way.

Binary data can be represented as integers, shorts, longs, floats, timestamps, hexpair strings, or more complex formats like C structures, disassembly listings, decompilation listing, be a result of an external processing...

Below is a list of available print modes listed by p?:

```
[0x00005310]> p?
|Usage: p[=68abcdDfiImrstuxz] [arg|len] [@addr]
| p[b|B|xb] [len] ([S])   bindump N bits skipping S bytes
| p[iI][df] [len]         print N ops/bytes (f=func) (see pi? and pdi)
| p[kK] [len]             print key in randomart (K is for mosaic)
| p-?[?][jh] [mode]      bar|json|histogram blocks (mode: e?search.in)
| p2 [len]                8x8 2bpp-tiles
| p3 [file]               print stereogram (3D)
| p6[de] [len]            base64 decode/encode
| p8[?][j] [len]          8bit hexpair list of bytes
| p=[?][bep] [N] [L] [b] show entropy/printable chars/chars bars
| pa[edD] [arg]           pa:assemble pa[dD]:disasm or pae: esil from hex
| pA[n_ops]              show n_ops address and type
| pb[?] [n]               bitstream of N bits
| pB[?] [n]               bitstream of N bytes
| pc[?][p] [len]          output C (or python) format
| pC[aAcDxw] [rows]       print disassembly in columns (see hex.cols and pdi)
| pd[?] [sz] [a] [b]      disassemble N opcodes (pd) or N bytes (pD)
| pf[?][.nam] [fmt]        print formatted data (pf.name, pf.name $<expr>)
| pF[?][apx]              print asn1, pkcs7 or x509
| pg[?][x y w h] [cmd]     create new visual gadget or print it (see pg? for details)
| ph[?][=|hash] ([len])   calculate hash for a block
| pj[?] [len]              print as indented JSON
| pm[?] [magic]            print libmagic data (see pm? and /m?)
```

po[?] hex	print operation applied to block (see po?)
pp[?][sz] [len]	print patterns, see pp? for more help
pq[?][is] [len]	print QR code with the first Nbytes
pr[?][glx] [len]	print N raw bytes (in lines or hexblocks, 'g'unzip)
ps[?][pwz] [len]	print pascal/wide/zero-terminated strings
pt[?][dn] [len]	print different timestamps
pu[?][w] [len]	print N url encoded bytes (w=wide)
pv[?][jh] [mode]	show variable/pointer/value in memory
pwd	display current working directory
px[?][owq] [len]	hexdump of N bytes (o=octal, w=32bit, q=64bit)
pz[?] [len]	print zoom view (see pz? for help)

[0x00005310]>

Tip: when using json output, you can append the ~{} to the command to get a pretty-printed version of the output:

```
[0x00000000]> oj
[{"raised":false,"fd":563280,"uri":"malloc://512","from":0,"writable":true,"size":512,"overlaps":false}]
[0x00000000]> oj~{}
[
  {
    "raised": false,
    "fd": 563280,
    "uri": "malloc://512",
    "from": 0,
    "writable": true,
    "size": 512,
    "overlaps": false
  }
]
```

For more on the magical powers of ~ see the help in ?@?, and the Command Format chapter earlier in the book.

Hexadecimal View

px gives a user-friendly output showing 16 pairs of numbers per row with offsets and raw representations:

hexprint

Show Hexadecimal Words Dump (32 bits) wordprint

8 bits Hexpair List of Bytes

```
[0x00404888]> p8 16
31ed4989d15e4889e24883e4f0505449
```

Show Hexadecimal Quad-words Dump (64 bits) pxq

Date/Time Formats

Currently supported timestamp output modes are:

```
[0x00404888]> pt?  
|Usage: pt [dn] print timestamps  
| pt. print current time  
| pt print UNIX time (32 bit `cfg.bigendian`) Since January 1, 1970  
| ptd print DOS time (32 bit `cfg.bigendian`) Since January 1, 1980  
| pth print HFS time (32 bit `cfg.bigendian`) Since January 1, 1904  
| ptn print NTFS time (64 bit `cfg.bigendian`) Since January 1, 1601
```

For example, you can ‘view’ the current buffer as timestamps in the ntfs time:

```
[0x08048000]> e cfg.bigendian = false  
[0x08048000]> pt 4  
29:04:32948 23:12:36 +0000  
[0x08048000]> e cfg.bigendian = true  
[0x08048000]> pt 4  
20:05:13001 09:29:21 +0000
```

As you can see, the endianness affects the result. Once you have printed a timestamp, you can grep the output, for example, by year:

```
[0x08048000]> pt ~1974 | wc -l  
15  
[0x08048000]> pt ~2022  
27:04:2022 16:15:43 +0000
```

The default date format can be configured using the `cfg.datefmt` variable. Formatting rules for it follow the well known `strftime(3)` format. Check the `manpage` for more details, but these are the most important:

```
%a The abbreviated name of the day of the week according to the current locale.  
%A The full name of the day of the week according to the current locale.  
%d The day of the month as a decimal number (range 01 to 31).  
%D Equivalent to %m/%d/%y. (Yecch-for Americans only).  
%H The hour as a decimal number using a 24-hour clock (range 00 to 23).  
%I The hour as a decimal number using a 12-hour clock (range 01 to 12).  
%m The month as a decimal number (range 01 to 12).  
%M The minute as a decimal number (range 00 to 59).  
%p Either "AM" or "PM" according to the given time value.  
%s The number of seconds since the Epoch, 1970-01-01 00:00:00 +0000 (UTC). (TZ)  
%S The second as a decimal number (range 00 to 60). (The range is up to 60 to allow for o  
%T The time in 24-hour notation (%H:%M:%S). (SU)  
%y The year as a decimal number without a century (range 00 to 99).  
%Y The year as a decimal number including the century.
```

%z The +hhmm or -hhmm numeric timezone (that is, the hour and minute offset from UTC). (SU)
 %Z The timezone name or abbreviation.

Basic Types

There are print modes available for all basic types. If you are interested in a more complex structure, type `pf??` for format characters and `pf???` for examples:

```
[0x00499999]> pf??
|pf: pf[.k[.f[=v]]|[v]]|[n]|[0|cnt][fmt] [a0 a1 ...]
| Format:
| b      byte (unsigned)
| B      resolve enum bitfield (see t?)
| c      char (signed byte)
| C      byte in decimal
| d      0xHEX value (4 bytes) (see 'i' and 'x')
| D      disassemble one opcode
| e      temporally swap endian
| E      resolve enum name (see t?)
| f      float value (4 bytes)
| F      double value (8 bytes)
| i      signed integer value (4 bytes) (see 'd' and 'x')
| n      next char specifies size of signed value (1, 2, 4 or 8 byte(s))
| N      next char specifies size of unsigned value (1, 2, 4 or 8 byte(s))
| o      octal value (4 byte)
| p      pointer reference (2, 4 or 8 bytes)
| q      quadword (8 bytes)
| r      CPU register `pf r (eax)plop`
| s      32bit pointer to string (4 bytes)
| S      64bit pointer to string (8 bytes)
| t      UNIX timestamp (4 bytes)
| T      show Ten first bytes of buffer
| u      uleb128 (variable length)
| w      word (2 bytes unsigned short in hex)
| x      0xHEX value and flag (fd @ addr) (see 'd' and 'i')
| X      show formatted hexpairs
| z      null terminated string
| Z      null terminated wide string
| ?      data structure `pf ? (struct_name)example_name`
| *      next char is pointer (honors asm.bits)
| +      toggle show flags for each offset
| :      skip 4 bytes
| .      skip 1 byte
| ;      rewind 4 bytes
| ,      rewind 1 byte
```

Use triple-question-mark `pf???` to get some examples using print format strings.


```
[0x00499999]> pf???
|pf: pf[.k[.f[=v]]|[v]]|[n]|[0|cnt][fmt] [a0 a1 ...]
| Examples:
| pf 3xi foo bar          3-array of struct, each with named fields: 'f'
| pf B (BitFldType)arg_name` bitfield type
| pf E (EnumType)arg_name`  enum type
| pf.obj xxdz prev next size name Define the obj format as xxdz
| pf obj=xxdz prev next size name Same as above
| pf *z*i*w nb name blob      Print the pointers with given labels
| pf iwq foo bar troll        Print the iwq format with foo, bar, troll as
| pf 0iwq foo bar troll       Same as above, but considered as a union (all
| pf.plop ? (troll)mystruct   Use structure troll previously defined
| pfj.plop @ 0x14             Apply format object at the given offset
| pf 10xiz pointer length string Print a size 10 array of the xiz struct with
| pf 5sqw string quad word    Print an array with sqw struct along with its
| pf {integer}? (bifc)        Print integer times the following format (bif
| pf [4]w[7]i                Print an array of 4 words and then an array of
| pf ic...?i foo bar "(pf xw yo foo)troll" yo Print nested anonymous structures
| pf ;..x                    Print value located 6 bytes from current offset
| pf [10]z[3]i[10]Zb         Print an fixed size str, widechar, and var
| pfj +F @ 0x14              Print the content at given offset with flag
| pf n2                      print signed short (2 bytes) value. Use N in
| pf [2]? (plop)structname @ 0 Prints an array of structs
| pf eqew bigWord beef       Swap endianness and print with given labels
| pf.foo rr (eax)reg1 (eip)reg2 Create object referencing to register values
| pf tt troll plop           print time stamps with labels troll and plop
```

Some examples are below:

```
[0x4A13B8C0]> pf i
0x00404888 = 837634441

[0x4A13B8C0]> pf
0x00404888 = 837634432.000000
```

High-level Languages Views

Valid print code formats for human-readable languages are:

- pc C
- pc* print 'wx' r2 commands
- pch C half-words (2 byte)
- pcw C words (4 byte)
- pcd C dwords (8 byte)
- pci C array of bytes with instructions
- pca GAS .byte blob
- pcA .bytes with instructions in comments
- pcs string

- pcS shellscript that reconstructs the bin
- pcj json
- pcJ javascript
- pco Objective-C
- pcp python
- pck kotlin
- pcr rust
- pcv JaVa
- pcV V (vlang.io)
- pcy yara
- pcz Swift

If we need to create a .c file containing a binary blob, use the `pc` command, that creates this output. The default size is like in many other commands: the block size, which can be changed with the `b` command.

We can also just temporarily override this block size by expressing it as an argument.

```
[0xB7F8E810]> pc 32
#define _BUFFER_SIZE 32
unsigned char buffer[_BUFFER_SIZE] = {
0x89, 0xe0, 0xe8, 0x49, 0x02, 0x00, 0x00, 0x89, 0xc7, 0xe8, 0xe2, 0xff, 0xff, 0xff, 0x81, 0x
```

That cstring can be used in many programming languages, not just C.

```
[0x7fcd6a891630]> pcs
"\x48\x89\xe7\xe8\x68\x39\x00\x00\x49\x89\xc4\x8b\x05\xef\x16\x22\x00\x5a\x48\x8d\x24\xc4\x2
```

Strings

Strings are probably one of the most important entry points when starting to reverse engineer a program because they usually reference information about functions' actions (asserts, debug or info messages...). Therefore, radare supports various string formats:

```
[0x00000000]> ps?
|Usage: ps[bijqpsuwWxz+] [N] Print String
| ps      print string
| ps+[j]  print libc++ std::string (same-endian, ascii, zero-terminated)
| psb     print strings in current block
| psi     print string inside curseek
| psj     print string in JSON format
| psp[j]  print pascal string
| psq     alias for pqs
| pss     print string in screen (wrap width)
| psu[zj] print utf16 unicode (json)
| psw[j]  print 16bit wide string
| psW[j]  print 32bit wide string
```

```
| psx      show string with escaped chars
| psz[j]   print zero-terminated string
```

Most strings are zero-terminated. Below there is an example using the debugger to continue the execution of a program until it executes the ‘open’ syscall. When we recover the control over the process, we get the arguments passed to the syscall, pointed by %ebx. In the case of the ‘open’ call, it is a zero terminated string which we can inspect using psz.

```
[0x4A13B8C0]> dcs open
0x4a14fc24 syscall(5) open ( 0x4a151c91 0x00000000 0x00000000 ) = 0xffffffffda
[0x4A13B8C0]> dr
    eax  0xffffffffda    esi  0xffffffff    eip    0x4a14fc24
    ebx  0x4a151c91     edi  0x4a151be1    oeax   0x00000005
    ecx  0x00000000     esp  0xbfbdb1c    eflags 0x200246
    edx  0x00000000     ebp  0xbfbdbbb0    cPaZstIdor0 (PZI)
[0x4A13B8C0]>
[0x4A13B8C0]> psz @ 0x4a151c91
/etc/ld.so.cache
```

Print Memory Contents

It is also possible to print various packed data types using the pf command:

```
[0xB7F08810]> pf xxS @ rsp
0x7fff0d29da30 = 0x00000001
0x7fff0d29da34 = 0x00000000
0x7fff0d29da38 = 0x7fff0d29da38 -> 0x0d29f7ee /bin/ls
```

This can be used to look at the arguments passed to a function. To achieve this, simply pass a ‘format memory string’ as an argument to pf, and temporally change the current seek position/offset using @. It is also possible to define arrays of structures with pf. To do this, prefix the format string with a numeric value. You can also define a name for each field of the structure by appending them as a space-separated arguments list.

```
[0x4A13B8C0]> pf 2*xw pointer type @ esp
0x00404888 [0] {
    pointer :
    (*0xfffffffff8949ed31)    type : 0x00404888 = 0x8949ed31
    0x00404890 = 0x48e2
}
0x00404892 [1] {
    (*0x50f0e483)    pointer : 0x00404892 = 0x50f0e483
    type : 0x0040489a = 0x2440
}
```

A practical example for using pf on a binary of a GStreamer plugin:

```
$ radare2 /usr/lib/gstreamer-1.0/libgstflv.so
[0x00006020]> aa; pdf @ sym.gst_plugin_flv_get_desc
[x] Analyze all flags starting with sym. and entry0 (aa)
sym.gst_plugin_flv_get_desc ();
[...]
```

0x00013830	488d0549db0000	lea rax, section..data.rel.ro ; 0x21380
0x00013837	c3	ret

```
[0x00006020]> s section..data.rel.ro
[0x00021380]> pf ii*z*zp*z*z*z*z*z*z major minor name desc init version license source package
major : 0x00021380 = 1
minor : 0x00021384 = 18
name : (*0x19cf2)0x00021388 = "flv"
desc : (*0x1b358)0x00021390 = "FLV muxing and demuxing plugin"
init : 0x00021398 = (qword)0x00000000000013460
version : (*0x19cae)0x000213a0 = "1.18.2"
license : (*0x19ce1)0x000213a8 = "LGPL"
source : (*0x19cd0)0x000213b0 = "gst-plugins-good"
package : (*0x1b378)0x000213b8 = "GStreamer Good Plugins (Arch Linux)"
origin : (*0x19cb5)0x000213c0 = "https://www.archlinux.org/"
release_datetime : (*0x19cf6)0x000213c8 = "2020-12-06"
```

Disassembly

The `pd` command is used to disassemble code. It accepts a numeric value to specify how many instructions should be disassembled. The `pD` command is similar but instead of a number of instructions, it decompiles a given number of bytes.

- `d` : disassembly N opcodes count of opcodes
- `D` : asm.arch disassembler bsize bytes

```
[0x00404888]> pd 1
;-- entry0:
0x00404888    31ed          xor ebp, ebp
```

Selecting Target Architecture

The architecture flavor for the disassembler is defined by the `asm.arch` eval variable. You can use `e asm.arch=??` to list all available architectures.

```
[0x00005310]> e asm.arch=??
_dAe _8_16      6502      LGPL3    6502/NES/C64/Tamagotchi/T-1000 CPU
_dAe _8          8051      PD         8051 Intel CPU
_dA_ _16_32     arc       GPL3      Argonaut RISC Core
a___ _16_32_64  arm.as    LGPL3    as ARM Assembler (use ARM_AS environment)
adAe _16_32_64  arm       BSD       Capstone ARM disassembler
_dA_ _16_32_64  arm.gnu   GPL3     Acorn RISC Machine CPU
_d__ _16_32     arm.winedbg LGPL2    WineDBG's ARM disassembler
```

adAe	_8_16	avr	GPL	AVR Atmel
adAe	_16_32_64	bf	LGPL3	Brainfuck
_dA	_32	chip8	LGPL3	Chip8 disassembler
_dA	_16	cr16	LGPL3	cr16 disassembly plugin
_dA	_32	cris	GPL3	Axis Communications 32-bit embedded processor
adA	_32_64	dalvik	LGPL3	AndroidVM Dalvik
ad__	_16	dcpu16	PD	Mojang's DCPU-16
_dA	_32_64	ebc	LGPL3	EFI Bytecode
adAe	_16	gb	LGPL3	GameBoy(TM) (z80-like)
_dAe	_16	h8300	LGPL3	H8/300 disassembly plugin
_dAe	_32	hexagon	LGPL3	Qualcomm Hexagon (QDSP6) V6
_d__	_32	hppa	GPL3	HP PA-RISC
_dAe	_0	i4004	LGPL3	Intel 4004 microprocessor
_dA	_8	i8080	BSD	Intel 8080 CPU
adA	_32	java	Apache	Java bytecode
_d__	_32	lanai	GPL3	LANAI
...				

Configuring the Disassembler

There are multiple options which can be used to configure the output of the disassembler. All these options are described in `e? asm`.

```
[0x00005310]> e? asm.
asm.anal: Analyze code and refs while disassembling (see anal.strings)
asm.arch: Set the arch to be used by asm
asm.assembler: Set the plugin name to use when assembling
asm.bbline: Show empty line after every basic block
asm.bits: Word size in bits at assembler
asm.bytes: Display the bytes of each instruction
asm.bytespace: Separate hexadecimal bytes with a whitespace
asm.calls: Show callee function related info as comments in disasm
asm.capitalize: Use camelcase at disassembly
asm.cmt.col: Column to align comments
asm.cmt.flgreffs: Show comment flags associated to branch reference
asm.cmt.fold: Fold comments, toggle with Vz
...
```

Currently there are 136 `asm.` configuration variables so we do not list them all.

Disassembly Syntax

The `asm.syntax` variable is used to change the flavor of the assembly syntax used by a disassembler engine. To switch between Intel and AT&T representations:

```
e asm.syntax = intel
e asm.syntax = att
```

You can also check `asm.pseudo`, which is an experimental pseudocode view, and `asm.esil` which outputs ESIL ('Evaluable Strings Intermediate Language'). ESIL's goal is to have a human-readable representation of every opcode semantics. Such representations can be evaluated (interpreted) to emulate effects of individual instructions.

SDB

SDB stands for String DataBase. It's a simple key-value database that only operates with strings created by `pancake`. It is used in many parts of `r2` to have a disk and in-memory database which is small and fast to manage using it as a hashtable on steroids.

SDB is a simple string key/value database based on `djb's cdb` disk storage and supports JSON and arrays introspection.

There's also the `sdbtypes`: a `vala` library that implements several data structures on top of an `sdb` or a `memcache` instance.

SDB supports:

- namespaces (multiple `sdb` paths)
- atomic database sync (never corrupted)
- bindings for `vala`, `luvit`, `newlisp` and `nodejs`
- commandline frontend for `sdb` databases
- `memcache` client and server with `sdb` backend
- arrays support (syntax sugar)
- `json` parser/getter

Usage example

Let's create a database!

```
$ sdb d hello=world
$ sdb d hello
world
```

Using arrays:

```
$ sdb - '[]list=1,2' '[0]list' '[0]list=foo' '[]list' '[+1]list=bar'
1
foo
2
foo
bar
2
```

Let's play with `json`:

```
$ sdb d g='{"foo":1,"bar":{"cow":3}}'
$ sdb d g?bar.cow
3
$ sdb - user='{"id":123}' user?id=99 user?id
99
```

Using the command line without any disk database:

```
$ sdb - foo=bar foo a=3 +a -a
bar
4
3
```

```
$ sdb -
foo=bar
foo
bar
a=3
+a
4
-a
3
```

Remove the database

```
$ rm -f d
```

So what ?

So, you can now do this inside your radare2 sessions!

Let's take a simple binary, and check what is already *sdbized*.

```
$ cat test.c
int main(){
    puts("Hello world\n");
}
$ gcc test.c -o test
$ r2 -A ./test
[0x08048320]> k **
bin
anal
syscall
debug
[0x08048320]> k bin/**
fd.6
[0x08048320]> k bin/fd.6/*
archs=0:0:x86:32
```

The file corresponding to the sixth file descriptor is a x86_32 binary.

```
[0x08048320]> k anal/meta/*
meta.s.0x80484d0=12,SGVsbG8gd29ybGQ=
[...]
[0x08048320]> ?b64- SGVsbG8gd29ybGQ=
Hello world
```

Strings are stored encoded in base64.

More Examples

List namespaces

```
k **
```

List sub-namespaces

```
k anal/**
```

List keys

```
k *
```

```
k anal/*
```

Set a key

```
k foo=bar
```

Get the value of a key

```
k foo
```

List all syscalls

```
k syscall/*~^0x
```

List all comments

```
k anal/meta/*~.C.
```

Show a comment at given offset:

```
k %anal/meta/[1]meta.C.0x100005000
```

Sections

The concept of sections is tied to the information extracted from the binary. We can display this information by using the `i` command.

Displaying information about sections:


```
[0x00005310]> iS
```

```
[Sections]
```

```
00 0x00000000      0 0x00000000      0 ----
01 0x00000238     28 0x00000238     28 -r-- .interp
02 0x00000254     32 0x00000254     32 -r-- .note.ABI_tag
03 0x00000278    176 0x00000278    176 -r-- .gnu.hash
04 0x00000328   3000 0x00000328   3000 -r-- .dynsym
05 0x00000ee0   1412 0x00000ee0   1412 -r-- .dynstr
06 0x00001464    250 0x00001464    250 -r-- .gnu.version
07 0x00001560    112 0x00001560    112 -r-- .gnu.version_r
08 0x000015d0   4944 0x000015d0   4944 -r-- .rela.dyn
09 0x00002920   2448 0x00002920   2448 -r-- .rela.plt
10 0x000032b0     23 0x000032b0     23 -r-x .init
...
```

As you may know, binaries have sections and maps. The sections define the contents of a portion of the file that can be mapped in memory (or not). What is mapped is defined by the segments.

Before the IO refactoring done by condret, the `S` command was used to manage what we now call maps. Currently the `S` command is deprecated because `iS` and `om` should be enough.

Firmware images, bootloaders and binary files usually place various sections of a binary at different addresses in memory. To represent this behavior, radare offers the `iS`. Use `iS?` to get the help message. To list all created sections use `iS` (or `iSj` to get the json format). The `iS=` will show the region bars in ascii-art.

You can create a new mapping using the `om` subcommand as follows:

```
om fd vaddr [size] [paddr] [rwx] [name]
```

For Example:

```
[0x0040100]> om 4 0x00000100 0x00400000 0x0001ae08 rwx test
```

You can also use `om` command to view information about mapped sections:

```
[0x00401000]> om
6 fd: 4 +0x0001ae08 0x00000100 - 0x004000ff rwx test
5 fd: 3 +0x00000000 0x00000000 - 0x0000055f r-- fmap.LOAD0
4 fd: 3 +0x00001000 0x00001000 - 0x000011e4 r-x fmap.LOAD1
3 fd: 3 +0x00002000 0x00002000 - 0x0000211f r-- fmap.LOAD2
2 fd: 3 +0x00002de8 0x00003de8 - 0x0000402f r-- fmap.LOAD3
1 fd: 4 +0x00000000 0x00004030 - 0x00004037 rw- mmap.LOAD3
```

Use `om?` to get all the possible subcommands. To list all the defined maps use `om` (or `omj` to get the json format or `om*` to get the r2 commands format). To get the ascii art view use `om=`.

It is also possible to delete the mapped section using the `om-mapid` command.

For Example:

```
[0x00401000]> om-6
```

Seeking

To move around the file we are inspecting we will need to change the offset at which we are using the `s` command.

The argument is a math expression that can contain flag names, parenthesis, addition, subtraction, multiplication of immediates of contents of memory using brackets.

Some example commands:

```
[0x00000000]> s 0x10
[0x00000010]> s+4
[0x00000014]> s-
[0x00000010]> s+
[0x00000014]>
```

Observe how the prompt offset changes. The first line moves the current offset to the address 0x10.

The second does a relative seek 4 bytes forward.

And finally, the last 2 commands are undoing, and redoing the last seek operations.

Instead of using just numbers, we can use complex expressions, or basic arithmetic operations to represent the address to seek.

To do this, check the `??` Help message which describes the internal variables that can be used in the expressions. For example, this is the same as doing `s+4`.

```
[0x00000000]> s $$+4
```

From the debugger (or when emulating) we can also use the register names as references. They are loaded as flags with the `.dr*` command, which happens under the hood.

```
[0x00000000]> s rsp+0x40
```

Here's the full help of the `s` command. We will explain in more detail below.

```
[0x00000000]> s?
```

```
Usage: s      # Help for the seek commands. See ?? to see all variables
| s          Print current address
| s.hexoff    Seek honoring a base from core->offset
| s:pad       Print current address with N padded zeros (defaults to 8)
| s addr      Seek to address
| s-          Undo seek
| s-*         Reset undo seek history
```

s- n	Seek n bytes backward
s--[n]	Seek blocksize bytes backward (/=n)
s+	Redo seek
s+ n	Seek n bytes forward
s++[n]	Seek blocksize bytes forward (/=n)
s[j*!=]	List undo seek history (JSON, =list, *r2, !=names, s==)
s/ DATA	Search for next occurrence of 'DATA'
s/x 9091	Search for next occurrence of \x90\x91
sa [[+--]a] [asz]	Seek asz (or bsize) aligned to addr
sb	Seek aligned to bb start
sC[?] string	Seek to comment matching given string
sf	Seek to next function (f->addr+f->size)
sf function	Seek to address of specified function
sf.	Seek to the beginning of current function
sg/sG	Seek begin (sg) or end (sG) of section or file
sl[?] [+--]line	Seek to line
sn/sp ([nkey])	Seek to next/prev location, as specified by scr.nkey
so [N]	Seek to N next opcode(s)
sr pc	Seek to register
ss	Seek silently (without adding an entry to the seek history)

```
> 3s++      ; 3 times block-seeking
> s 10+0x80  ; seek at 0x80+10
```

If you want to inspect the result of a math expression, you can evaluate it using the `? command`. Simply pass the expression as an argument. The result can be displayed in hexadecimal, decimal, octal or binary formats.

```
> ? 0x100+200
0x1C8 ; 456d ; 710o ; 1100 1000
```

There are also subcommands of `? command` that display the output in one specific format (base 10, base 16 ,...). See `?v` and `?vi`.

In the visual mode, you can press `u` (undo) or `U` (redo) inside the seek history to return back to previous or forward to the next location.

Open file

As a test file, let's use a simple `hello_world.c` compiled in Linux ELF format. After we compile it let's open it with `radare2`:

```
$ r2 hello_world
```

Now we have the command prompt:

```
[0x00400410]>
```

And it is time to go deeper.

Seeking at any position

All seeking commands that take an address as a command parameter can use any numeral base such as hex, octal, binary or decimal.

Seek to an address 0x0. An alternative command is simply 0x0

```
[0x00400410]> s 0x0
[0x00000000]>
```

Print current address:

```
[0x00000000]> s
0x0
[0x00000000]>
```

There is an alternate way to print current position: ?v \$\$.

Seek N positions forward, space is optional:

```
[0x00000000]> s+ 128
[0x00000080]>
```

Undo last two seeks to return to the initial address:

```
[0x00000080]> s-
[0x00000000]> s-
[0x00400410]>
```

We are back at *0x00400410*.

There's also a command to show the seek history:

```
[0x00400410]> s*
f undo_3 @ 0x400410
f undo_2 @ 0x40041a
f undo_1 @ 0x400410
f undo_0 @ 0x400411
# Current undo/redo position.
f redo_0 @ 0x4005b4
```

Working with data types

Radare2 can also work with data types. You can use standard C data types or define your own using C. Currently, there is a support for structs, unions, function signatures, and enums.

```
[0x00000000]> t?
```

Usage: t # cparse types commands

t	List all loaded types
tj	List all loaded types as json
t <type>	Show type in 'pf' syntax
t*	List types info in r2 commands

t- <name>	Delete types by its name
t-*	Remove all types
tail [filename]	Output the last part of files
tc [type.name]	List all/given types in C output format
te[?]	List all loaded enums
td[?] <string>	Load types from string
tf	List all loaded functions signatures
tk <sdb-query>	Perform sdb query
tl[?]	Show/Link type to an address
tn[?] [-][addr]	manage noreturn function attributes and marks
to -	Open cfg.editor to load types
to <path>	Load types from C header file
toe [type.name]	Open cfg.editor to edit types
tos <path>	Load types from parsed Sdb database
tp <type> [addr varname]	cast data at <address> to <type> and print it (XXX: type can co
tpv <type> @ [value]	Show offset formatted for given type
tpx <type> <hexpairs>	Show value for type with specified byte sequence (XXX: type can
ts[?]	Print loaded struct types
tu[?]	Print loaded union types
tx[f?]	Type xrefs
tt[?]	List all loaded typedefs

Defining new types

There are three different methods to define new types:

1. Defining a new type from r2 shell immediately, to do this you will use **td** command, and put the whole line between double quotes. For example:

```
"td struct person {int age; char *name; char *address;};"
```

2. You can also use **to -** to open a text editor and write your own types in there. This is preferable when you got too many types to define.
3. Radare2 also supports loading header files using the command **to** followed by a path to the header file you want to load.

You can View loaded types in r2 using **ts** for structures, **tu** for unions, **tf** for function signatures, **te** for enums.

You can also cast pointers to data types and view data in there accordingly with **tp**. EX:

```
[0x00400511]> tp person = 0x7fff170a46b0
    age : 0x7fff170a46b0 = 20
    name : (*0x4005b0) 0x7fff170a46b4 = My name
    address : (*0x4005b8) 0x7fff170a46bc = My age
[0x00400511]>
```

Writing Data

Radare can manipulate a loaded binary file in many ways. You can resize the file, move and copy/paste bytes, insert new bytes (shifting data to the end of the block or file), or simply overwrite bytes. New data may be given as a wide-string, assembler instructions, or the data may be read in from another file.

Resize the file using the `r` command. It accepts a numeric argument. A positive value sets a new size for the file. A negative one will truncate the file to the current seek position minus N bytes.

```
r 1024      ; resize the file to 1024 bytes
r -10 @ 33  ; strip 10 bytes at offset 33
```

Write bytes using the `w` command. It accepts multiple input formats like inline assembly, endian-friendly dwords, files, hexpair files, wide strings:

```
[0x00404888]> w?
Usage: w[x] [str] [<file] [<<EOF] [@addr]
| w[1248] [+ -] [n]      increment/decrement byte,word..
| w foobar              write string 'foobar'
| w0 [len]              write 'len' bytes with value 0x00
| w6[de] base64/hex     write base64 [d]ecoded or [e]ncoded string
| wa[?] push ebp        write opcode, separated by ';' (use '"' around the command)
| waf f.asm             assemble file and write bytes
| waF f.asm             assemble file and write bytes and show 'wx' op with hexpair bytes of
| wao[?] op             modify opcode (change conditional of jump. nop, etc)
| wA[?] r 0             alter/modify opcode at current seek (see wA?)
| wb 010203             fill current block with cyclic hexpairs
| wB[-]0xVALUE          set or unset bits with given value
| wc                   list all write changes
| wc[?][jir+*?]        write cache undo/commit/reset/list (io.cache)
| wd [off] [n]          duplicate N bytes from offset at current seek (memcpy) (see y?)
| we[?] [nNsxX] [arg]   extend write operations (insert instead of replace)
| wf[fs] -|file         write contents of file at current offset
| wh r2                 whereis/which shell command
| wm f0ff               set binary mask hexpair to be used as cyclic write mask
| wo[?] hex             write in block with operation. 'wo?' fmi
| wp[?] -|file          apply radare patch file. See wp? fmi
| wr 10                 write 10 random bytes
| ws pstring            write 1 byte for length and then the string
| wt[f][?] file [sz]    write to file (from current seek, blocksize or sz bytes)
| wts host:port [sz]    send data to remote host:port via tcp://
| ww foobar             write wide string 'f\x00o\x00o\x00b\x00a\x00r\x00'
| wx[?][fs] 9090        write two intel nops (from wxfile or wxseek)
| wv[?] eip+34          write 32-64 bit value honoring cfg.bigendian
| wz string             write zero terminated string (like w + \x00)
```

Some examples:

```
[0x00000000]> wx 123456 @ 0x8048300
[0x00000000]> wv 0x8048123 @ 0x8049100
[0x00000000]> wa jmp 0x8048320
```

Write Over

The **wo** command (write over) has many subcommands, each combines the existing data with the new data using an operator. The command is applied to the current block. Supported operators include XOR, ADD, SUB...

```
[0x4A13B8C0]> wo?
|Usage: wo[asmdxoArl24] [hexpairs] @ addr[:bsize]
|Example:
|  wox 0x90    ; xor cur block with 0x90
|  wox 90      ; xor cur block with 0x90
|  wox 0x0203  ; xor cur block with 0203
|  woa 02 03   ; add [0203][0203][...] to curblk
|  woe 02 03   ; create sequence from 2 to 255 with step 3
|Supported operations:
|  wow == write looped value (alias for 'wb')
|  woa += addition
|  wos -= subtraction
|  wom *= multiply
|  wod /= divide
|  wox ^= xor
|  woo |= or
|  woA &= and
|  woR random bytes (alias for 'wr $b')
|  wor >>= shift right
|  wol <<= shift left
|  wo2 2= 2 byte endian swap
|  wo4 4= 4 byte endian swap
```

It is possible to implement cipher-algorithms using radare core primitives and **wo**. A sample session performing **xor(90) + add(01, 02)**:

```
[0x7fcd6a891630]> px
- offset -      0 1  2 3  4 5  6 7  8 9  A B  C D  E F
0x7fcd6a891630  4889 e7e8 6839 0000 4989 c48b 05ef 1622
0x7fcd6a891640  005a 488d 24c4 29c2 5248 89d6 4989 e548
0x7fcd6a891650  83e4 f048 8b3d 061a 2200 498d 4cd5 1049
0x7fcd6a891660  8d55 0831 ede8 06e2 0000 488d 15cf e600
[0x7fcd6a891630]> wox 90
[0x7fcd6a891630]> px
- offset -      0 1  2 3  4 5  6 7  8 9  A B  C D  E F
0x7fcd6a891630  d819 7778 d919 541b 90ca d81d c2d8 1946
```

```

0x7fcd6a891640  1374 60d8 b290 d91d 1dc5 98a1 9090 d81d
0x7fcd6a891650  90dc 197c 9f8f 1490 d81d 95d9 9f8f 1490
0x7fcd6a891660  13d7 9491 9f8f 1490 13ff 9491 9f8f 1490
[0x7fcd6a891630]> woa 01 02
[0x7fcd6a891630]> px
- offset -      0 1  2 3  4 5  6 7  8 9  A B  C D  E F
0x7fcd6a891630  d91b 787a 91cc d91f 1476 61da 1ec7 99a3
0x7fcd6a891640  91de 1a7e d91f 96db 14d9 9593 1401 9593
0x7fcd6a891650  c4da 1a6d e89a d959 9192 9159 1cb1 d959
0x7fcd6a891660  9192 79cb 81da 1652 81da 1456 a252 7c77

```

Yank/Paste

Radare2 has an internal clipboard to save and write portions of memory loaded from the current io layer.

This clipboard can be manipulated with the y command.

The two basic operations are

- copy (yank)
- paste

The yank operation will read N bytes (specified by the argument) into the clipboard. We can later use the yy command to paste what we read before into a file.

You can yank/paste bytes in visual mode selecting them with the cursor mode (Vc) and then using the y and Y key bindings which are aliases for y and yy commands of the command-line interface.

```

[0x00000000]> y?
Usage: y[ptxy] [len] [[@]addr]  # See wd? for memcpy, same as 'yf'.
| y!                               open cfg.editor to edit the clipboard
| y 16 0x200                       copy 16 bytes into clipboard from 0x200
| y 16 @ 0x200                     copy 16 bytes into clipboard from 0x200
| y 16                             copy 16 bytes into clipboard
| y                                show yank buffer information (srcoff len bytes)
| y*                               print in r2 commands what's been yanked
| yf 64 0x200                     copy file 64 bytes from 0x200 from file
| yfa file copy                   copy all bytes from file (opens w/ io)
| yfx 10203040                   yank from hexpairs (same as ywx)
| yj                               print in JSON commands what's been yanked
| yp                               print contents of clipboard
| yq                               print contents of clipboard in hexpairs
| ys                               print contents of clipboard as string
| yt 64 0x200                     copy 64 bytes from current seek to 0x200
| ytf file                       dump the clipboard to given file
| yw hello world                 yank from string

```



```
| ywx 10203040    yank from hexpairs (same as yfx)
| yx              print contents of clipboard in hexadecimal
| yy 0x3344       paste clipboard
| yz [len]        copy nul-terminated string (up to blocksize) into clipboard
```

Sample session:

```
[0x00000000]> s 0x100    ; seek at 0x100
[0x00000100]> y 100      ; yanks 100 bytes from here
[0x00000200]> s 0x200    ; seek 0x200
[0x00000200]> yy         ; pastes 100 bytes
```

You can perform a yank and paste in a single line by just using the `yt` command (yank-to). The syntax is as follows:

```
[0x4A13B8C0]> x
  offset  0 1  2 3  4 5  6 7  8 9  A B  0123456789AB
0x4A13B8C0, 89e0 e839 0700 0089 c7e8 e2ff ...9.....
0x4A13B8CC, ffff 81c3 eea6 0100 8b83 08ff .....
0x4A13B8D8, ffff 5a8d 2484 29c2          ..Z.$.).
```

```
[0x4A13B8C0]> yt 8 0x4A13B8CC @ 0x4A13B8C0
```

```
[0x4A13B8C0]> x
  offset  0 1  2 3  4 5  6 7  8 9  A B  0123456789AB
0x4A13B8C0, 89e0 e839 0700 0089 c7e8 e2ff ...9.....
0x4A13B8CC, 89e0 e839 0700 0089 8b83 08ff ...9.....
0x4A13B8D8, ffff 5a8d 2484 29c2          ..Z.$.).
```

Zoom

The zoom is a print mode that allows you to get a global view of the whole file or a memory map on a single screen. In this mode, each byte represents `file_size/block_size` bytes of the file. Use the `pz` command, or just use `Z` in the visual mode to toggle the zoom mode.

The cursor can be used to scroll faster through the zoom out view. Pressing `z` again will zoom-in at the cursor position.

```
[0x004048c5]> pz?
|Usage: pz [len] print zoomed blocks (filesize/N)
| e zoom.maxsz  max size of block
| e zoom.from   start address
| e zoom.to     end address
| e zoom.byte   specify how to calculate each byte
| pzp          number of printable chars
| pzf          count of flags in block
| pzs          strings in range
| pz0          number of bytes with value '0'
```

```

| pzF          number of bytes with value 0xFF
| pze          calculate entropy and expand to 0-255 range
| pzh          head (first byte value); This is the default mode

```

Let's see some examples:

```

[0x08049790]> e zoom.byte=h
[0x08049790]> pz // or default pzh
0x00000000  7f00 0000 e200 0000 146e 6f74 0300 0000
0x00000010  0000 0000 0068 2102 00ff 2024 e8f0 007a
0x00000020  8c00 18c2 ffff 0080 4421 41c4 1500 5dff
0x00000030  ff10 0018 0fc8 031a 000c 8484 e970 8648
0x00000040  d68b 3148 348b 03a0 8b0f c200 5d25 7074
0x00000050  7500 00e1 ffe8 58fe 4dc4 00e0 dbc8 b885

```

```

[0x08049790]> e zoom.byte=p
[0x08049790]> pz // or pzp
0x00000000  2f47 0609 070a 0917 1e9e a4bd 2a1b 2c27
0x00000010  322d 5671 8788 8182 5679 7568 82a2 7d89
0x00000020  8173 7f7b 727a 9588 a07b 5c7d 8daf 836d
0x00000030  b167 6192 a67d 8aa2 6246 856e 8c9b 999f
0x00000040  a774 96c3 b1a4 6c8e a07c 6a8f 8983 6a62
0x00000050  7d66 625f 7ea4 7ea6 b4b6 8b57 a19f 71a2

```

```

[0x08049790]> eval zoom.byte = flags
[0x08049790]> pz // or pzf
0x00406e65  48d0 80f9 360f 8745 ffff ffeb ae66 0f1f
0x00406e75  4400 0083 f801 0f85 3fff ffff 410f b600
0x00406e85  3c78 0f87 6301 0000 0fb6 c8ff 24cd 0026
0x00406e95  4100 660f 1f84 0000 0000 0084 c074 043c
0x00406ea5  3a75 18b8 0500 0000 83f8 060f 95c0 e9cd
0x00406eb5  feff ff0f 1f84 0000 0000 0041 8801 4983
0x00406ec5  c001 4983 c201 4983 c101 e9ec feff ff0f

```

```

[0x08049790]> e zoom.byte=F
[0x08049790]> p0 // or pzF
0x00000000  0000 0000 0000 0000 0000 0000 0000 0000
0x00000010  0000 2b5c 5757 3a14 331f 1b23 0315 1d18
0x00000020  222a 2330 2b31 2e2a 1714 200d 1512 383d
0x00000030  1e1a 181b 0a10 1a21 2a36 281e 1d1c 0e11
0x00000040  1b2a 2f22 2229 181e 231e 181c 1913 262b
0x00000050  2b30 4741 422f 382a 1e22 0f17 0f10 3913

```

You can limit zooming to a range of bytes instead of the whole bytespace. Change `zoom.from` and `zoom.to` eval variables:

```

[0x00003a04]> e? zoom.
zoom.byte: Zoom callback to calculate each byte (See pz? for help)
zoom.from: Zoom start address

```

```
zoom.in: Specify boundaries for zoom
zoom.maxsz: Zoom max size of block
zoom.to: Zoom end address
[0x00003a04]> e zoom.
zoom.byte = h
zoom.from = 0
zoom.in = io.map
zoom.maxsz = 512
zoom.to = 0
```

Charset

Sometimes, a binary can contain custom characters in another charsets than `ascii`. It is as example the case of the gameboy. The gameboy has his own custom charset for each game. To select as example a custom charset corresponding to the game pokemon red and blue of the gameboy, do:

```
e cfg.charset=pokered;
```

You can now do commands such as `ps` and `w` with custom charset.

Sadly, as you can imagine, very often the charset will be missing. There can not be as many plugin as there are games for gameboy. If the charset does not exists yet, fell free to write it. See the chapters `plugins` and then ‘charset’ to read about charsets.

It works of course for each architecture. Not only the gameboy.

You can also use visual mode with the custom charset for single char encodings.

Colors

Console access is wrapped in API that permits to show the output of any command as ANSI, W32 Console or HTML formats. This allows radare’s core to run inside environments with limited displaying capabilities, like kernels or embedded devices. It is still possible to receive data from it in your favorite format.

To enable colors support by default, add a corresponding configuration option to the `.radare2` configuration file:

```
$ echo 'e scr.color=1' >> ~/.radare2rc
```

Note that enabling colors is not a boolean option. Instead, it is a number because there are different color depth levels. This is:

- 0: black and white
- 1: 16 basic ANSI colors
- 2: 256 scale colors
- 3: 24bit true color

The reason for having such user-defined options is because there's no standard or portable way for the terminal programs to query the console to determine the best configuration, same goes for charset encodings, so r2 allows you to choose that by hand.

Usually, serial consoles may work with 0 or 1, while xterms may support up to 3. RCons will try to find the closest color scheme for your theme when you choose a different them with the `eco` command.

It is possible to configure the color of almost any element of disassembly output. For *NIX terminals, r2 accepts color specification in RGB format. To change the console color palette use `ec` command.

Type `ec` to get a list of all currently used colors. Type `ecs` to show a color palette to pick colors from:

img

Themes

You can create your own color theme, but radare2 have its own predefined ones. Use the `eco` command to list or select them.

After selecting one, you can compare between the color scheme of the shell and the current theme by pressing Ctrl-Shift and then right arrow key for the toggle.

In visual mode use the `R` key to randomize colors or choose the next theme in the list.

Configuration Variables

Below is a list of the most frequently used configuration variables. You can get a complete list by issuing `e` command without arguments. For example, to see all variables defined in the "cfg" namespace, issue `e cfg.` (mind the ending dot). You can get help on any eval configuration variable by using `e? cfg.`

The `e??` command to get help on all the evaluable configuration variables of radare2. As long as the output of this command is pretty large you can combine it with the internal `grep ~` to filter for what you are looking for:

`e??~color`

The Visual mode has an eval browser that is accessible through the `Vbe` command.

asm.arch

Defines the target CPU architecture used for disassembling (`pd`, `pD` commands) and code analysis (`a` command). You can find the list of possible values by looking at the result of `e asm.arch=?` or `rasm2 -L`. It is quite simple to add new architectures for disassembling and analyzing code. There is an interface for

that. For x86, it is used to attach a number of third-party disassembler engines, including GNU binutils, Udis86 and a few handmade ones.

asm.bits

Determines width in bits of registers for the current architecture. Supported values: 8, 16, 32, 64. Note that not all target architectures support all combinations for asm.bits.

asm.syntax

Changes syntax flavor for disassembler between Intel and AT&T. At the moment, this setting affects Udis86 disassembler for Intel 32/Intel 64 targets only. Supported values are `intel` and `att`.

asm.pseudo

A boolean value to set the psuedo syntax in the disassembly. “False” indicates a native one, defined by the current architecture, “true” activates a pseudocode strings format. For example, it’ll transform :

	0x080483ff	e832000000	call 0x8048436
	0x08048404	31c0	xor eax, eax
	0x08048406	0205849a0408	add al, byte [0x8049a84]
	0x0804840c	83f800	cmp eax, 0
	0x0804840f	7405	je 0x8048416

to

	0x080483ff	e832000000	0x8048436 ()
	0x08048404	31c0	eax = 0
	0x08048406	0205849a0408	al += byte [0x8049a84]
	0x0804840c	83f800	var = eax - 0
	0x0804840f	7405	if (!var) goto 0x8048416

It can be useful while disassembling obscure architectures.

asm.os

Selects a target operating system of currently loaded binary. Usually, OS is automatically detected by `rabin -rI`. Yet, `asm.os` can be used to switch to a different syscall table employed by another OS.

asm.flags

If defined to “true”, disassembler view will have flags column.

asm.lines.call

If set to “true”, draw lines at the left of the disassemble output (**pd**, **pD** commands) to graphically represent control flow changes (jumps and calls) that are targeted inside current block. Also, see **asm.lines.out**.

asm.lines.out

When defined as “true”, the disassembly view will also draw control flow lines that go outside of the block.

asm.linestyle

A boolean value which changes the direction of control flow analysis. If set to “false”, it is done from top to bottom of a block; otherwise, it goes from bottom to top. The “false” setting seems to be a better choice for improved readability and is the default one.

asm.offset

Boolean value which controls the visibility of offsets for individual disassembled instructions.

asm.trace

A boolean value that controls displaying of tracing information (sequence number and counter) at the left of each opcode. It is used to assist with programs trace analysis.

asm.bytes

A boolean value used to show or hide displaying of raw bytes of instructions.

asm.sub.reg

A boolean value used to replace register names with arguments or their associated role alias.

For example, if you have something like this:

	0x080483ea	83c404	add esp, 4
	0x080483ed	68989a0408	push 0x8049a98
	0x080483f7	e870060000	call sym.imp.scanf
	0x080483fc	83c408	add esp, 8
	0x08048404	31c0	xor eax, eax

This variable changes it to:

	0x080483ea	83c404	add SP, 4
	0x080483ed	68989a0408	push 0x8049a98

	0x080483f7	e870060000	call sym.imp.scanf
	0x080483fc	83c408	add SP, 8
	0x08048404	31c0	xor A0, A0

asm.sub.jump

A boolean value used to substitute jump, call and branch targets in disassembly.

For example, when turned on, it'd display `jal 0x80001a40` as `jal fcn.80001a40` in the disassembly.

asm.sub.rel

A boolean value which substitutes pc relative expressions in disassembly. When turned on, it shows the references as string references.

For example:

```
0x5563844a0181      488d3d7c0e00. lea rdi, [rip + 0xe7c]      ; str.argv__2d:__s
```

When turned on, this variable lets you display the above instruction as:

```
0x5563844a0181      488d3d7c0e00. lea rdi, str.argv__2d:__s      ; 0x5563844a1004 ; "argv[%2
```

asm.sub.section

Boolean which shows offsets in disassembly prefixed with the name of the section or map.

That means, from something like:

```
0x000067ea      488d0def0c01. lea rcx, [0x000174e0]
```

to the one below, when toggled on.

```
0x000067ea      488d0def0c01. lea rcx, [fmap.LOAD1.0x000174e0]
```

asm.sub.varonly

Boolean which substitutes the variable expression with the local variable name.

For example: `var_14h` as `rbp - var_14h`, in the disassembly.

cfg.bigendian

Change endianness. “true” means big-endian, “false” is for little-endian. “file.id” and “file.flag” both to be true.

cfg.newtab

If this variable is enabled, help messages will be displayed along with command names in tab completion for commands.

scr.color

This variable specifies the mode for colorized screen output: “false” (or 0) means no colors, “true” (or 1) means 16-colors mode, 2 means 256-colors mode, 3 means 16 million-colors mode. If your favorite theme looks weird, try to bump this up.

scr.seek

This variable accepts a full-featured expression or a pointer/flag (eg. eip). If set, radare will set seek position to its value on startup.

scr.scrollbar

If you have set up any flagzones (fz?), this variable will let you display the scrollbar with the flagzones, in Visual mode. Set it to 1 to display the scrollbar at the right end, 2 for the top and 3 to display it at the bottom.

scr.utf8

A boolean variable to show UTF-8 characters instead of ANSI.

cfg.fortunes

Enables or disables “fortune” messages displayed at each radare start.

cfg.fortunes.type

Fortunes are classified by type. This variable determines which types are allowed for displaying when **cfg.fortunes** is **true**, so they can be fine-tuned on what’s appropriate for the intended audience. Current types are **tips**, **fun**, **nsfw**, **creepy**.

stack.size

This variable lets you set the size of stack in bytes.

Files

Use **r2 -H** to list all the environment variables that matter to know where it will be looking for files. Those paths depend on the way (and operating system) you have built r2 for.

```
R2_PREFIX=/usr
MAGICPATH=/usr/share/radare2/2.8.0-git/magic
PREFIX=/usr
INCDIR=/usr/include/libr
LIBDIR=/usr/lib64
LIBEXT=so
RCONFIGHOME=/home/user/.config/radare2
```



```
RDATAHOME=/home/user/.local/share/radare2
RCACHEHOME=/home/user/.cache/radare2
LIBR_PLUGINS=/usr/lib/radare2/2.8.0-git
USER_PLUGINS=/home/user/.local/share/radare2/plugins
USER_ZIGNS=/home/user/.local/share/radare2/zigns
```

RC Files

RC files are r2 scripts that are loaded at startup time. Those files must be in 3 different places:

System

radare2 will first try to load `/usr/share/radare2/radare2rc`

Your Home

Each user in the system can have its own r2 scripts to run on startup to select the color scheme, and other custom options by having r2 commands in there.

- `~/.radare2rc`
- `~/.config/radare2/radare2rc`
- `~/.config/radare2/radare2rc.d/`

Target file

If you want to run a script everytime you open a file, just create a file with the same name of the file but appending `.r2` to it.

Configuration

The core reads `~/.config/radare2/radare2rc` while starting. You can add `e` commands to this file to tune the radare2 configuration to your taste.

To prevent radare2 from parsing this file at startup, pass it the `-N` option.

All the configuration of radare2 is done with the `eval` commands. A typical startup configuration file looks like this:

```
$ cat ~/.radare2rc
e scr.color = 1
e dbg.bep    = loader
```

The configuration can also be changed with `-e <config=value>` command-line option. This way you can adjust configuration from the command line, keeping the `.radare2rc` file intact. For example, to start with empty configuration and then adjust `scr.color` and `asm.syntax` the following line may be used:

```
$ radare2 -N -e scr.color=1 -e asm.syntax=intel -d /bin/ls
```

Internally, the configuration is stored in a hash table. The variables are grouped in namespaces: `cfg.`, `file.`, `dbg.`, `scr.` and so on.

To get a list of all configuration variables just type `e` in the command line prompt. To limit the output to a selected namespace, pass it with an ending dot to `e`. For example, `e file.` will display all variables defined inside the “file” namespace.

To get help about `e` command type `e?`:

```
Usage: e [var[=value]]  Evaluable vars
| e?asm.bytes          show description
| e??                  list config vars with description
| e a                  get value of var 'a'
| e a=b                set var 'a' the 'b' value
| e var=?              print all valid values of var
| e var=??             print all valid values of var with description
| e.a=b                same as 'e a=b' but without using a space
| e,k=v,k=v,k=v       comma separated k[=v]
| e-                   reset config vars
| e*                   dump config vars in r commands
| e!a                  invert the boolean value of 'a' var
| ec [k] [color]       set color for given key (prompt, offset, ...)
| eevar                open editor to change the value of var
| ed                   open editor to change the ~/.radare2rc
| ej                   list config vars in JSON
| env [k[=v]]          get/set environment variable
| er [key]              set config key as readonly. no way back
| es [space]           list all eval spaces [or keys]
| et [key]             show type of given config variable
| ev [key]             list config vars in verbose format
| evj [key]            list config vars in verbose format in JSON
```

A simpler alternative to the `e` command is accessible from the visual mode. Type `Ve` to enter it, use arrows (up, down, left, right) to navigate the configuration, and `q` to exit it. The start screen for the visual configuration edit looks like this:

```
[EvalSpace]

> anal
  asm
  scr
  asm
  bin
  cfg
  diff
  dir
  dbg
  cmd
```

```
fs
hex
http
graph
hud
scr
search
io
```

For configuration values that can take one of several values, you can use the `=?` operator to get a list of valid values:

```
[0x00000000]> e scr.nkey = ?
scr.nkey = fun, hit, flag
```

Crackmes

Crackmes (from “crack me” challenge) are the training ground for reverse engineering people. This section will go over tutorials on how to defeat various crackmes using `r2`.

Authors & Contributors

This book wouldn’t be possible without the help of a large list of contributors who have been reviewing, writing and reporting bugs and stuff in the `radare2` project as well as in this book.

The `radare2` book

This book was started by `mai jin` as a new version of the original `radare` book written by `pancake`.

- Old `radare1` book <http://www.radare.org/get/radare.pdf>

Many thanks to everyone who has been involved with the gitbook:

Adrian Studer, Ahmed Mohamed Abd El-MAwgood, Akshay Krishnan R, Andrew Hoog, Anton Kochkov, Antonio Sánchez, Austin Hartzheim, Aswin C (official-cjunior), Bob131, DZ_ruyk, David Tomaschik, Eric, Fangrui Song, Francesco Tamagni, FreeArtMan, Gerardo García Peña, Giuseppe, Grigory Rechistov, Hui Peng, ITAYC0HEN, Itay Cohen, Jeffrey Crowell, John, Judge Dredd (key 6E23685A), Jupiter, Kevin Grandemange, Kevin Laeuffer, Luca Di Bartolomeo, Lukas Dresel, Maijin, Michael Scherer, Mike, Nikita Abdullin, Paul, Paweł Łukasik, Peter C, RandomLive, Ren Kimura, Reto Schneider, SchumBlubBlub, SkUaTeR, Solomon, Srimanta Barua, Sushant Dinesh, TDKPS, Thanat0s, Vanellope, Vex Woo, Vorlent, XYlearn, Yuri Slobodyanyuk, ali, aoighost, condret, hdznrrd, izhuer, jvoisin, kij, madblobfish, muzlightbeer, pancake, polym (Tim), puddl3glum, radare, sghctoma, shakreiner, sivaramaaa, taiyu, vane1lope, xarkes.

Files

The radare2 debugger allows the user to list and manipulate the file descriptors from the target process.

This is a useful feature, which is not found in other debuggers, the functionality is similar to the lsof command line tool, but have extra subcommands to change the seek, close or duplicate them.

So, at any time in the debugging session you can replace the stdio file descriptors to use network sockets created by r2, or replace a network socket connection to hijack it.

This functionality is also available in r2frida by using the dd command prefixed with a backslash. In r2 you may want to see the output of dd? for proper details.

Getting Started

Small session in radare2 debugger

- **r2 -d /bin/ls**: Opens radare2 with file **/bin/ls** in debugger mode using the radare2 native debugger, but does not run the program. You'll see a prompt (radare2) - all examples are from this prompt.
- **db flag**: place a breakpoint at flag, where flag can be either an address or a function name
- **db - flag**: remove the breakpoint at flag, where flag can be either an address or a function name
- **db**: show list of breakpoint
- **dc**: run the program
- **dr**: Show registers state
- **dr**: Show registers references (telescoping) (like peda)
- **ds**: Step into instruction
- **dso**: Step over instruction
- **dbt**: Display backtrace
- **dm**: Show memory maps
- **dk <signal>**: Send KILL signal to child
- **ood**: reopen in debug mode
- **ood arg1 arg2**: reopen in debug mode with arg1 and arg2

Heap

radare2's `dm` subcommands can also display a map of the heap which is useful for those who are interested in inspecting the heap and its content. Simply execute `dmh` to show a map of the heap:

```
[0x7fae46236ca6]> dmh
  Malloc chunk @ 0x55a7ecbce250 [size: 0x411][allocated]
  Top chunk @ 0x55a7ecbce660 - [brk_start: 0x55a7ecbce000, brk_end: 0x55a7ecbef000]
```

You can also see a graph layout of the heap:

```
[0x7fae46236ca6]> dmhg
Heap Layout
```

```

  .------.
  |   Malloc chunk @ 0x55a7ecbce000   |
  | size: 0x251                       |
  | fd: 0x0, bk: 0x0                 |
  `-----'
  |
  .---'
  |
  |
  .------.
  |   Malloc chunk @ 0x55a7ecbce250   |
  | size: 0x411                       |
  | fd: 0x57202c6f6c6c6548, bk: 0xa21646c726f |
  `-----'
  |
  .---'
  |
  |
  .------.
  | Top chunk @ 0x55a7ecbce660         |
  | [brk_start:0x55a7ecbce000, brk_end:0x55a7ecbef000] |
  `-----'
```

Another heap commands can be found under `dmh`, check `dmh?` for the full list.

```
[0x00000000]> dmh?
|Usage:  dmh # Memory map heap
| dmh                List chunks in heap segment
| dmh [malloc_state] List heap chunks of a particular arena
| dmha               List all malloc_state instances in application
| dmhb               Display all parsed Double linked list of main_arena's bins instance
| dmhb [bin_num|bin_num:malloc_state]      Display parsed double linked list of bins in
| dmhbg [bin_num]    Display double linked list graph of main_arena's bin [Under developem
| dmhc @[chunk_addr] Display malloc_chunk struct for a given malloc chunk
```

```

| dmhf          Display all parsed fastbins of main_arena's fastbinY instance
| dmhf [fastbin_num|fastbin_num:malloc_state] Display parsed single linked list in fastbinY
| dmhg          Display heap graph of heap segment
| dmhg [malloc_state] Display heap graph of a particular arena
| dmhi @[malloc_state] Display heap_info structure/structures for a given arena
| dmhm          List all elements of struct malloc_state of main thread (main_arena)
| dmhm [malloc_state] List all malloc_state instance of a particular arena
| dmht          Display all parsed thread cache bins of main_arena's tcache instance
| dmh?          Show map heap help

```

To print safe-linked lists (glibc >= 2.32) with demangled pointers, the variable `dbg.glibc.demangle` must be true.

Debugger

Debuggers are implemented as IO plugins. Therefore, radare can handle different URI types for spawning, attaching and controlling processes. The complete list of IO plugins can be viewed with `r2 -L`. Those that have “d” in the first column (“rwd”) support debugging. For example:

```

r_d  debug      Debug a program or pid. dbg:///bin/ls, dbg://1388 (LGPL3)
rwd  gdb        Attach to gdbserver, 'qemu -s', gdb://localhost:1234 (LGPL3)

```

There are different backends for many target architectures and operating systems, e.g., GNU/Linux, Windows, MacOS X, (Net,Free,Open)BSD and Solaris.

Process memory is treated as a plain file. All mapped memory pages of a debugged program and its libraries can be read and interpreted as code or data structures.

Communication between radare and the debugger IO layer is wrapped into `system()` calls, which accept a string as an argument, and executes it as a command. An answer is then buffered in the output console, its contents can be additionally processed by a script. Access to the IO system is achieved with `!=`. Most IO plugins provide help with `!=?` or `!=help`. For example:

```

$ r2 -d /bin/ls
...
[0x7fc15afa3cc0]> !=help
Usage: !=cmd args
!=ptrace  - use ptrace io
!=mem     - use /proc/pid/mem io if possible
!=pid     - show targeted pid
!=pid <#> - select new pid

```

In general, debugger commands are portable between architectures and operating systems. Still, as radare tries to support the same functionality for all target architectures and operating systems, certain things have to be handled separately.

They include injecting shellcodes and handling exceptions. For example, in MIPS targets there is no hardware-supported single-stepping feature. In this case, radare2 provides its own implementation for single-step by using a mix of code analysis and software breakpoints.

To get basic help for the debugger, type 'd?':

```
Usage: d    # Debug commands
| db[?]          Breakpoints commands
| dbt[?]         Display backtrace based on dbg.btdepth and dbg.btalgo
| dc[?]         Continue execution
| dd[?]         File descriptors (!fd in r1)
| de[-sc] [perm] [rm] [e] Debug with ESIL (see de?)
| dg <file>      Generate a core-file (WIP)
| dh [handler]   Transplant process to a new handler
| di[?]         Show debugger backend information (See dh)
| dk[?]         List, send, get, set, signal handlers of child
| dL[?]         List or set debugger handler
| dm[?]         Show memory maps
| do[?]         Open process (reload, alias for 'oo')
| doo[args]      Reopen in debug mode with args (alias for 'ood')
| doof[file]     Reopen in debug mode from file (alias for 'oodf')
| doc           Close debug session
| dp[?]         List, attach to process or thread id
| dr[?]         Cpu registers
| ds[?]         Step, over, source line
| dt[?]         Display instruction traces
| dw <pid>      Block prompt until pid dies
| dx[?]         Inject and run code on target process (See gs)
```

To restart your debugging session, you can type `oo` or `oo+`, depending on desired behavior.

```
oo           reopen current file (kill+fork in debugger)
oo+          reopen current file in read-write
```

Memory Maps

The ability to understand and manipulate the memory maps of a debugged program is important for many different Reverse Engineering tasks. radare2 offers a rich set of commands to handle memory maps in the binary. This includes listing the memory maps of the currently debugged binary, removing memory maps, handling loaded libraries and more.

First, let's see the help message for `dm`, the command which is responsible for handling memory maps:

```
[0x55f2104cf620]> dm?
```

```

Usage: dm    # Memory maps commands
| dm                    List memory maps of target process
| dm address size      Allocate <size> bytes at <address> (anywhere if address 1
| dm=                  List memory maps of target process (ascii-art bars)
| dm.                  Show map name of current address
| dm*                  List memmaps in radare commands
| dm- address          Deallocate memory map of <address>
| dmd[a] [file]        Dump current (all) debug map region to a file (from-to.dm
| dmh[?]               Show map of heap
| dmi [addr|libname] [symname] List symbols of target lib
| dmi* [addr|libname] [symname] List symbols of target lib in radare commands
| dmi.                 List closest symbol to the current address
| dmiv                 Show address of given symbol for given lib
| dmj                  List memmaps in JSON format
| dml <file>           Load contents of file into the current map region
| dmm[?][j*]           List modules (libraries, binaries loaded in memory)
| dmp[?] <address> <size> <perms> Change page at <address> with <size>, protection <perms>
| dms[?] <id> <mapaddr> Take memory snapshot
| dms- <id> <mapaddr> Restore memory snapshot
| dmS [addr|libname] [sectname] List sections of target lib
| dmS* [addr|libname] [sectname] List sections of target lib in radare commands
| dml address size      Allocate <size> bytes at <address> and promote to huge pa

```

In this chapter, we'll go over some of the most useful subcommands of `dm` using simple examples. For the following examples, we'll use a simple `helloworld` program for Linux but it'll be the same for every binary.

First things first - open a program in debugging mode:

```

$ r2 -d helloworld
Process with PID 20304 started...
= attach 20304 20304
bin.baddr 0x56136b475000
Using 0x56136b475000
asm.bits 64
[0x7f133f022fb0]>

```

Note that we passed "helloworld" to radare2 without `./`. radare2 will try to find this program in the current directory and then in `$PATH`, even if no `./` is passed. This is contradictory with UNIX systems, but makes the behaviour consistent for windows users

Let's use `dm` to print the memory maps of the binary we've just opened:

```

[0x7f133f022fb0]> dm
0x0000563a0113a000 - usr    4K s r-x /tmp/helloworld /tmp/helloworld ; map.tmp_helloworld.r_x
0x0000563a0113a000 - usr    8K s rw- /tmp/helloworld /tmp/helloworld ; map.tmp_helloworld.rw
0x00007f133f022000 * usr 148K s r-x /usr/lib/ld-2.27.so /usr/lib/ld-2.27.so ; map.usr_lib_ld
0x00007f133f246000 - usr    8K s rw- /usr/lib/ld-2.27.so /usr/lib/ld-2.27.so ; map.usr_lib_ld

```



```

0x00007f133f248000 - usr    4K s rw- unk0 unk0 ; map.unk0.rw
0x00007ffffd25ce000 - usr 132K s rw- [stack] [stack] ; map.stack_.rw
0x00007ffffd25f6000 - usr  12K s r-- [vvar] [vvar] ; map.vvar_.r
0x00007ffffd25f9000 - usr   8K s r-x [vdso] [vdso] ; map.vdso_.r_x
0xfffffffff600000 - usr   4K s r-x [vsyscall] [vsyscall] ; map.vsyscall_.r_x

```

For those of you who prefer a more visual way, you can use `dm=` to see the memory maps using an ASCII-art bars. This will be handy when you want to see how these maps are located in the memory.

If you want to know the memory-map you are currently in, use `dm.:`

```

[0x7f133f022fb0]> dm.
0x00007f947eed9000 # 0x00007f947eefe000 * usr    148K s r-x /usr/lib/ld-2.27.so /usr/lib/ld-2.27.so

```

Using `dmm` we can “List modules (libraries, binaries loaded in memory)”, this is quite a handy command to see which modules were loaded.

```

[0x7fa80a19dfb0]> dmm
0x55ca23a4a000 /tmp/helloworld
0x7fa80a19d000 /usr/lib/ld-2.27.so

```

Note that the output of `dm` subcommands, and `dmm` specifically, might be different in various systems and different binaries.

We can see that along with our `helloworld` binary itself, another library was loaded which is `ld-2.27.so`. We don’t see `libc` yet and this is because `radare2` breaks before `libc` is loaded to memory. Let’s use `dcu` (**d**ebug **c**ontinue **u**ntil) to execute our program until the entry point of the program, which `radare` flags as `entry0`.

```

[0x7fa80a19dfb0]> dcu entry0
Continue until 0x55ca23a4a520 using 1 bpsize
hit breakpoint at: 55ca23a4a518
[0x55ca23a4a520]> dmm
0x55ca23a4a000 /tmp/helloworld
0x7fa809de1000 /usr/lib/libc-2.27.so
0x7fa80a19d000 /usr/lib/ld-2.27.so

```

Now we can see that `libc-2.27.so` was loaded as well, great!

Speaking of `libc`, a popular task for binary exploitation is to find the address of a specific symbol in a library. With this information in hand, you can build, for example, an exploit which uses ROP. This can be achieved using the `dmi` command. So if we want, for example, to find the address of `system()` in the loaded `libc`, we can simply execute the following command:

```

[0x55ca23a4a520]> dmi libc system
514 0x00000000 0x7fa809de1000 LOCAL FILE    0 system.c
515 0x00043750 0x7fa809e24750 LOCAL FUNC 1221 do_system
4468 0x001285a0 0x7fa809f095a0 LOCAL FUNC 100 svcerr_systemerr

```

```

5841 0x001285a0 0x7fa809f095a0 LOCAL FUNC 100 svcerr_systemerr
6427 0x00043d10 0x7fa809e24d10 WEAK FUNC 45 system
7094 0x00043d10 0x7fa809e24d10 GLOBAL FUNC 45 system
7480 0x001285a0 0x7fa809f095a0 GLOBAL FUNC 100 svcerr_systemerr

```

Similar to the `dm.` command, with `dmi.` you can see the closest symbol to the current address.

Another useful command is to list the sections of a specific library. In the following example we'll list the sections of `ld-2.27.so`:

```

[0x55a7ebf09520]> dmS ld-2.27
[Sections]
00 0x00000000      0 0x00000000      0 ---- ld-2.27.so.
01 0x000001c8     36 0x4652d1c8     36 -r-- ld-2.27.so..note.gnu.build_id
02 0x000001f0    352 0x4652d1f0    352 -r-- ld-2.27.so..hash
03 0x00000350    412 0x4652d350    412 -r-- ld-2.27.so..gnu.hash
04 0x000004f0    816 0x4652d4f0    816 -r-- ld-2.27.so..dynsym
05 0x00000820    548 0x4652d820    548 -r-- ld-2.27.so..dynstr
06 0x00000a44     68 0x4652da44     68 -r-- ld-2.27.so..gnu.version
07 0x00000a88    164 0x4652da88    164 -r-- ld-2.27.so..gnu.version_d
08 0x00000b30   1152 0x4652db30   1152 -r-- ld-2.27.so..rela.dyn
09 0x00000fb0 11497 0x4652dfb0 11497 -r-x ld-2.27.so..text
10 0x0001d0e0 17760 0x4654a0e0 17760 -r-- ld-2.27.so..rodata
11 0x00021640   1716 0x4654e640   1716 -r-- ld-2.27.so..eh_frame_hdr
12 0x00021cf8   9876 0x4654ecf8   9876 -r-- ld-2.27.so..eh_frame
13 0x00024660   2020 0x46751660   2020 -rw- ld-2.27.so..data.rel.ro
14 0x00024e48    336 0x46751e48    336 -rw- ld-2.27.so..dynamic
15 0x00024f98     96 0x46751f98     96 -rw- ld-2.27.so..got
16 0x00025000   3960 0x46752000   3960 -rw- ld-2.27.so..data
17 0x00025f78      0 0x46752f80    376 -rw- ld-2.27.so..bss
18 0x00025f78     17 0x00000000     17 ---- ld-2.27.so..comment
19 0x00025fa0     63 0x00000000     63 ---- ld-2.27.so..gnu.warning.llseek
20 0x00025fe0 13272 0x00000000 13272 ---- ld-2.27.so..symtab
21 0x000293b8   7101 0x00000000   7101 ---- ld-2.27.so..strtab
22 0x0002af75    215 0x00000000    215 ---- ld-2.27.so..shstrtab

```

Migration from ida, GDB or WinDBG

How to run the program using the debugger

`r2 -d /bin/ls - start in debugger mode => [video]`

How do I attach/detach to running process ? (gdb -p)

`r2 -d <pid> - attach to process`

```
r2 ptrace://pid - same as above, but only for io (not debugger backend hooked)
[0x7fff6ad90028]> o-225 - close fd=225 (listed in o~[1]:0)
r2 -D gdb gdb://localhost:1234 - attach to gdbserver
```

How to set args/environment variable/load a specific libraries for the debugging session of radare

Use rarun2 (libpath=\$PWD:/tmp/lib, arg2=hello, setenv=F00=BAR ...) see rarun2 -h / man rarun2

How to script radare2 ?

```
r2 -i <scriptfile> ... - run a script after loading the file => [video]
r2 -I <scriptfile> ... - run a script before loading the file
r2 -c $@ | awk $@ - run through awk to get asm from function => [link]
[0x80480423]> . scriptfile - interpret this file => [video]
[0x80480423]> #!c - enter C repl (see #! to list all available RLang plugins)
=> [video], everything have to be done in a oneliner or a .c file must be passed
as an argument.
```

To get #!python and much more, just build radare2-bindings

How to list Source code as in gdb list ?

CL @ sym.main - though the feature is highly experimental

shortcuts

Command	IDA Pro	radare2	r2 (visual mode)	GDB	WinDbg
Analysis					
Analysis of everything	Automatically launched when opening a binary	aaa or -A (aaaa or -AA for even experimental analysis)	N/A	N/A	N/A
Navigation					
xref to	x	axt	x	N/A	N/A
xref from	ctrl + j	axf	X	N/A	N/A
xref to graph	?	agt [offset]	?	N/A	N/A
xref from graph	?	agf [offset]	?	N/A	N/A

Command	IDA Pro	radare2	r2 (visual mode)	GDB	WinDbg
list functions	alt + 1	afl;is	t	N/A	N/A
listing	alt + 2	pdf	p	N/A	N/A
hex mode	alt + 3	pxa	P	N/A	N/A
imports	alt + 6	ii	:ii	N/A	N/A
exports	alt + 7	is~FUNC	?	N/A	N/A
follow	enter	s offset	enter or 0-9	N/A	N/A
jmp/call					
undo seek	esc	s-	u	N/A	N/A
redo seek	ctrl+enter	s+	U	N/A	N/A
show graph	space	agv	V	N/A	N/A
Edit					
rename	n	afn	dr	N/A	N/A
graph view	space	agv	V	N/A	N/A
define as data	d	Cd [size]	dd,db,dw,dW	N/A	N/A
define as code	c	C- [size]	d- or du	N/A	N/A
define as undefined	u	C- [size]	d- or du	N/A	N/A
define as string	A	Cs [size]	ds	N/A	N/A
define as struct	Alt+Q	Cf [size]	dF	N/A	N/A
Debugger					
Start Process/Continue execution	F9	dc	F9	r and c	g
Terminate Process	Ctrl+F2	dk 9	?	kill	q
Detach	?	o-	?	detach	
step into	F7	ds	s	n	t
step into 4 instructions	?	ds 4	F7	n 4	t 4
step over	F8	dso	S	s	p
step until a specific address	?	dsu <addr>	?	s	g <addr>
Run until return	Ctrl+F7	dcr	?	finishgu	
Run until cursor	F4	#249	#249	N/A	N/A

Command	IDA Pro	radare2	r2 (visual mode)	GDB	WinDbg
Show Backtrace	?	dbt	?	bt	
display Register	On register Windows	dr all	Shown in Visual mode	info r registers	
display eax	On register Windows	dr?eax	Shown in Visual mode	info r registers eax	
display old state of all registers	?	dro	?	? ?	
display function addr + N	?	afi \$\$ - display function information of current offset (\$\$)	?	? ?	
display frame state	?	pxw rbp-rsp@rsp	?	i f ?	
How to step until condition is true	?	dsi	?	? ?	
Update a register value	?	dr rip=0x456	?	set r \$rip=0x456	
Disassembly					
disassembly forward	N/A	pd	Vp	disas uf, u	
disassembly N instructions	N/A	pd X	Vp	x/i u <addr> LX	
disassembly N (backward)	N/A	pd -X	Vp	disas ub <a-o> <a>	
Information on the bin					
Sections/regions	Menu sections	iS or S (append j for json)	N/A	maint !address info sec- tions	
Load symbol file					

Command	IDA Pro	radare2	r2 (visual mode)	GDB	WinDbg
Sections/regions	pdb menu	asm.dwarf.file,N/A pdb.XX)		add- r symbol- file	
BackTrace					
Stack Trace	N/A	dbt	N/A	bt	k
Stack Trace in Json	N/A	dbtj	N/A		
Partial Backtrace (innermost)	N/A	dbt (dbg.btdepth dbg.btalgo)	N/A	bt	k
Partial Backtrace (outermost)	N/A	dbt (dbg.btdepth dbg.btalgo)	N/A	bt -	
Stacktrace for all threads	N/A	dbt@t	N/A	thread~* apply k all bt	
Breakpoints					
Breakpoint list	Ctrl+Alt+B	db	?	info bl	
add breakpoint	F2	db [offset]	F2	breakpoints break bp	
Threads					
Switch to thread	Thread menu	dp	N/A	thread~<N>s <N>	
Frames					
Frame Numbers	N/A	?	N/A	any kn bt command	
Select Frame	N/A	?	N/A	frame .frame	
Parameters/Locals					
Display parameters	N/A	afv	N/A	info dv args /t /i /V	
Display parameters	N/A	afv	N/A	info dv locals/t /i /V	
Display parame- ters/locals in json	N/A	afvj	N/A	info dv locals/t /i /V	

Command	IDA Pro	radare2	r2 (visual mode)	GDB	WinDbg
list addresses where vars are ac- cessed(R/W)	N/A	afvR/afvW	N/A	?	?
Project Related					
open project		Po [file]		?	
save project	automatic	Ps [file]		?	
show project in- formations		Pi [file]		?	
Miscellaneous					
Dump byte char array options	N/A	pc? (json, C, char, etc.)	Vpppp	x/bc	db
	option menu	e?	e		
search	search menu	/?	Select the zone with the cursor c then /		s

Equivalent of “set-follow-fork-mode” gdb command

This can be done using 2 commands:

1. **dcf** - until a fork happen
2. then use **dp** to select what process you want to debug.

Common features

- r2 accepts FLIRT signatures
- r2 can connect to GDB, LLVM and WinDbg
- r2 can write/patch in place
- r2 have fortunes and [s]easter eggs[/s]balls of steel
- r2 can do basic loading of ELF core files from the box and MDMP (Windows minidumps)

Registers

The registers are part of a user area stored in the context structure used by the scheduler. This structure can be manipulated to get and set the values of those

registers, and, for example, on Intel hosts, it is possible to directly manipulate DR0-DR7 hardware registers to set hardware breakpoints.

There are different commands to get values of registers. For the General Purpose ones use:

```
[0x4A13B8C0]> dr
r15 = 0x00000000
r14 = 0x00000000
r13 = 0x00000000
r12 = 0x00000000
rbp = 0x00000000
rbx = 0x00000000
r11 = 0x00000000
r10 = 0x00000000
r9 = 0x00000000
r8 = 0x00000000
rax = 0x00000000
rcx = 0x00000000
rdx = 0x00000000
rsi = 0x00000000
rdi = 0x00000000
oeax = 0x0000003b
rip = 0x7f20bf5df630
rsp = 0x7fff515923c0
```

```
[0x7f0f2dbae630]> dr rip ; get value of 'rip'
0x7f0f2dbae630
```

```
[0x4A13B8C0]> dr rip = esp ; set 'rip' as esp
```

Interaction between a plugin and the core is done by commands returning radare instructions. This is used, for example, to set flags in the core to set values of registers.

```
[0x7f0f2dbae630]> dr* ; Appending '*' will show radare commands
f r15 1 0x0
f r14 1 0x0
f r13 1 0x0
f r12 1 0x0
f rbp 1 0x0
f rbx 1 0x0
f r11 1 0x0
f r10 1 0x0
f r9 1 0x0
f r8 1 0x0
f rax 1 0x0
f rcx 1 0x0
```



```

f rdx 1 0x0
f rsi 1 0x0
f rdi 1 0x0
f oeax 1 0x3b
f rip 1 0x7fff73557940
f rflags 1 0x200
f rsp 1 0x7fff73557940

```

```
[0x4A13B8C0]> .dr* ; include common register values in flags
```

An old copy of registers is stored all the time to keep track of the changes done during execution of a program being analyzed. This old copy can be accessed with `oregs`.

```

[0x7f1fab84c630]> dro
r15 = 0x00000000
r14 = 0x00000000
r13 = 0x00000000
r12 = 0x00000000
rbp = 0x00000000
rbx = 0x00000000
r11 = 0x00000000
r10 = 0x00000000
r9 = 0x00000000
r8 = 0x00000000
rax = 0x00000000
rcx = 0x00000000
rdx = 0x00000000
rsi = 0x00000000
rdi = 0x00000000
oeax = 0x0000003b
rip = 0x7f1fab84c630
rflags = 0x00000200
rsp = 0x7fff386b5080

```

Current state of registers

```

[0x7f1fab84c630]> dr
r15 = 0x00000000
r14 = 0x00000000
r13 = 0x00000000
r12 = 0x00000000
rbp = 0x00000000
rbx = 0x00000000
r11 = 0x00000000
r10 = 0x00000000
r9 = 0x00000000
r8 = 0x00000000

```

```

rax = 0x00000000
rcx = 0x00000000
rdx = 0x00000000
rsi = 0x00000000
rdi = 0x7fff386b5080
oeax = 0xfffffffffffffff
rip = 0x7f1fab84c633
rflags = 0x00000202
rsp = 0x7fff386b5080

```

Values stored in eax, oeax and eip have changed.

To store and restore register values you can just dump the output of 'dr*' command to disk and then re-interpret it again:

```

[0x4A13B8C0]> dr* > regs.saved ; save registers
[0x4A13B8C0]> drp regs.saved ; restore

```

EFLAGS can be similarly altered. E.g., setting selected flags:

```

[0x4A13B8C0]> dr eflags = pst
[0x4A13B8C0]> dr eflags = azsti

```

You can get a string which represents latest changes of registers using drd command (diff registers):

```

[0x4A13B8C0]> drd
oeax = 0x0000003b was 0x00000000 delta 59
rip = 0x7f00e71282d0 was 0x00000000 delta -418217264
rflags = 0x00000200 was 0x00000000 delta 512
rsp = 0x7fffe85a09c0 was 0x00000000 delta -396752448

```

Debugging with gdbserver

radare2 allows remote debugging over the gdb remote protocol. So you can run a gdbserver and connect to it with radare2 for remote debugging. The syntax for connecting is:

```
$ r2 -d gdb://<host>:<port>
```

Note that the following command does the same, r2 will use the debug plugin specified by the uri if found.

```
$ r2 -D gdb gdb://<host>:<port>
```

The debug plugin can be changed at runtime using the dL or Ld commands.

Or if the gdbserver is running in extended mode, you can attach to a process on the host with:

```
$ r2 -d gdb://<host>:<port>/<pid>
```

It is also possible to start debugging after analyzing a file using the `doof` command which rebases the current session's data after opening gdb

```
[0x00404870]> doof gdb://<host>:<port>/<pid>
```

After connecting, you can use the standard r2 debug commands as normal.

radare2 does not yet load symbols from gdbserver, so it needs the binary to be locally present to load symbols from it. In case symbols are not loaded even if the binary is present, you can try specifying the path with `e dbg.exe.path:`

```
$ r2 -e dbg.exe.path=<path> -d gdb://<host>:<port>
```

If symbols are loaded at an incorrect base address, you can try specifying the base address too with `e bin.baddr:`

```
$ r2 -e bin.baddr=<baddr> -e dbg.exe.path=<path> -d gdb://<host>:<port>
```

Usually the gdbserver reports the maximum packet size it supports. Otherwise, radare2 resorts to sensible defaults. But you can specify the maximum packet size with the environment variable `R2_GDB_PKT SZ`. You can also check and set the max packet size during a session with the IO system, `!=`.

```
$ export R2_GDB_PKT SZ=512
$ r2 -d gdb://<host>:<port>
= attach <pid> <tid>
Assuming filepath <path/to/exe>
[0x7ff659d9fcc0]> !=pkt sz
packet size: 512 bytes
[0x7ff659d9fcc0]> !=pkt sz 64
[0x7ff659d9fcc0]> !=pkt sz
packet size: 64 bytes
```

The gdb IO system provides useful commands which might not fit into any standard radare2 commands. You can get a list of these commands with `!=?`. (Remember, `!=` accesses the underlying IO plugin's `system()`).

```
[0x7ff659d9fcc0]> !=?
Usage: !=cmd args
!=pid          - show targeted pid
!=pkt s        - send packet 's'
!=monitor cmd  - hex-encode monitor command and pass to target interpreter
!=rd           - show reverse debugging availability
!=dsb          - step backwards
!=dcb          - continue backwards
!=detach [pid] - detach from remote/detach specific pid
!=inv.reg      - invalidate reg cache
!=pkt sz       - get max packet size used
!=pkt sz bytes - set max. packet size as 'bytes' bytes
!=exec_file [pid] - get file which was executed for current/specified pid
```

Note that `=!dsb` and `=!dcb` are only available in special gdbserver implementations such as Mozilla's rr, the default gdbserver doesn't include remote reverse debugging support. Use `=!rd` to print the currently available reverse debugging capabilities.

If you are interested in debugging radare2's interaction with gdbserver you can use `=!monitor set remote-debug 1` to turn on logging of gdb's remote protocol packets in gdbserver's console and `=!monitor set debug 1` to show general debug messages from gdbserver in it's console.

radare2 also provides its own gdbserver implementation:

```
$ r2 -
[0x00000000]> =g?
|Usage:  =[g] [...] # gdb server
| gdbserver:
| =g port file [args]  listen on 'port' debugging 'file' using gdbserver
| =g! port file [args]  same as above, but debug protocol messages (like gdbserver --remote-
```

So you can start it as:

```
$ r2 -
[0x00000000]> =g 8000 /bin/radare2 -
```

And then connect to it like you would to any gdbserver. For example, with radare2:

```
$ r2 -d gdb://localhost:8000
```

Remote Access Capabilities

Radare can be run locally, or it can be started as a server process which is controlled by a local radare2 process. This is possible because everything uses radare's IO subsystem which abstracts access to `system()`, `cmd()` and all basic IO operations so to work over a network.

Help for commands useful for remote access to radare:

```
[0x00405a04]> =?
Usage:  =[:!+--=ghH] [...]  # connect with other instances of r2
```

remote commands:

=	list all open connections
=<[fd] cmd	send output of local command to remote fd
=[fd] cmd	exec cmd at remote 'fd' (last open is default one)
=! cmd	run command via r_io_system
+= [proto://]host:port	connect to remote host:port (*rap://, raps://, tcp://, udp://)
--[fd]	remove all hosts or host 'fd'
==[fd]	open remote session with host 'fd', 'q' to quit

```

| !==                disable remote cmd mode
| !=!                enable remote cmd mode

servers:
| .:9000              start the tcp server (echo x|nc ::1 9090 or curl ::1:9090/cn
| =:port              start the rap server (o rap://9999)
| =g[?]              start the gdbserver
| =h[?]              start the http webserver
| =H[?]              start the http webserver (and launch the web browser)

other:
| =&:port             start rap server in background (same as '&_h')
| =:host:port cmd     run 'cmd' command on remote server

examples:
| +=tcp://localhost:9090/ connect to: r2 -c.:9090 ./bin
| +=rap://localhost:9090/ connect to: r2 rap://:9090
| +=http://localhost:9090/cmd/ connect to: r2 -c='h 9090' bin
| o rap://:9090/        start the rap server on tcp port 9090

```

You can learn radare2 remote capabilities by displaying the list of supported IO plugins: `radare2 -L`.

A little example should make this clearer. A typical remote session might look like this:

At the remote host1:

```
$ radare2 rap://:1234
```

At the remote host2:

```
$ radare2 rap://:1234
```

At localhost:

```
$ radare2 -
```

Add hosts

```
[0x004048c5]> += rap://<host1>:1234//bin/ls
Connected to: <host1> at port 1234
waiting... ok
```

```
[0x004048c5]> =
0 - rap://<host1>:1234//bin/ls
```

You can open remote files in debug mode (or using any IO plugin) specifying URI when adding hosts:

```
[0x004048c5]> += += rap://<host2>:1234/dbg:///bin/ls
Connected to: <host2> at port 1234
```

```
waiting... ok
0 - rap://<host1>:1234//bin/ls
1 - rap://<host2>:1234/dbg:///bin/ls
```

To execute commands on host1:

```
[0x004048c5]> =0 px
[0x004048c5]> = s 0x666
```

To open a session with host2:

```
[0x004048c5]> ==1
fd:6> pi 1
...
fd:6> q
```

To remove hosts (and close connections):

```
[0x004048c5]> ==
```

You can also redirect radare output to a TCP or UDP server (such as `nc -l`). First, Add the server with `'+= tcp://'` or `'+= udp://'`, then you can redirect the output of a command to be sent to the server:

```
[0x004048c5]> += tcp://<host>:<port>/
Connected to: <host> at port <port>
5 - tcp://<host>:<port>/
[0x004048c5]> =<5 cmd...
```

The `=<` command will send the output from the execution of `cmd` to the remote connection number `N` (or the last one used if no id specified).

Reverse Debugging

Radare2 has reverse debugger, that can seek the program counter backward. (e.g. `reverse-next`, `reverse-continue` in `gdb`) Firstly you need to save program state at the point that you want to start recording. The syntax for recording is:

```
[0x004028a0]> dts+
```

You can use `dts` commands for recording and managing program states. After recording the states, you can seek pc back and forth to any points after saved address. So after recording, you can try single step back:

```
[0x004028a0]> 2dso
[0x004028a0]> dr rip
0x004028ae
[0x004028a0]> dsb
continue until 0x004028a2
hit breakpoint at: 4028a2
[0x004028a0]> dr rip
```

0x004028a2

When you run **dsb**, reverse debugger restore previous recorded state and execute program from it until desired point.

Or you can also try continue back:

```
[0x004028a0]> db 0x004028a2
[0x004028a0]> 10dso
[0x004028a0]> dr rip
0x004028b9
[0x004028a0]> dcb
[0x004028a0]> dr rip
0x004028a2
```

dcb seeks program counter until hit the latest breakpoint. So once set a breakpoint, you can back to it any time.

You can see current recorded program states using **dts**:

```
[0x004028a0]> dts
session: 0    at:0x004028a0    ""
session: 1    at:0x004028c2    ""
```

NOTE: Program records can be saved at any moments. These are diff style format that save only different memory area from previous. It saves memory space rather than entire dump.

And also can add comment:

```
[0x004028c2]> dtsC 0 program start
[0x004028c2]> dtsC 1 decryption start
[0x004028c2]> dts
session: 0    at:0x004028a0    "program start"
session: 1    at:0x004028c2    "decryption start"
```

You can leave notes for each records to keep in your mind. **dsb** and **dcb** commands restore the program state from latest record if there are many records.

Program records can exported to file and of course import it. Export/Import records to/from file:

```
[0x004028c2]> dtst records_for_test
Session saved in records_for_test.session and dump in records_for_test.dump
[0x004028c2]> dtsf records_for_test
session: 0, 0x4028a0 diffs: 0
session: 1, 0x4028c2 diffs: 0
```

Moreover, you can do reverse debugging in ESIL mode. In ESIL mode, program state can be managed by **aets** commands.

```
[0x00404870]> aets+
```

And step back by `aesb`:

```
[0x00404870]> aer rip
0x00404870
[0x00404870]> 5aeso
[0x00404870]> aer rip
0x0040487d
[0x00404870]> aesb
[0x00404870]> aer rip
0x00404879
```

In addition to the native reverse debugging capabilities in radare2, it's also possible to use gdb's remote protocol to reverse debug a target gdbserver that supports it. `=!dsb` and `=!dcb` are available as `dsb` and `dcb` replacements for this purpose, see remote gdb's documentation for more information.

WinDBG Kernel-mode Debugging (KD)

The WinDBG KD interface support for r2 allows you to attach to VM running Windows and debug its kernel over a serial port or network.

It is also possible to use the remote GDB interface to connect and debug Windows kernels without depending on Windows capabilities.

Bear in mind that WinDBG KD support is still work-in-progress, and this is just an initial implementation which will get better in time.

Setting Up KD on Windows

For a complete walkthrough, refer to Microsoft's documentation.

Serial Port

Enable KD over a serial port on Windows Vista and higher like this:

```
bcdedit /debug on
bcdedit /dbgsettings serial debugport:1 baudrate:115200
```

Or like this for Windows XP: Open boot.ini and add `/debug /debugport=COM1 /baudrate=115200`:

```
[boot loader]
timeout=30
default=multi(0)disk(0)rdisk(0)partition(1)\WINDOWS
[operating systems]
multi(0)disk(0)rdisk(0)partition(1)\WINDOWS="Debugging with Cable" /fastdetect /debug /debugport=COM1 /baudrate=115200
```

In case of VMWare


```

Virtual Machine Settings -> Add -> Serial Port
Device Status:
[v] Connect at power on
Connection:
[v] Use socket (named pipe)
[_/tmp/winkd.pipe_-----]
From: Server To: Virtual Machine

```

Configure the VirtualBox Machine like this:

```

Preferences -> Serial Ports -> Port 1

[v] Enable Serial Port
Port Number: [_COM1_-----[v]]
Port Mode:   [_Host_Pipe__[v]]
             [v] Create Pipe
Port/File Path: [_/tmp/winkd.pipe_---]

```

Or just spawn the VM with qemu like this:

```

$ qemu-system-x86_64 -chardev socket,id=serial0,\
    path=/tmp/winkd.pipe,nowait,server \
    -serial chardev:serial0 -hda Windows7-VM.vdi

```

Network

Enable KD over network (KDNet) on Windows 7 or later likes this:

```

bcdedit /debug on
bcdedit /dbgsettings net hostip:w.x.y.z port:n

```

Starting from Windows 8 there is no way to enforce debugging for every boot, but it is possible to always show the advanced boot options, which allows to enable kernel debugging:

```

bcdedit /set {globalsettings} advancedoptions true

```

Connecting to KD interface on r2

Serial Port

Radare2 will use the `winkd` io plugin to connect to a socket file created by virtualbox or qemu. Also, the `winkd` debugger plugin and we should specify the x86-32 too. (32 and 64 bit debugging is supported)

```

$ r2 -a x86 -b 32 -D winkd winkd:///tmp/winkd.pipe

```

On Windows you should run the following line:

```

$ radare2 -D winkd winkd://\\.\pipe\com_1

```

Network

```
$ r2 -a x86 -b 32 -d winkd://<hostip>:<port>:w.x.y.z
```

Using KD

When connecting to a KD interface, r2 will send a breakin packet to interrupt the target and we will get stuck here:

```
[0x828997b8]> pd 20
;-- eip:
0x828997b8    cc          int3
0x828997b9    c20400         ret 4
0x828997bc    cc          int3
0x828997bd    90           nop
0x828997be    c3           ret
0x828997bf    90           nop
```

In order to skip that trap we will need to change eip and run 'dc' twice:

```
dr eip=eip+1
dc
dr eip=eip+1
dc
```

Now the Windows VM will be interactive again. We will need to kill r2 and attach again to get back to control the kernel.

In addition, the **dp** command can be used to list all processes, and **dpa** or **dp=** to attach to the process. This will display the base address of the process in the physical memory layout.

WinDBG Backend for Windows (DbgEng)

On Windows, radare2 can use **DbgEng.dll** as a debugging backend, allowing it to make use of WinDBG's capabilities, supporting dump files, local and remote user and kernel mode debugging.

You can use the debugging DLLs included on Windows or get the latest version from Microsoft's download page (recommended).

You cannot use DLLs from the Microsoft Store's **WinDbg Preview** app folder directly as they are not marked as executable for normal users.

radare2 will try to load **dbgeng.dll** from the **_NT_DEBUGGER_EXTENSION_PATH** environment variable before using Windows' default library search path.

Using the plugin

To use the `windbg` plugin, pass the same command-line options as you would for WinDBG or `kd` (see Microsoft's documentation), quoting/escaping when necessary:

```
> r2 -d "windbg://-remote tcp:server=Server,port=Socket"
> r2 -d "windbg://MyProgram.exe \"my arg\""
> r2 -d "windbg://-k net:port=<n>,key=<MyKey>"
> r2 -d "windbg://-z MyDumpFile.dmp"
```

You can then debug normally (see `d?` command) or interact with the backend shell directly with the `=!` command:

```
[0x7ffcac9fcea0]> dcu 0x0007ffc98f42190
Continue until 0x7ffc98f42190 using 1 bpsize
ModLoad: 00007ffc`ab6b0000 00007ffc`ab6e0000 C:\WINDOWS\System32\IMM32.DLL
Breakpoint 1 hit
hit breakpoint at: 0x7ffc98f42190

[0x7fffcf232190]> =!k4
Child-SP      RetAddr      Call Site
00000033`73b1f618 00007ff6`c67a861d r_main!r_main_radare2
00000033`73b1f620 00007ff6`c67d0019 radare2!main+0x8d
00000033`73b1f720 00007ff6`c67cfebe radare2!invoke_main+0x39
00000033`73b1f770 00007ff6`c67cfd7e radare2!__scrt_common_main_seh+0x12e
```

Windows Messages

On Windows, you can use `dbW` while debugging to set a breakpoint for the message handler of a specific window.

Get a list of the current process windows with `dW` :

```
[0x7ffe885c1164]> dW
-----
| Handle      | PID   | TID   | Class Name      |
)-----(
| 0x0023038e  | 9432  | 22432 | MSCTFIME UI     |
| 0x0029049e  | 9432  | 22432 | IME              |
| 0x002c048a  | 9432  | 22432 | Edit             |
| 0x000d0474  | 9432  | 22432 | msctls_statusbar32 |
| 0x00070bd6  | 9432  | 22432 | Notepad          |
`-----`
```

Set the breakpoint with a message type, together with either the window class name or its handle:

```
[0x7ffe885c1164]> dbW WM_KEYDOWN Edit
Breakpoint set.
```

Or

```
[0x7ffe885c1164]> dbW WM_KEYDOWN 0x002c048a
Breakpoint set.
```

If you aren't sure which window you should put a breakpoint on, use `dWi` to identify it with your mouse:

```
[0x7ffe885c1164]> dWi
Move cursor to the window to be identified. Ready? y
Try to get the child? y
```

```
-----
| Handle      | PID   | TID   | Class Name |
------(
| 0x002c048a  | 9432  | 22432 | Edit       |
-----'
```

Adding Metadata to Disassembly

The typical work involved in reversing binary files makes powerful annotation capabilities essential. Radare offers multiple ways to store and retrieve such metadata.

By following common basic UNIX principles, it is easy to write a small utility in a scripting language which uses `objdump`, `otool` or any other existing utility to obtain information from a binary and to import it into radare. For example, take a look at `idc2r.py` shipped with `radare2ida`. To use it, invoke it as `idc2r.py file.idc > file.r2`. It reads an IDC file exported from an IDA Pro database and produces an `r2` script containing the same comments, names of functions and other data. You can import the resulting 'file.r2' by using the dot `.` command of radare:

```
[0x00000000]> . file.r2
```

The `.` command is used to interpret Radare commands from external sources, including files and program output. For example, to omit generation of an intermediate file and import the script directly you can use this combination:

```
[0x00000000]> .!idc2r.py < file.idc
```

Please keep in mind that importing IDA Pro metadata from IDC dump is deprecated mechanism and might not work in the future. The recommended way to do it - use python-idb-based `ida2r2.py` which opens IDB files directly without IDA Pro installed.

The `C` command is used to manage comments and data conversions. You can define a range of program's bytes to be interpreted as either code, binary data or string. It is also possible to execute external code at every specified flag location

in order to fetch some metadata, such as a comment, from an external file or database.

There are many different metadata manipulation commands, here is the glimpse of all of them:

```
[0x00404cc0]> C?
| Usage: C[-LCvsdfm*?][*?] [...] # Metadata management
| C                                list meta info in human friendly form
| C*                              list meta info in r2 commands
| C*.                             list meta info of current offset in r2 commands
| C- [len] [[@]addr]             delete metadata at given address range
| C.                              list meta info of current offset in human friendly form
| CC! [@addr]                    edit comment with $EDITOR
| CC[?] [-] [comment-text] [@addr] add/remove comment
| CC. [addr]                     show comment in current address
| CCa[-at] | [at] [text] [@addr] add/remove comment at given address
| CCu [comment-text] [@addr]    add unique comment
| CF[sz] [fcn-sign..] [@addr]   function signature
| CL[-] [*] [file:line] [addr]  show or add 'code line' information (bininfo)
| CS[-] [space]                 manage meta-spaces to filter comments, etc.
| C[Cthsdmf]                    list comments/types/hidden/strings/data/magic
| C[Cthsdmf]*                   list comments/types/hidden/strings/data/magic
| Cd[-] [size] [repeat] [@addr] hexdump data array (Cd 4 10 == dword [10])
| Cd. [@addr]                   show size of data at current address
| Cf[?] [-] [sz] [0|cnt] [fmt] [a0 a1...] [@addr] format memory (see pf?)
| Ch[-] [size] [@addr]         hide data
| Cm[-] [sz] [fmt..] [@addr]   magic parse (see pm?)
| Cs[?] [-] [size] [@addr]     add string
| Ct[?] [-] [comment-text] [@addr] add/remove type analysis comment
| Ct. [@addr]                  show comment at current or specified address
| Cv[bsr] [?]                  add comments to args
| Cz[@addr]                    add string (see Cs?)
```

Simply to add the comment to a particular line/address you can use Ca command:

```
[0x00000000]> CCa 0x00000002 this guy seems legit
[0x00000000]> pd 2
0x00000000      0000      add [rax], al
;      this guy seems legit
0x00000002      0000      add [rax], al
```

The C? family of commands lets you mark a range as one of several kinds of types. Three basic types are: code (disassembly is done using asm.arch), data (an array of data elements) or string. Use the Cs comand to define a string, use the Cd command for defining an array of data elements, and use the Cf command to define more complex data structures like structs.

Annotating data types is most easily done in visual mode, using the “d” key,

short for “data type change”. First, use the cursor to select a range of bytes (press `c` key to toggle cursor mode and use HJKL keys to expand selection), then press ‘`d`’ to get a menu of possible actions/types. For example, to mark the range as a string, use the ‘`s`’ option from the menu. You can achieve the same result from the shell using the `Cs` command:

```
[0x00000000]> f string_foo @ 0x800
[0x00000000]> Cs 10 @ string_foo
```

The `Cf` command is used to define a memory format string (the same syntax used by the `pf` command). Here’s an example:

```
[0x7fd9f13ae630]> Cf 16 2xi foo bar
[0x7fd9f13ae630]> pd
;-- rip:
0x7fd9f13ae630 format 2xi foo bar {
0x7fd9f13ae630 [0] {
    foo : 0x7fd9f13ae630 = 0xe8e78948
    bar : 0x7fd9f13ae634 = 14696
}
0x7fd9f13ae638 [1] {
    foo : 0x7fd9f13ae638 = 0x8bc48949
    bar : 0x7fd9f13ae63c = 571928325
}
} 16
0x7fd9f13ae633      e868390000      call 0x7fd9f13b1fa0
0x7fd9f13ae638      4989c4          mov r12, rax
```

The `[sz]` argument to `Cf` is used to define how many bytes the struct should take up in the disassembly, and is completely independent from the size of the data structure defined by the format string. This may seem confusing, but has several uses. For example, you may want to see the formatted structure displayed in the disassembly, but still have those locations be visible as offsets and with raw bytes. Sometimes, you find large structures, but only identified a few fields, or only interested in specific fields. Then, you can tell `r2` to display only those fields, using the format string and using ‘`skip`’ fields, and also have the disassembly continue after the entire structure, by giving it full size using the `sz` argument.

Using `Cf`, it’s easy to define complex structures with simple oneliners. See `pf?` for more information. Remember that all these `C` commands can also be accessed from the visual mode by pressing the `d` (data conversion) key. Note that unlike `t` commands `Cf` doesn’t change analysis results. It is only a visual boon.

Sometimes just adding a single line of comments is not enough, in this case `radare2` allows you to create a link for a particular text file. You can use it with `CC`, command or by pressing `,` key in the visual mode. This will open an `$EDITOR` to create a new file, or if filename does exist, just will create a link. It will be shown in the disassembly comments:

```
[0x00003af7 11% 290 /bin/ls]> pd $r @ main+55 # 0x3af7
|0x00003af7 call sym.imp.setlocale ;[1] ; ,(locale-help.txt) ; char *setlocale(int c
|0x00003afc lea rsi, str.usr_share_locale ; 0x179cc ; "/usr/share/locale"
|0x00003b03 lea rdi, [0x000179b2] ; "coreutils"
|0x00003b0a call sym.imp.bindtextdomain ;[2] ; char *bindtextdomain(char *domainname, cha
```

Note `,(locale-help.txt)` appeared in the comments, if we press `,` again in the visual mode, it will open the file. Using this mechanism we can create a long descriptions of some particular places in disassembly, link datasheets or related articles.

ESIL

ESIL stands for ‘Evaluable Strings Intermediate Language’. It aims to describe a Forth-like representation for every target CPU opcode semantics. ESIL representations can be evaluated (interpreted) in order to emulate individual instructions. Each command of an ESIL expression is separated by a comma. Its virtual machine can be described as this:

```
while ((word=haveCommand())) {
    if (word.isOperator()) {
        esilOperators[word](esil);
    } else {
        esil.push (word);
    }
    nextCommand();
}
```

As we can see ESIL uses a stack-based interpreter similar to what is commonly used for calculators. You have two categories of inputs: values and operators. A value simply gets pushed on the stack, an operator then pops values (its arguments if you will) off the stack, performs its operation and pushes its results (if any) back on. We can think of ESIL as a post-fix notation of the operations we want to do.

So let’s see an example:

```
4,esp,-=,ebp,esp,[4]
```

Can you guess what this is? If we take this post-fix notation and transform it back to in-fix we get

```
esp -= 4
4bytes(dword) [esp] = ebp
```

We can see that this corresponds to the x86 instruction `push ebp`! Isn’t that cool? The aim is to be able to express most of the common operations performed by CPUs, like binary arithmetic operations, memory loads and stores, processing syscalls. This way if we can transform the instructions to ESIL we can see what

a program does while it is running even for the most cryptic architectures you definitely don't have a device to debug on for.

Using ESIL

r2's visual mode is great to inspect the ESIL evaluations.

There are 3 environment variables that are important for watching what a program does:

```
[0x00000000]> e emu.str = true
```

`asm.emu` tells r2 if you want ESIL information to be displayed. If it is set to true, you will see comments appear to the right of your disassembly that tell you how the contents of registers and memory addresses are changed by the current instruction. For example, if you have an instruction that subtracts a value from a register it tells you what the value was before and what it becomes after. This is super useful so you don't have to sit there yourself and track which value goes where.

One problem with this is that it is a lot of information to take in at once and sometimes you simply don't need it. r2 has a nice compromise for this. That is what the `emu.str` variable is for (`asm.emustr` on ≤ 2.2). Instead of this super verbose output with every register value, this only adds really useful information to the output, e.g., strings that are found at addresses a program uses or whether a jump is likely to be taken or not.

The third important variable is `asm.esil`. This switches your disassembly to no longer show you the actual disassembled instructions, but instead now shows you corresponding ESIL expressions that describe what the instruction does. So if you want to take a look at how instructions are expressed in ESIL simply set "asm.esil" to true.

```
[0x00000000]> e asm.esil = true
```

In visual mode you can also toggle this by simply typing 0.

ESIL Commands

- "ae" : Evaluate ESIL expression.

```
[0x00000000]> "ae 1,1,+"  
0x2  
[0x00000000]>
```

- "aes" : ESIL Step.

```
[0x00000000]> aes  
[0x00000000]>10aes
```

- "aeso" : ESIL Step Over.


```
[0x00000000]> aeso
[0x00000000]>10aeso
```

- “aeso” : ESIL Step Until.

```
[0x00001000]> aesu 0x1035
ADDR BREAK
[0x00001019]>
```

- “ar” : Show/modify ESIL registry.

```
[0x00001ec7]> ar r_00 = 0x1035
[0x00001ec7]> ar r_00
0x00001035
[0x00001019]>
```

ESIL Instruction Set

Here is the complete instruction set used by the ESIL VM:

ESIL Opcode	Operands	Name	Operation	example
TRAP	src	Trap	Trap signal	

```
** |src|Interrupt|interrupt|0x80, ()** | src | Syscall | syscall | rax,() $$ | src |
Instruction address | Get address of current instructionstack=instruction address
| == | src,dst | Compare | stack = (dst == src) ; update_eflags(dst - src) | < |
src,dst | Smaller (signed comparison) | stack = (dst < src) ; update_eflags(dst -
src) | [0x00000000]> “ae 1,5,<” 0x0> “ae 5,5”0x0” <= | src,dst | Smaller or
Equal (signed comparison) | stack = (dst <= src) ; update_eflags(dst - src)
| [0x00000000]> “ae 1,5,<” 0x0> “ae 5,5”0x1” > | src,dst | Bigger (signed
comparison) | stack = (dst > src) ; update_eflags(dst - src) | > “ae 1,5,>”0x1>
“ae 5,5,>”0x0 >= | src,dst | Bigger or Equal (signed comparison) | stack = (dst
>= src) ; update_eflags(dst - src) | > “ae 1,5,>=”0x1> “ae 5,5,>=”0x1 « |
src,dst | Shift Left | stack = dst « src | > “ae 1,1,«”0x2> “ae 2,1,«”0x4 » | src,dst
| Shift Right | stack = dst » src | > “ae 1,4,»”0x2> “ae 2,4,»”0x1 « < | src,dst |
Rotate Left | stack=dst ROL src | > “ae 31,1,«<”0x80000000> “ae 32,1,«<”0x1
»> | src,dst | Rotate Right | stack=dst ROR src | > “ae 1,1,»>”0x80000000>
“ae 32,1,»>”0x1 & | src,dst | AND | stack = dst & src | > “ae 1,1,&”0x1> “ae
1,0,&”0x0> “ae 0,1,&”0x0> “ae 0,0,&”0x0 | | src,dst | OR | stack = dst | src | >
“ae 1,1,|”0x1> “ae 1,0,|”0x1> “ae 0,1,|”0x1> “ae 0,0,|”0x0 ^ | src,dst | XOR |
stack = dst ^src | > “ae 1,1,^”0x0> “ae 1,0,^”0x1> “ae 0,1,^”0x1> “ae 0,0,^”0x0
+ | src,dst | ADD | stack = dst + src | > “ae 3,4,+”0x7> “ae 5,5,+”0xa
- | src,dst | SUB | stack = dst - src | > “ae 3,4,-”0x1> “ae 5,5,-”0x0> “ae
4,3,-”0xffffffff * | src,dst | MUL | stack = dst * src | > “ae 3,4,*”0xc> “ae
5,5,*”0x19 / | src,dst | DIV | stack = dst / src | > “ae 2,4,/”0x2> “ae 5,5,/”0x1>
“ae 5,9,/”0x1 % | src,dst | MOD | stack = dst % src | > “ae 2,4,%”0x0> “ae
5,5,%”0x0> “ae 5,9,%”0x4 ~ | bits,src | SIGNEXT | stack = src sign extended |
```

```

> "ae 8,0x80,~"0xfffffffffff80 ~/ | src,dst | SIGNED DIV | stack = dst / src
(signed) | > "ae 2,-4,~/ "0xfffffffffff80 ~% | src,dst | SIGNED MOD | stack =
dst % src (signed) | > "ae 2,-5,~%"0xfffffffffff80 ! | src | NEG | stack = !!src |
> "ae 1,! "0x0> "ae 4,! "0x0> "ae 0,! "0x1 ++ | src | INC | stack = src++ | >
ar r_00=0;ar r_000x00000000> "ae r_00,++ "0x1> ar r_000x00000000> "ae
1,++ "0x2 - | src | DEC | stack = src- | > ar r_00=5;ar r_000x00000005> "ae
r_00,- "0x4> ar r_000x00000005> "ae 5,- "0x4 = | src,reg | EQU | reg = src | >
"ae 3,r_00,="> aer r_000x00000003> "ae r_00,r_01,="> aer r_010x00000003
:= | src,reg | weak EQU | reg = src without side effects | > "ae 3,r_00,;=">
aer r_000x00000003> "ae r_00,r_01,;="> aer r_010x00000003 += | src,reg |
ADD eq | reg = reg + src | > ar r_01=5;ar r_00=0;ar r_000x00000000> "ae
r_01,r_00,+="> ar r_000x00000005> "ae 5,r_00,+="> ar r_000x0000000a -=
| src,reg | SUB eq | reg = reg - src | > "ae r_01,r_00,;="> ar r_000x00000004>
"ae 3,r_00,;="> ar r_000x00000001 *= | src,reg | MUL eq | reg = reg *
src | > ar r_01=3;ar r_00=5;ar r_000x00000005> "ae r_01,r_00,*="> ar
r_000x0000000f> "ae 2,r_00,*="> ar r_000x00000001e /= | src,reg | DIV
eq | reg = reg / src | > ar r_01=3;ar r_00=6;ar r_000x00000006> "ae
r_01,r_00,/="> ar r_000x00000002> "ae 1,r_00,/="> ar r_000x00000002
%= | src,reg | MOD eq | reg = reg % src | > ar r_01=3;ar r_00=7;ar r_00
0x00000007 > "ae r_01,r_00,%=" > ar r_00 0x00000001 > ar r_00=9;ar
r_00 0x00000009 > "ae 5,r_00,%=" > ar r_00 0x00000004 «= | src,reg | Shift
Left eq | reg = reg « src | > ar r_00=1;ar r_01=1;ar r_010x00000001> "ae
r_00,r_01,«="> ar r_010x00000002> "ae 2,r_01,«="> ar r_010x00000008
»= | src,reg | Shift Right eq | reg = reg « src | > ar r_00=1;ar r_01=8;ar
r_010x00000008> "ae r_00,r_01,»="> ar r_010x00000004> "ae 2,r_01,»=">
ar r_010x00000001 &= | src,reg | AND eq | reg = reg & src | > ar r_00=2;ar
r_01=6;ar r_010x00000006> "ae r_00,r_01,&="> ar r_010x00000002> "ae
2,r_01,&="> ar r_010x00000002> "ae 1,r_01,&="> ar r_010x00000000 |=
| src,reg | OR eq | reg = reg | src | > ar r_00=2;ar r_01=1;ar r_010x00000001>
"ae r_00,r_01,|= "> ar r_010x00000003> "ae 4,r_01,|= "> ar r_010x00000007
^= | src,reg | XOR eq | reg = reg ^ src | > ar r_00=2;ar r_01=0xab;ar
r_010x0000000ab> "ae r_00,r_01,^="> ar r_010x0000000a9> "ae 2,r_01,^=">
ar r_010x0000000ab ++= | reg | INC eq | reg = reg + 1 | > ar r_00=4;ar
r_000x00000004> "ae r_00,++="> ar r_000x00000005 -= | reg | DEC eq | reg
= reg - 1 | > ar r_00=4;ar r_000x00000004> "ae r_00,-="> ar r_000x00000003
!= | reg | NOT eq | reg = !reg | > ar r_00=4;ar r_000x00000004> "ae
r_00,!="> ar r_000x00000000> "ae r_00,!="> ar r_000x00000001 — | — |
— | — | ————— =[]=[*]=[1]=[2]=[4]=[8] | src,dst |
poke [*dst=src | > "ae 0xdeadbeef,0x10000,[4],"> pxw 4@0x100000x00010000
0xdeadbeef ...> "ae 0x0,0x10000,[4],"> pxw 4@0x100000x00010000
0x00000000 [][*][1][2][4][8] | src | peek | stack=*src | > w test@0x10000>
"ae 0x10000,[4], "0x74736574> ar r_00=0x10000> "ae r_00,[4], "0x74736574
|=[]|=1|=2|=4|=8 | reg | nombre | code | > > SWAP | | Swap | Swap two
top elements | SWAP DUP | | Duplicate | Duplicate top element in stack |
DUP NUM | | Numeric | If top element is a reference (register name, label, etc),
dereference it and push its real value | NUM CLEAR | | Clear | Clear stack |

```

CLEAR BREAK | | Break | Stops ESIL emulation | BREAK GOTO | n | Goto |
 Jumps to Nth ESIL word | GOTO 5 TODO | | To Do | Stops execution (reason:
 ESIL expression not completed) | TODO

ESIL Flags

ESIL VM provides by default a set of helper operations for calculating flags. They fulfill their purpose by comparing the old and the new value of the dst operand of the last performed eq-operation. On every eq-operation (e.g. =) ESIL saves the old and new value of the dst operand. Note, that there also exist weak eq operations (e.g. :=), which do not affect flag operations. The == operation affects flag operations, despite not being an eq operation. Flag operations are prefixed with \$ character.

```
z      - zero flag, only set if the result of an operation is 0
b      - borrow, this requires to specify from which bit (example: 4,$b - checks if borrow from bit 4)
c      - carry, same like above (example: 7,$c - checks if carry from bit 7)
o      - overflow
p      - parity
r      - regsize ( asm.bits/8 )
s      - sign
ds     - delay slot state
jt     - jump target
js     - jump target set
```

Syntax and Commands

A target opcode is translated into a comma separated list of ESIL expressions.

```
xor eax, eax    ->    0,eax,=,1,zf,=
```

Memory access is defined by brackets operation:

```
mov eax, [0x80480]    ->    0x80480,[],eax,=
```

Default operand size is determined by size of operation destination.

```
movb $0, 0x80480      ->    0,0x80480,=[1]
```

The ? operator uses the value of its argument to decide whether to evaluate the expression in curly braces.

1. Is the value zero? -> Skip it.
2. Is the value non-zero? -> Evaluate it.

```
cmp eax, 123    ->    123,eax,==,$z,zf,=
jz eax          ->    zf,{eax,eip,=,}
```

If you want to run several expressions under a conditional, put them in curly braces:

```
zf,{eip,esp,=[],eax,eip,=,$r,esp,-=,}
```

Whitespaces, newlines and other chars are ignored. So the first thing when processing a ESIL program is to remove spaces:

```
esil = r_str_replace (esil, " ", "", R_TRUE);
```

Syscalls need special treatment. They are indicated by '\$' at the beginning of an expression. You can pass an optional numeric value to specify a number of syscall. An ESIL emulator must handle syscalls. See (r_esil_syscall).

Arguments Order for Non-associative Operations

As discussed on IRC, the current implementation works like this:

```
a,b,-      b - a
a,b,/=     b /= a
```

This approach is more readable, but it is less stack-friendly.

Special Instructions

NOPs are represented as empty strings. As it was said previously, interrupts are marked by '*command.Forexample,'0x80,'*. It delegates emulation from the ESIL machine to a callback which implements interrupt handler for a specific OS/kernel/platform.

Traps are implemented with the TRAP command. They are used to throw exceptions for invalid instructions, division by zero, memory read error, or any other needed by specific architectures.

Quick Analysis

Here is a list of some quick checks to retrieve information from an ESIL string. Relevant information will be probably found in the first expression of the list.

```
indexOf('[')    -> have memory references
indexOf("=[")   -> write in memory
indexOf("pc,=") -> modifies program counter (branch, jump, call)
indexOf("sp,=") -> modifies the stack (what if we found sp+= or sp-=?)
indexOf("=")    -> retrieve src and dst
indexOf(":")    -> unknown esil, raw opcode ahead
indexOf("$")    -> accesses internal esil vm flags ex: $z
indexOf("$")    -> syscall ex: 1,$
indexOf("TRAP") -> can trap
indexOf('++')   -> has iterator
indexOf('--')   -> count to zero
indexOf("?{")   -> conditional
equalsTo("")    -> empty string, aka nop (wrong, if we append pc+=x)
```

Common operations: * Check dstreg * Check srcreg * Get destinaion * Is jump
* Is conditional * Evaluate * Is syscall

CPU Flags

CPU flags are usually defined as single bit registers in the RReg profile. They are sometimes found under the ‘flg’ register type.

Variables

Properties of the VM variables:

1. They have no predefined bit width. This way it should be easy to extend them to 128, 256 and 512 bits later, e.g. for MMX, SSE, AVX, Neon SIMD.
2. There can be unbound number of variables. It is done for SSA-form compatibility.
3. Register names have no specific syntax. They are just strings.
4. Numbers can be specified in any base supported by RNum (dec, hex, oct, binary ...).
5. Each ESIL backend should have an associated RReg profile to describe the ESIL register specs.

Bit Arrays

What to do with them? What about bit arithmetics if use variables instead of registers?

Arithmetics

1. ADD (“+”)
2. MUL (“*”)
3. SUB (“-”)
4. DIV (“/”)
5. MOD (“%”)

Bit Arithmetics

1. AND “&”
2. OR “|”
3. XOR “^”
4. SHL “«”
5. SHR “»”
6. ROL “«<”
7. ROR “»>”
8. NEG “!”

Floating Point Unit Support

At the moment of this writing, ESIL does not yet support FPU. But you can implement support for unsupported instructions using r2pipe. Eventually we will get proper support for multimedia and floating point.

Handling x86 REP Prefix in ESIL

ESIL specifies that the parsing control-flow commands must be uppercase. Bear in mind that some architectures have uppercase register names. The corresponding register profile should take care not to reuse any of the following:

```
3,SKIP    - skip N instructions. used to make relative forward GOTOs
3,GOTO    - goto instruction 3
LOOP      - alias for 0,GOTO
BREAK     - stop evaluating the expression
STACK     - dump stack contents to screen
CLEAR     - clear stack
```

Usage Example: `rep cmpsb`

```
cx,!,{BREAK},esi,[1],edi,[1],==,{BREAK},esi,++,edi,++,cx,--,0,GOTO
```

Unimplemented/Unhandled Instructions

Those are expressed with the ‘TODO’ command. They act as a ‘BREAK’, but displays a warning message describing that an instruction is not implemented and will not be emulated. For example:

```
fmulp ST(1), ST(0)      =>      TODO,fmulp ST(1),ST(0)
```

ESIL Disassembly Example:

```
[0x1000010f8]> e asm.esil=true
[0x1000010f8]> pd $r @ entry0
0x1000010f8    55                8,rsp,--,rbp,rsp,=[8]
0x1000010f9    4889e5            rsp,rbp,=
0x1000010fc    4883c768          104,rdi,+=
0x100001100    4883c668          104,rsi,+=
0x100001104    5d                rsp,[8],rbp,=,8,rsp,+=
0x100001105    e950350000       0x465a,rip,= ;[1]
0x10000110a    55                8,rsp,--,rbp,rsp,=[8]
0x10000110b    4889e5            rsp,rbp,=
0x10000110e    488d4668          rsi,104,+,rax,=
0x100001112    488d7768          rdi,104,+,rsi,=
0x100001116    4889c7            rax,rdi,=
0x100001119    5d                rsp,[8],rbp,=,8,rsp,+=
0x10000111a    e93b350000       0x465a,rip,= ;[1]
```

0x10000111f	55	8, rsp, -=, rbp, rsp, = [8]
0x100001120	4889e5	rsp, rbp, =
0x100001123	488b4f60	rdi, 96, +, [8], rcx, =
0x100001127	4c8b4130	rcx, 48, +, [8], r8, =
0x10000112b	488b5660	rsi, 96, +, [8], rdx, =
0x10000112f	b801000000	1, eax, =
0x100001134	4c394230	rdx, 48, +, [8], r8, ==, cz, ?=
0x100001138	7f1a	sf, of, !, ^, zf, !, &, ?{, 0x1154, rip, =, } ; [2]
0x10000113a	7d07	of, !, sf, ^, ?{, 0x1143, rip, } ; [3]
0x10000113c	b8ffffff	0xffffffff, eax, = ; 0xffffffff
0x100001141	eb11	0x1154, rip, = ; [2]
0x100001143	488b4938	rcx, 56, +, [8], rcx, =
0x100001147	48394a38	rdx, 56, +, [8], rcx, ==, cz, ?=

Introspection

To ease ESIL parsing we should have a way to express introspection expressions to extract the data that we want. For example, we may want to get the target address of a jump. The parser for ESIL expressions should offer an API to make it possible to extract information by analyzing the expressions easily.

```
> ao~esil, opcode
opcode: jmp 0x10000465a
esil: 0x10000465a, rip, =
```

We need a way to retrieve the numeric value of 'rip'. This is a very simple example, but there are more complex, like conditional ones. We need expressions to be able to get:

- opcode type
- destination of a jump
- condition depends on
- all regs modified (write)
- all regs accessed (read)

API HOOKS

It is important for emulation to be able to setup hooks in the parser, so we can extend it to implement analysis without having to change it again and again. That is, every time an operation is about to be executed, a user hook is called. It can be used for example to determine if RIP is going to change, or if the instruction updates the stack. Later, we can split that callback into several ones to have an event-based analysis API that may be extended in JavaScript like this:

```
esil.on('regset', function(){..
esil.on('syscall', function(){esil.regset('rip'
```

For the API, see the functions `hook_flag_read()`, `hook_execute()` and `hook_mem_read()`. A callback should return true or 1 if you want to override the action that it takes. For example, to deny memory reads in a region, or voiding memory writes, effectively making it read-only. Return false or 0 if you want to trace ESIL expression parsing.

Other operations require bindings to external functionalities to work. In this case, `r_ref` and `r_io`. This must be defined when initializing the ESIL VM.

- Io Get/Set

```
Out ax, 44
44,ax,:ou
```

- Selectors (cs,ds,gs...)

```
Mov eax, ds:[ebp+8]
Ebp,8,+, :ds,eax,=
```

Disassembling

Disassembling in radare is just a way to represent an array of bytes. It is handled as a special print mode within `p` command.

In the old times, when the radare core was smaller, the disassembler was handled by an external `rsc` file. That is, radare first dumped current block into a file, and then simply called `objdump` configured to disassemble for Intel, ARM or other supported architectures.

It was a working and unix friendly solution, but it was inefficient as it repeated the same expensive actions over and over, because there were no caches. As a result, scrolling was terribly slow.

So there was a need to create a generic disassembler library to support multiple plugins for different architectures. We can list the current loaded plugins with

```
$ rasm2 -L
```

Or from inside radare2:

```
> e asm.arch=??
```

This was many years before capstone appeared. So r2 was using `udis86` and `olly` disassemblers, many `gnu` (from `binutils`).

Nowadays, the disassembler support is one of the basic features of radare. It now has many options, endianness, including target architecture flavor and disassembler variants, among other things.

To see the disassembly, use the `pd` command. It accepts a numeric argument to specify how many opcodes of current block you want to see. Most of the commands in radare consider the current block size as the default limit for data

input. If you want to disassemble more bytes, set a new block size using the `b` command.

```
[0x00000000]> b 100      ; set block size to 100
[0x00000000]> pd         ; disassemble 100 bytes
[0x00000000]> pd 3       ; disassemble 3 opcodes
[0x00000000]> pD 30      ; disassemble 30 bytes
```

The `pD` command works like `pd` but accepts the number of input bytes as its argument, instead of the number of opcodes.

The “pseudo” syntax may be somewhat easier for a human to understand than the default assembler notations. But it can become annoying if you read lots of code. To play with it:

```
[0x00405e1c]> e asm.pseudo = true
[0x00405e1c]> pd 3
          ; JMP XREF from 0x00405dfa (fcn.00404531)
0x00405e1c  488b9424a80. rdx = [rsp+0x2a8]
0x00405e24  64483314252. rdx ^= [fs:0x28]
0x00405e2d  4889d8      rax = rbx

[0x00405e1c]> e asm.syntax = intel
[0x00405e1c]> pd 3
          ; JMP XREF from 0x00405dfa (fcn.00404531)
0x00405e1c  488b9424a80. mov rdx, [rsp+0x2a8]
0x00405e24  64483314252. xor rdx, [fs:0x28]
0x00405e2d  4889d8      mov rax, rbx

[0x00405e1c]> e asm.syntax=att
[0x00405e1c]> pd 3
          ; JMP XREF from 0x00405dfa (fcn.00404531)
0x00405e1c  488b9424a80. mov 0x2a8(%rsp), %rdx
0x00405e24  64483314252. xor %fs:0x28, %rdx
0x00405e2d  4889d8      mov %rbx, %rax
```

Basic Debugger Session

To debug a program, start `radare` with the `-d` option. Note that you can attach to a running process by specifying its PID, or you can start a new program by specifying its name and parameters:

```
$ pidof mc
32220
$ r2 -d 32220
$ r2 -d /bin/ls
$ r2 -a arm -b 16 -d gdb://192.168.1.43:9090
```

...

In the second case, the debugger will fork and load the debuggee `ls` program in memory.

It will pause its execution early in `ld.so` dynamic linker. As a result, you will not yet see the entrypoint or any shared libraries at this point.

You can override this behavior by setting another name for an entry breakpoint. To do this, add a radare command `e dbg.bep=entry` or `e dbg.bep=main` to your startup script, usually it is `~/.config/radare2/radare2rc`.

Another way to continue until a specific address is by using the `dcu` command. Which means: “debug continue until” taking the address of the place to stop at. For example:

```
dcu main
```

Be warned that certain malware or other tricky programs can actually execute code before `main()` and thus you’ll be unable to control them. (Like the program constructor or the `tls` initializers)

Below is a list of most common commands used with debugger:

```
> d?          ; get help on debugger commands
> ds 3        ; step 3 times
> db 0x8048920 ; setup a breakpoint
> db -0x8048920 ; remove a breakpoint
> dc          ; continue process execution
> dcs         ; continue until syscall
> dd          ; manipulate file descriptors
> dm          ; show process maps
> dmp A S rwx  ; change permissions of page at A and size S
> dr eax=33    ; set register value. eax = 33
```

There is another option for debugging in radare, which may be easier: using visual mode.

That way you will neither need to remember many commands nor to keep program state in your mind.

To enter visual debugger mode use `Vpp`:

```
[0xb7f0c8c0]> Vpp
```

The initial view after entering visual mode is a hexdump view of the current target program counter (e.g., EIP for x86). Pressing `p` will allow you to cycle through the rest of visual mode views. You can press `p` and `P` to rotate through the most commonly used print modes. Use `F7` or `s` to step into and `F8` or `S` to step over current instruction. With the `c` key you can toggle the cursor mode to mark a byte range selection (for example, to later overwrite them with `nop`). You can set breakpoints with `F2` key.

In visual mode you can enter regular radare commands by prepending them with `:.` . For example, to dump a one block of memory contents at ESI:

```
<Press ':'>
x @ esi
```

To get help on visual mode, press `?`. To scroll the help screen, use arrows. To exit the help view, press `q`.

A frequently used command is `dr`, which is used to read or write values of the target's general purpose registers. For a more compact register value representation you might use `dr=` command. You can also manipulate the hardware and the extended/floating point registers.

Command Format

A general format for radare2 commands is as follows:

```
[.][times][cmd][~grep][@[@iter]addr!size][|>pipe] ;
```

People who use Vim daily and are familiar with its commands will find themselves at home. You will see this format used throughout the book. Commands are identified by a single case-sensitive character `[a-zA-Z]`.

To repeatedly execute a command, prefix the command with a number:

```
px      # run px
3px     # run px 3 times
```

The `!` prefix is used to execute a command in shell context. If you want to use the `cmd` callback from the I/O plugin you must prefix with `=!`.

Note that a single exclamation mark will run the command and print the output through the RCons API. This means that the execution will be blocking and not interactive. Use double exclamation marks – `!!` – to run a standard system call.

All the socket, filesystem and execution APIs can be restricted with the `cfg.sandbox` configuration variable.

A few examples:

```
ds                ; call the debugger's 'step' command
px 200 @ esp      ; show 200 hex bytes at esp
pc > file.c       ; dump buffer as a C byte array to file.c
wx 90 @@ sym.*    ; write a nop on every symbol
pd 2000 | grep eax ; grep opcodes that use the 'eax' register
px 20 ; pd 3 ; px 40 ; multiple commands in a single line
```

The standard UNIX pipe `|` is also available in the radare2 shell. You can use it to filter the output of an `r2` command with any shell program that reads from stdin, such as `grep`, `less`, `wc`. If you do not want to spawn anything, or

you can't, or the target system does not have the basic UNIX tools you need (Windows or embedded users), you can also use the built-in grep (~).

See ~? for help.

The ~ character enables internal grep-like function used to filter output of any command:

```
pd 20~call          ; disassemble 20 instructions and grep output for 'call'
```

Additionally, you can grep either for columns or for rows:

```
pd 20~call:0        ; get first row
pd 20~call:1        ; get second row
pd 20~call[0]       ; get first column
pd 20~call[1]       ; get second column
```

Or even combine them:

```
pd 20~call:0[0]     ; grep the first column of the first row matching 'call'
```

This internal grep function is a key feature for scripting radare2, because it can be used to iterate over a list of offsets or data generated by disassembler, ranges, or any other command. Refer to the loops section (iterators) for more information.

The @ character is used to specify a temporary offset at which the command to its left will be executed. The original seek position in a file is then restored.

For example, `pd 5 @ 0x100000fce` to disassemble 5 instructions at address 0x100000fce.

Most of the commands offer autocompletion support using <TAB> key, for example `seek` or `flags` commands. It offers autocompletion using all possible values, taking flag names in this case. Note that it is possible to see the history of the commands using the `!~...` command - it offers a visual mode to scroll through the radare2 command history.

To extend the autocompletion support to handle more commands or enable autocompletion to your own commands defined in core, I/O plugins you must use the `!!!` command.

Command-line Options

The radare core accepts many flags from the command line.

This is an excerpt from the usage help message:

```
$ radare2 -h
Usage: r2 [-ACdfLMnQStuvwzX] [-P patch] [-p prj] [-a arch] [-b bits] [-i file]
          [-s addr] [-B baddr] [-m maddr] [-c cmd] [-e k=v] file|pid|--|=
--          run radare2 without opening any file
-          same as 'r2 malloc://512'
```

```

=          read file from stdin (use -i and -c to run cmds)
-=        perform !=! command to run all commands remotely
-0        print \x00 after init and every command
-2        close stderr file descriptor (silent warning messages)
-a [arch]  set asm.arch
-A        run 'aaa' command to analyze all referenced code
-b [bits]  set asm.bits
-B [baddr] set base address for PIE binaries
-c 'cmd..' execute radare command
-C        file is host:port (alias for -c+=http://%s/cmd/)
-d        debug the executable 'file' or running process 'pid'
-D [backend] enable debug mode (e cfg.debug=true)
-e k=v     evaluate config var
-f        block size = file size
-F [binplug] force to use that rbin plugin
-h, -hh    show help message, -hh for long
-H ([var]) display variable
-i [file]  run script file
-I [file]  run script file before the file is opened
-k [OS/kern] set asm.os (linux, macos, w32, netbsd, ...)
-l [lib]   load plugin file
-L        list supported IO plugins

-m [addr]  map file at given address (loadaddr)
-M        do not demangle symbol names
-n, -nn    do not load RBin info (-nn only load bin structures)
-N        do not load user settings and scripts
-q        quiet mode (no prompt) and quit after -i
-Q        quiet mode (no prompt) and quit faster (quickLeak=true)
-p [prj]   use project, list if no arg, load if no file
-P [file]  apply rapatch file and quit
-r [rarun2] specify rarun2 profile to load (same as -e dbg.profile=X)
-R [rr2rule] specify custom rarun2 directive
-s [addr]  initial seek
-S        start r2 in sandbox mode
-t        load rabin2 info in thread
-u        set bin.filter=false to get raw sym/sec/cls names
-v, -V    show radare2 version (-V show lib versions)
-w        open file in write mode
-x        open without exec-flag (asm.emu will not work), See io.exec
-X        same as -e bin.usextr=false (useful for dyldcache)
-z, -zz   do not load strings or load them even in raw

```

Common usage patterns

Open a file in write mode without parsing the file format headers.

```
$ r2 -nw file
```

Quickly get into an r2 shell without opening any file.

```
$ r2 -
```

Specify which sub-binary you want to select when opening a fatbin file:

```
$ r2 -a ppc -b 32 ls.fat
```

Run a script before showing interactive command-line prompt:

```
$ r2 -i patch.r2 target.bin
```

Execute a command and quit without entering the interactive mode:

```
$ r2 -qc ij hi.bin > imports.json
```

Set the configuration variable:

```
$ r2 -e scr.color=0 blah.bin
```

Debug a program:

```
$ r2 -d ls
```

Use an existing project file:

```
$ r2 -p test
```

Android

Radare2 can be cross-compiled for other architectures/systems as well, like Android.

Prerequisites

- Python 3
- Meson
- Ninja
- Git
- Android NDK

Step-by-step

Download and extract the Android NDK Download the Android NDK from the official site and extract it somewhere on your system (e.g. /tmp/android-ndk)

Make

Specify NDK base path

```
$ echo NDK=/tmp/android-ndk > ~/.r2androidrc
```

Compile + create tar.gz + push it to connected android device

```
./sys/android-build.sh arm64-static
```

You can build for different architectures by changing the argument to `./sys/android-build.sh`. Run the script without any argument to see the accepted values.

Meson

Create a cross-file for meson Meson needs a configuration file that describes the cross compilation environment (e.g. `meson-android.ini`). You can adjust it as necessary, but something like the following should be a good starting point:

```
[binaries]
c      = '/tmp/android-ndk/toolchains/llvm/prebuilt/linux-x86_64/bin/aarch64-linux-android28'
cpp    = '/tmp/android-ndk/toolchains/llvm/prebuilt/linux-x86_64/bin/aarch64-linux-android28'
ar      = '/tmp/android-ndk/toolchains/llvm/prebuilt/linux-x86_64/bin/aarch64-linux-android28'
as      = '/tmp/android-ndk/toolchains/llvm/prebuilt/linux-x86_64/bin/aarch64-linux-android28'
ranlib  = '/tmp/android-ndk/toolchains/llvm/prebuilt/linux-x86_64/bin/aarch64-linux-android28'
ld      = '/tmp/android-ndk/toolchains/llvm/prebuilt/linux-x86_64/bin/aarch64-linux-android28'
strip   = '/tmp/android-ndk/toolchains/llvm/prebuilt/linux-x86_64/bin/aarch64-linux-android28'
pkgconfig = 'false'

[properties]
sys_root = '/tmp/android-ndk/sysroot'

[host_machine]
system = 'android'
cpu_family = 'arm'
cpu = 'aarch64'
endian = 'little'
```

Compile with meson + ninja Now setup the build directory with meson as usual:

```
$ CFLAGS="-static" LDFLAGS="-static" meson --default-library static --prefix=/tmp/android-d
```

A bit of explanation about all the options: `* CFLAGS="-static", LDFLAGS="-static", --default-library static`: this ensure that libraries and binaries are statically compiled, so you do not need to properly set `LD_*` environment variables in your Android environment to make it find the right libraries. Binaries have everything they need inside. `* -Dblob=true`: it tells meson to compile just one binary with all the needed code for running `radare2`, `rabin2`, `rasm2`, etc. and creates symbolic links to those names. This avoids creating many statically compiled large binaries and just create one that provides all features. You will still have `rabin2`, `rasm2`, `rax2`, etc. but

they are just symlinks to **radare2**. * `--cross-file ./meson-android.ini`: it describes how to compile radare2 for Android

Then compile and install the project:

```
$ ninja -C build
$ ninja -C build install
```

Move files to your android device and enjoy At this point you can copy the generated files in `/tmp/android-dir` to your Android device and running radare2 from it. For example:

```
$ cd /tmp && tar -cvf radare2-android.tar.gz android-dir
$ adb push radare2-android.tar.gz /data/local/tmp
$ adb shell
DEVICE:/ $ cd /data/local/tmp
DEVICE:/data/local/tmp $ tar xvf radare2-android.tar.gz
DEVICE:/data/local/tmp $ ./android-dir/bin/radare2
Usage: r2 [-ACdfLMnQStuvwzX] [-P patch] [-p prj] [-a arch] [-b bits] [-i file]
        [-s addr] [-B baddr] [-m maddr] [-c cmd] [-e k=v] file|pid|--|=
```

Compilation and Portability

Currently the core of radare2 can be compiled on many systems and architectures, but the main development is done on GNU/Linux with GCC, and on MacOS X with clang. Radare is also known to compile on many different systems and architectures (including TCC and SunStudio).

People often want to use radare as a debugger for reverse engineering. Currently, the debugger layer can be used on Windows, GNU/Linux (Intel x86 and x86_64, MIPS, and ARM), OS X, FreeBSD, NetBSD, and OpenBSD (Intel x86 and x86_64)..

Compared to core, the debugger feature is more restrictive portability-wise. If the debugger has not been ported to your favorite platform, you can disable the debugger layer with the `--without-debugger` **configure** script option when compiling radare2.

Note that there are I/O plugins that use GDB, WinDbg, or Wine as back-ends, and therefore rely on presence of corresponding third-party tools (in case of remote debugging - just on the target machine).

To build on a system using **acr** and **GNU Make** (e.g. on *BSD systems):

```
$ ./configure --prefix=/usr
$ gmake
$ sudo gmake install
```

There is also a simple script to do this automatically:

```
$ sys/install.sh
```


Static Build

You can build radare2 statically along with all other tools with the command:

```
$ sys/static.sh
```

Meson build

You can use meson + ninja to build:

```
$ sys/meson.py --prefix=/usr --shared --install
```

If you want to build locally:

```
$ sys/meson.py --prefix=/home/$USER/r2meson --local --shared --install
```

Docker

Radare2 repository ships a Dockerfile that you can use with Docker.

This dockerfile is also used by Remnux distribution from SANS, and is available on the docker registryhub.

Cleaning Up Old Radare2 Installations

```
./configure --prefix=/old/r2/prefix/installation  
make purge
```

Contributing

Radare2 Book

If you want to contribute to the Radare2 book, you can do it at the Github repository. Suggested contributions include:

- Crackme writeups
- CTF writeups
- Documentation on how to use Radare2
- Documentation on developing for Radare2
- Conference presentations/workshops using Radare2
- Missing content from the Radare1 book updated to Radare2

Please get permission to port any content you do not own/did not create before you put it in the Radare2 book.

See <https://github.com/radareorg/radare2/blob/master/DEVELOPERS.md> for general help on contributing to radare2.

Expressions

Expressions are mathematical representations of 64-bit numerical values. They can be displayed in different formats, be compared or used with all commands


```
double: 0.000000
trits   0t10
```

Numbers can be displayed in several formats:

```
0x033   : hexadecimal can be displayed
3334    : decimal
sym.fo   : resolve flag offset
10K      : KBytes   10*1024
10M      : MBytes   10*1024*1024
```

You can also use variables and seek positions to build complex expressions.

Use the `?$?` command to list all the available commands or read the `refcard` chapter of this book.

```
$$      here (the current virtual seek)
$l      opcode length
$s      file size
$j      jump address (e.g. jmp 0x10, jz 0x10 => 0x10)
$f      jump fail address (e.g. jz 0x10 => next instruction)
$m      opcode memory reference (e.g. mov eax,[0x10] => 0x10)
$b      block size
```

Some more examples:

```
[0x4A13B8C0]> ? $m + $l
140293837812900 0x7f98b45df4a4 03771426427372244 130658.0G 8b45d000:04a4 140293837812900 103
[0x4A13B8C0]> pd 1 @ +$l
0x4A13B8C2      call 0x4a13c000
```

Downloading radare2

You can get radare from the GitHub repository: <https://github.com/radareorg/radare2>

Binary packages are available for a number of operating systems (Ubuntu, Maemo, Gentoo, Windows, iPhone, and so on). But you are highly encouraged to get the source and compile it yourself to better understand the dependencies, to make examples more accessible and, of course, to have the most recent version.

A new stable release is typically published every month.

The radare development repository is often more stable than the ‘stable’ releases. To obtain the latest version:

```
$ git clone https://github.com/radareorg/radare2.git
```

This will probably take a while, so take a coffee break and continue reading this book.

To update your local copy of the repository, use `git pull` anywhere in the radare2 source code tree:

```
$ git pull
```

If you have local modifications of the source, you can revert them (and lose them!) with:

```
$ git reset --hard HEAD
```

Or send us a patch:

```
$ git diff > radare-foo.patch
```

The most common way to get r2 updated and installed system wide is by using:

```
$ sys/install.sh
```

Building with meson + ninja

There is also a work-in-progress support for Meson.

Using clang and ld.gold makes the build faster:

```
CC=clang LDFLAGS=-fuse-ld=gold meson . release --buildtype=release --prefix ~/.local/stow/r2
ninja -C release
# ninja -C release install
```

Helper Scripts

Take a look at the scripts in `sys/`, they are used to automate stuff related to syncing, building and installing r2 and its bindings.

The most important one is `sys/install.sh`. It will pull, clean, build and symstall r2 system wide.

Symstalling is the process of installing all the programs, libraries, documentation and data files using symlinks instead of copying the files.

By default it will be installed in `/usr/local`, but you can specify a different prefix using the argument `--prefix`.

This is useful for developers, because it permits them to just run ‘make’ and try changes without having to run make install again.

Cleaning Up

Cleaning up the source tree is important to avoid problems like linking to old objects files or not updating objects after an ABI change.

The following commands may help you to get your git clone up to date:

```
$ git clean -xdf
$ git reset --hard origin/master
$ git pull
```

If you want to remove previous installations from your system, you must run the following commands:

```
$ ./configure --prefix=/usr/local
$ make purge
```

History

In 2006, Sergi Álvarez (aka pancake) was working as a forensic analyst. Since he wasn't allowed to use the company software for his personal needs, he decided to write a small tool-a hexadecimal editor-with very basic characteristics:

- be extremely portable (unix friendly, command line, c, small)
- open disk devices, this is using 64bit offsets
- search for a string or hexpair
- review and dump the results to disk

The editor was originally designed to recover a deleted file from an HFS+ partition.

After that, pancake decided to extend the tool to have a pluggable io to be able to attach to processes and implemented the debugger functionalities, support for multiple architectures, and code analysis.

Since then, the project has evolved to provide a complete framework for analyzing binaries, while making use of basic UNIX concepts. Those concepts include the famous “everything is a file”, “small programs that interact using stdin/stdout”, and “keep it simple” paradigms.

The need for scripting showed the fragility of the initial design: a monolithic tool made the API hard to use, and so a deep refactoring was needed. In 2009 radare2 (r2) was born as a fork of radare1. The refactor added flexibility and dynamic features. This enabled much better integration, paving the way to use r2 from different programming languages. Later on, the r2pipe API allowed access to radare2 via pipes from any language.

What started as a one-man project, with some eventual contributions, gradually evolved into a big community-based project around 2014. The number of users was growing fast, and the author-and main developer-had to switch roles from coder to manager in order to integrate the work of the different developers that were joining the project.

Instructing users to report their issues allows the project to define new directions to evolve in. Everything is managed in radare2's GitHub and discussed in the Telegram channel.

The project remains active at the time of writing this book, and there are several side projects that provide, among other things, a graphical user interface (Cutter), a decompiler (r2dec, radeco), Frida integration (r2frida), Yara, Unicorn, Keystone, and many other projects indexed in the r2pm (the radare2 package manager).

Since 2016, the community gathers once a year in r2con, a congress around radare2 that takes place in Barcelona.

Basic Radare2 Usage

The learning curve is usually somewhat steep at the beginning. Although after an hour of using it you should easily understand how most things work, and how to combine the various tools radare offers. You are encouraged to read the rest of this book to understand how some non-trivial things work, and to ultimately improve your skills.

learning_curve

Navigation, inspection and modification of a loaded binary file is performed using three simple actions: seek (to position), print (buffer), and alternate (write, append).

The ‘seek’ command is abbreviated as **s** and accepts an expression as its argument. The expression can be something like 10, +0x25, or [0x100+ptr_table]. If you are working with block-based files, you may prefer to set the block size to a required value with **b** command, and seek forward or backwards with positions aligned to it. Use **s++** and **s--** commands to navigate this way.

If radare2 opens an executable file, by default it will open the file in Virtual Addressing (VA) mode and the sections will be mapped to their virtual addresses. In VA mode, seeking is based on the virtual address and the starting position is set to the entry point of the executable. Using **-n** option you can suppress this default behavior and ask radare2 to open the file in non-VA mode for you. In non-VA mode, seeking is based on the offset from the beginning of the file.

The ‘print’ command is abbreviated as **p** and has a number of submodes — the second letter specifying a desired print mode. Frequent variants include **px** to print in hexadecimal, and **pd** for disassembling.

To be allowed to write files, specify the **-w** option to radare2 when opening a file. The **w** command can be used to write strings, hexpairs (**x** subcommand), or even assembly opcodes (**a** subcommand). Examples:

```
> w hello world          ; string
> wx 90 90 90 90         ; hexpairs
> wa jmp 0x8048140        ; assemble
> wf inline.bin          ; write contents of file
```

Appending a `?` to a command will show its help message, for example, `p?`. Appending `?*` will show commands starting with the given string, e.g. `p?*`.

To enter visual mode, press `V<enter>`. Use `q` to quit visual mode and return to the prompt.

In visual mode you can use HJKL keys to navigate (left, down, up, and right, respectively). You can use these keys in cursor mode toggled by `c` key. To select a byte range in cursor mode, hold down **SHIFT** key, and press navigation keys HJKL to mark your selection.

While in visual mode, you can also overwrite bytes by pressing `i`. You can press **TAB** to switch between the hex (middle) and string (right) columns. Pressing `q` inside the hex panel returns you to visual mode. By pressing `p` or `P` you can scroll different visual mode representations. There is a second most important visual mode - curses-like panels interface, accessible with `V!` command.

The Framework

The Radare2 project is a set of small command-line utilities that can be used together or independently.

This chapter will give you a quick understanding of them, but you can check the dedicated sections for each tool at the end of this book.

radare2

The main tool of the whole framework. It uses the core of the hexadecimal editor and debugger. `radare2` allows you to open a number of input/output sources as if they were simple, plain files, including disks, network connections, kernel drivers, processes under debugging, and so on.

It implements an advanced command line interface for moving around a file, analyzing data, disassembling, binary patching, data comparison, searching, replacing, and visualizing. It can be scripted with a variety of languages, including Python, Ruby, JavaScript, Lua, and Perl.

rabin2

A program to extract information from executable binaries, such as ELF, PE, Java CLASS, Mach-O, plus any format supported by `r2` plugins. `rabin2` is used by the core to get data like exported symbols, imports, file information, cross references (xrefs), library dependencies, and sections.

rasm2

A command line assembler and disassembler for multiple architectures (including Intel x86 and x86-64, MIPS, ARM, PowerPC, Java, and myriad of others).

Examples

```
$ rasm2 -a java 'nop'
00

$ rasm2 -a x86 -d '90'
nop

$ rasm2 -a x86 -b 32 'mov eax, 33'
b821000000

$ echo 'push eax;nop;nop' | rasm2 -f -
509090
```

rahash2

An implementation of a block-based hash tool. From small text strings to large disks, rahash2 supports multiple algorithms, including MD4, MD5, CRC16, CRC32, SHA1, SHA256, and others. rahash2 can be used to check the integrity or track changes of big files, memory dumps, or disks.

Examples

```
$ rahash2 file
file: 0x00000000-0x00000007 sha256: 887cfbd0d44aaff69f7bdbedebd282ec96191cce9d7fa7336298a18e

$ rahash2 -a md5 file
file: 0x00000000-0x00000007 md5: d1833805515fc34b46c2b9de553f599d
```

radiff2

A binary diffing utility that implements multiple algorithms. It supports byte-level or delta diffing for binary files, and code-analysis diffing to find changes in basic code blocks obtained from the radare code analysis.

rafind2

A program to find byte patterns in files.

ragg2

A frontend for r_egg. ragg2 compiles programs written in a simple high-level language into tiny binaries for x86, x86-64, and ARM.

Examples

```
$ cat hi.r
/* hello world in r_egg */
write@syscall(4); //x64 write@syscall(1);
exit@syscall(1); //x64 exit@syscall(60);
```



```

main@global(128) {
    .var0 = "hi!\n";
    write(1,.var0, 4);
    exit(0);
}
$ ragg2 -0 -F hi.r
$ ./hi
hi!

$ cat hi.c
main@global(0,6) {
    write(1, "Hello0", 6);
    exit(0);
}
$ ragg2 hi.c
$ ./hi.c.bin
Hello

```

rarun2

A launcher for running programs within different environments, with different arguments, permissions, directories, and overridden default file descriptors. rarun2 is useful for:

- Solving crackmes
- Fuzzing
- Test suites

Sample rarun2 script

```

$ cat foo.rr2
#!/usr/bin/rarun2
program=./pp400
arg0=10
stdin=foo.txt
chdir=/tmp
#chroot=.
./foo.rr2

```

Connecting a Program with a Socket

```

$ nc -l 9999
$ rarun2 program=/bin/ls connect=localhost:9999

```

Debugging a Program Redirecting the stdio into Another Terminal

1 - open a new terminal and type ‘tty’ to get a terminal name:

```
$ tty ; clear ; sleep 999999  
/dev/ttyS010
```

2 - Create a new file containing the following rarun2 profile named foo.rr2:

```
#!/usr/bin/rarun2  
program=/bin/ls  
stdio=/dev/ttys010
```

3 - Launch the following radare2 command:

```
r2 -r foo.rr2 -d /bin/ls
```

rax2

A minimalistic mathematical expression evaluator for the shell that is useful for making base conversions between floating point values, hexadecimal representations, hexpair strings to ASCII, octal to integer, and more. It also supports endianness settings and can be used as an interactive shell if no arguments are given.

Examples

```
$ rax2 1337  
0x539
```

```
$ rax2 0x400000  
4194304
```

```
$ rax2 -b 01111001  
y
```

```
$ rax2 -S radare2  
72616461726532
```

```
$ rax2 -s 617765736f6d65  
awesome
```

User Interfaces

Radare2 has seen many different user interfaces being developed over the years.

Maintaining a GUI is far from the scope of developing the core machinery of a reverse engineering toolkit: it is preferred to have a separate project and community, allowing both projects to collaborate and to improve together - rather than forcing cli developers to think in gui problems and having to jump back and forth between the graphic aspect and the low level logic of the implementations.

In the past, there have been at least 5 different native user interfaces (ragui, r2gui, gradare, r2net, bokken) but none of them got enough maintenance power to take off and they all died.

In addition, r2 has an embedded webserver and ships some basic user interfaces written in html/js. You can start them like this:

```
$ r2 -c=H /bin/ls
```

After 3 years of private development, Hugo Teso; the author of Bokken (python-gtk gui of r2) released to the public another frontend of r2, this time written in c++ and qt, which has been very welcomed by the community.

This GUI was named Iaito, but as long as he preferred not to keep maintaining it, Xarkes decided to fork it under the name of Cutter (name voted by the community), and lead the project. This is how it looks:

- <https://github.com/radareorg/cutter>.

Cutter

Windows

To build r2 on Windows you have to use the Meson build system. Despite being able to build r2 on Windows using cygwin, mingw or wsl using the acr/make build system it is not the recommended/official/supported method and may result on unexpected results.

Binary builds can be downloaded from the release page or the github CI artifacts from every single commit for 32bit and 64bit Windows.

- <https://github.com/radareorg/radare2/releases>

Prerequisites

- 3 GB of free disk space
- Visual Studio 2019 (or higher)
- Python 3
- Meson
- Ninja
- Git

Step-by-Step

Install Visual Studio 2015 (or higher) Visual Studio must be installed with a Visual C++ compiler, supporting C++ libraries, and the appropriate Windows SDK for the target platform version.

- Ensure **Programming Languages > Visual C++** is selected

If you need a copy of Visual Studio, the Community versions are free and work great.

- Download Visual Studio 2019

Install Python 3 and Meson via Conda Conda is our probably the best Python distribution for Windows. But you can skip the next steps if you have Python installed already

- <https://docs.conda.io/en/latest/miniconda.html>
- <https://repo.anaconda.com/archive/>

Create a Python Environment for Radare2 Follow these steps to create and activate a Conda environment named *r2*. All instructions from this point on will assume this name matches your environment, but you may change this if desired.

1. Start > Anaconda Prompt
2. `conda create -n r2 python=3`
3. `activate r2`

Any time you wish to enter this environment, open the Anaconda Prompt and re-issue `activate r2`. Conversely, `deactivate` will leave the environment.

Install Meson Ensure Meson is version 0.48 or higher (`meson -v`)

`pip install meson`

Install Git for Windows All Radare2 code is managed via the Git version control system and hosted on GitHub.

Follow these steps to install Git for Windows.

Download Git for Windows

- <https://git-scm.com/download/win>

Check the following options during the Wizard steps.

- Use a TrueType font in all console windows
- Use Git from the Windows Command Prompt
- Use the native Windows Secure Channel library (instead of OpenSSL)
- Checkout Windows-style, commit Unix-style line endings (`core.autocrlf=true`)
- Use Windows' default console window (instead of Mintty)
- Ensure `git --version` works after install

Get Radare2 Code Follow these steps to clone the Radare2 git repository.

`git clone https://github.com/radareorg/radare2`

Compile Radare2 Code Follow these steps to compile the Radare2 Code.

* **Visual Studio 2017:**

Note 1: Change `Community` to either `Professional` or `Enterprise` in the command below

Note 2: Change `vcvars32.bat` to `vcvars64.bat` in the command below for the 64-bit version

`"%ProgramFiles(x86)%\Microsoft Visual Studio\2017\Community\VC\Auxiliary\Build\vcvars32.bat"

4. Generate the build system with Meson:

Meson takes some arguments to configure the build type, but in short you should be able to build r2 without any meson flag. For Visual Studio. Note: Change **Debug** to **Release** in the command below depending on the version desired.

```
meson b --buildtype debug --backend vs2019 --prefix %cd%\dest
msbuild build\radare2.sln /p:Configuration=Debug /m
```

For Ninja (no visual studio interface required, just msvc compiler toolchain installed):

```
meson b
ninja -C b
```

Finally run this line to install r2 into the given absolute prefix directory:

```
meson install -C build --no-rebuild
```

Build options notes The `/m[axcpucount]` switch creates one MSBuild worker process per logical processor on your machine. You can specify a numeric value (e.g. `/m:2`) to limit the number of worker processes if needed. (This should not be confused with the Visual C++ Compiler switch `/MP`.)

If you get an error with the 32-bit install that says something along the lines of **error MSB4126: The specified solution configuration "Debug|x86" is invalid**. Get around this by adding the following argument to the command: `/p:Platform=Win32`

Check your Radare2 version: `dest\bin\radare2.exe -v`

Check That Radare2 Runs From All Locations Note that `r2` in UNIX systems is just a symlink to the `radare2` executable. So, in case you want to have it in Windows you can just `copy radare2.exe r2.exe` and add the directory into the system-wide PATH env var in the **File Explorer** settings.

Open the `cmd.exe` console and type `r2 -v` to confirm the whole process was successful.

Notes about setting up the system-wide env var

1. In the file explorer go to the folder Radare2 was just installed in.
2. From this folder go to **dest > bin** and keep this window open.
3. Go to System Properties: In the Windows search bar enter **sysdm.cpl**.
4. Go to **Advanced > Environment Variables**.
5. Click on the PATH variable and then click edit (if it exists within both the user and system variables, look at the user version).
6. Ensure the file path displayed in the window left open is listed within the PATH variable. If it is not add it and click **ok**.
7. Log out of your Windows session.
8. Open up a new Windows Command Prompt: type **cmd** in the search bar. Ensure that the current path is not in the Radare2 folder.
9. Check Radare2 version from Command Prompt Window: **radare2 -v**

Debugging

It is common to have an issues when you write a plugin, especially if you do this for the first time. This is why debugging them is very important. The first step for debugging is to set an environment variable when running radare2 instance:

```
R_DEBUG=yes r2 /bin/ls
Loading /usr/local/lib/radare2/2.2.0-git//bin_xtr_dyldcache.so
Cannot find symbol 'radare_plugin' in library '/usr/local/lib/radare2/2.2.0-git//bin_xtr_dyldcache.so'
Cannot open /usr/local/lib/radare2/2.2.0-git//2.2.0-git
Loading /home/user/.config/radare2/plugins/asm_mips_ks.so
PLUGIN OK 0x55b205ea6070 fcn 0x7f298de08762
Loading /home/user/.config/radare2/plugins/asm_sparc_ks.so
PLUGIN OK 0x55b205ea6070 fcn 0x7f298de08762
Cannot open /home/user/.config/radare2/plugins/pimp
Cannot open /home/user/.config/radare2/plugins/yara
Loading /home/user/.config/radare2/plugins/asm_arm_ks.so
PLUGIN OK 0x55b205ea6070 fcn 0x7f298de08762
Loading /home/user/.config/radare2/plugins/core_yara.so
Module version mismatch /home/user/.config/radare2/plugins/core_yara.so (2.1.0) vs (2.2.0-git)
Loading /home/user/.config/radare2/plugins/asm_ppc_ks.so
PLUGIN OK 0x55b205ea6070 fcn 0x7f298de08762
Loading /home/user/.config/radare2/plugins/lang_python3.so
PLUGIN OK 0x55b205ea5ed0 fcn 0x7f298de08692
Loading /usr/local/lib/radare2/2.2.0-git/bin_xtr_dyldcache.so
Cannot find symbol 'radare_plugin' in library '/usr/local/lib/radare2/2.2.0-git/bin_xtr_dyldcache.so'
Cannot open /usr/local/lib/radare2/2.2.0-git/2.2.0-git
Cannot open directory '/usr/local/lib/radare2-extras/2.2.0-git'
Cannot open directory '/usr/local/lib/radare2-bindings/2.2.0-git'
USER CONFIG loaded from /home/user/.config/radare2/radare2rc
-- In visual mode press 'c' to toggle the cursor mode. Use tab to navigate
```

```
[0x00005520]>
```

Implementing a new analysis plugin

After implementing disassembly plugin, you might have noticed that output is far from being good - no proper highlighting, no reference lines and so on. This is because radare2 requires every architecture plugin to provide also analysis information about every opcode. At the moment the implementation of disassembly and opcodes analysis is separated between two modules - RAsm and RAnal. Thus we need to write an analysis plugin too. The principle is very similar - you just need to create a C file and corresponding Makefile.

They structure of RAnal plugin looks like

```
RAnalPlugin r_anal_plugin_v810 = {
    .name = "mycpu",
    .desc = "MYCPU code analysis plugin",
    .license = "LGPL3",
    .arch = "mycpu",
    .bits = 32,
    .op = mycpu_op,
    .esil = true,
    .set_reg_profile = set_reg_profile,
};
```

Like with disassembly plugin there is a key function - `mycpu_op` which scans the opcode and builds `RAnalOp` structure. On the other hand, in this example analysis plugins also performs uplifting to ESIL, which is enabled in `.esil = true` statement. Thus, `mycpu_op` obliged to fill the corresponding `RAnalOp` ESIL field for the opcodes. Second important thing for ESIL uplifting and emulation - register profile, like in debugger, which is set within `set_reg_profile` function.

Makefile

```
NAME=anal_snes
R2_PLUGIN_PATH=$(shell r2 -H R2_USER_PLUGINS)
LIBEXT=$(shell r2 -H LIBEXT)
CFLAGS=-g -fPIC $(shell pkg-config --cflags r_anal)
LDFLAGS=-shared $(shell pkg-config --libs r_anal)
OBJS=$(NAME).o
LIB=$(NAME).$(LIBEXT)

all: $(LIB)

clean:
    rm -f $(LIB) $(OBJS)

$(LIB): $(OBJS)
```

```

$(CC) $(CFLAGS) $(LDFLAGS) $(OBJS) -o $(LIB)

install:
    cp -f anal_snes.$(SO_EXT) $(R2_PLUGIN_PATH)

uninstall:
    rm -f $(R2_PLUGIN_PATH)/anal_snes.$(SO_EXT)

anal_snes.c:
    /* radare - LGPL - Copyright 2015 - condret */

#include <string.h>
#include <r_types.h>
#include <r_lib.h>
#include <r_asm.h>
#include <r_anal.h>
#include "snes_op_table.h"

static int snes_anop(RAnal *anal, RAnalOp *op, ut64 addr, const ut8 *data, int len) {
    memset (op, '\0', sizeof (RAnalOp));
    op->size = snes_op[data[0]].len;
    op->addr = addr;
    op->type = R_ANAL_OP_TYPE_UNK;
    switch (data[0]) {
        case 0xea:
            op->type = R_ANAL_OP_TYPE_NOP;
            break;
    }
    return op->size;
}

struct r_anal_plugin_t r_anal_plugin_snes = {
    .name = "snes",
    .desc = "SNES analysis plugin",
    .license = "LGPL3",
    .arch = R_SYS_ARCH_NONE,
    .bits = 16,
    .init = NULL,
    .fini = NULL,
    .op = &snes_anop,
    .set_reg_profile = NULL,
    .fingerprint_bb = NULL,
    .fingerprint_fcn = NULL,
    .diff_bb = NULL,
    .diff_fcn = NULL,
    .diff_eval = NULL

```



```
};

#ifdef R2_PLUGIN_INCORE
R_API RLibStruct radare_plugin = {
    .type = R_LIB_TYPE_ANAL,
    .data = &r_anal_plugin_snes,
    .version = R2_VERSION
};
#endif
```

After compiling radare2 will list this plugin in the output:

```
_dA_ _8_16      snes      LGPL3    SuperNES CPU
```

`snes_op_table.h`: https://github.com/radareorg/radare2/blob/master/libr/asm/arch/snes/snes_op_table.h

Example:

- **6502**: <https://github.com/radareorg/radare2/commit/64636e9505f9ca8b408958d3c01ac8e3ce254a9b>
- **SNES**: <https://github.com/radareorg/radare2/commit/60d6e5a1b9d244c7085b22ae8985d00027624b49>

Implementing a new disassembly plugin

Radare2 has modular architecture, thus adding support for a new architecture is very easy, if you are fluent in C. For various reasons it might be easier to implement it out of the tree. For this we will need to create single C file, called `asm_mycpu.c` and makefile for it.

The key thing of RAsm plugin is a structure

```
RAsmPlugin r_asm_plugin_mycpu = {
    .name = "mycpu",
    .license = "LGPL3",
    .desc = "MYCPU disassembly plugin",
    .arch = "mycpu",
    .bits = 32,
    .endian = R_SYS_ENDIAN_LITTLE,
    .disassemble = &disassemble
};
```

where `.disassemble` is a pointer to disassembly function, which accepts the bytes buffer and length:

```
static int disassemble(RAsm *a, RAsmOp *op, const ut8 *buf, int len)
```

Makefile

```
NAME=asm_snes
R2_PLUGIN_PATH=$(shell r2 -H R2_USER_PLUGINS)
LIBEXT=$(shell r2 -H LIBEXT)
CFLAGS=-g -fPIC $(shell pkg-config --cflags r_anal)
```

```

LDFLAGS=-shared $(shell pkg-config --libs r_anal)
OBJS=$(NAME).o
LIB=$(NAME).$(LIBEXT)

all: $(LIB)

clean:
    rm -f $(LIB) $(OBJS)

$(LIB): $(OBJS)
    $(CC) $(CFLAGS) $(LDFLAGS) $(OBJS) -o $(LIB)

install:
    cp -f asm_mycpu.$(SO_EXT) $(R2_PLUGIN_PATH)

uninstall:
    rm -f $(R2_PLUGIN_PATH)/asm_mycpu.$(SO_EXT)

asm_mycpu.c
/* radare - LGPL - Copyright 2018 - user */

#include <stdio.h>
#include <string.h>
#include <r_types.h>
#include <r_lib.h>
#include <r_asm.h>

static int disassemble(RAsm *a, RAsmOp *op, const ut8 *buf, int len) {
    struct op_cmd cmd = {
        .instr = "",
        .operands = ""
    };
    if (len < 2) return -1;
    int ret = decode_opcode (buf, len, &cmd);
    if (ret > 0) {
        snprintf (op->buf_asm, R_ASM_BUFSIZE, "%s %s",
            cmd.instr, cmd.operands);
    }
    return op->size = ret;
}

RAsmPlugin r_asm_plugin_mycpu = {
    .name = "mycpu",
    .license = "LGPL3",
    .desc = "MYCPU disassembly plugin",
    .arch = "mycpu",

```

```

        .bits = 32,
        .endian = R_SYS_ENDIAN_LITTLE,
        .disassemble = &disassemble
};

#ifdef R2_PLUGIN_INCORE
R_API RLibStruct radare_plugin = {
    .type = R_LIB_TYPE_ASM,
    .data = &r_asm_plugin_mycpu,
    .version = R2_VERSION
};
#endif

```

After compiling radare2 will list this plugin in the output:

```
_d__  _8_32      mycpu      LGPL3    MYCPU
```

Moving plugin into the tree

Pushing a new architecture into the main branch of r2 requires to modify several files in order to make it fit into the way the rest of plugins are built.

List of affected files:

- `plugins.def.cfg` : add the `asm.mycpu` plugin name string in there
- `libr/asm/p/mycpu.mk` : build instructions
- `libr/asm/p/asm_mycpu.c` : implementation
- `libr/include/r_asm.h` : add the struct definition in there

Check out how the NIOS II CPU disassembly plugin was implemented by reading those commits:

Implement RAsm plugin: <https://github.com/radareorg/radare2/commit/933dc0ef6ddfe44c88bbb261165bf8b55>

Implement RAnal plugin: <https://github.com/radareorg/radare2/commit/ad430f0d52f9e933e0830c49ee607e9b0e>

Implementing a new format

To enable virtual addressing

In `info` add `et->has_va = 1;` and `ptr->srwx` with the `R_BIN_SCN_MAP`; attribute

Create a folder with file format name in `libr/bin/format`

Makefile:

```

NAME=bin_nes
R2_PLUGIN_PATH=$(shell r2 -H R2_USER_PLUGINS)
LIBEXT=$(shell r2 -H LIBEXT)
CFLAGS=-g -fPIC $(shell pkg-config --cflags r_bin)

```

```

LDFLAGS=-shared $(shell pkg-config --libs r_bin)
OBJS=$(NAME).o
LIB=$(NAME).$(LIBEXT)

```

```

all: $(LIB)

```

```

clean:
    rm -f $(LIB) $(OBJS)

```

```

$(LIB): $(OBJS)
    $(CC) $(CFLAGS) $(LDFLAGS) $(OBJS) -o $(LIB)

```

```

install:
    cp -f $(NAME).$(SO_EXT) $(R2_PLUGIN_PATH)

```

```

uninstall:
    rm -f $(R2_PLUGIN_PATH)/$(NAME).$(SO_EXT)

```

```

bin_nes.c:

```

```

#include <r_util.h>
#include <r_bin.h>

```

```

static bool load_buffer(RBinFile *bf, void **bin_obj, RBuffer *b, ut64 loadaddr, Sdb *sdb) {
    ut64 size;
    const ut8 *buf = r_buf_data (b, &size);
    r_return_val_if_fail (buf, false);
    *bin_obj = r_bin_internal_nes_load (buf, size);
    return *bin_obj != NULL;
}

```

```

static void destroy(RBinFile *bf) {
    r_bin_free_all_nes_obj (bf->o->bin_obj);
    bf->o->bin_obj = NULL;
}

```

```

static bool check_buffer(RBuffer *b) {
    if (!buf || length < 4) return false;
    return (!memcmp (buf, "\x4E\x45\x53\x1A", 4));
}

```

```

static RBinInfo* info(RBinFile *arch) {
    RBinInfo *ret = R_NEW0 (RBinInfo);
    if (!ret) return NULL;

    if (!arch || !arch->buf) {
        free (ret);
    }
}

```

```

        return NULL;
    }
    ret->file = strdup (arch->file);
    ret->type = strdup ("ROM");
    ret->machine = strdup ("Nintendo NES");
    ret->os = strdup ("nes");
    ret->arch = strdup ("6502");
    ret->bits = 8;

    return ret;
}

struct r_bin_plugin_t r_bin_plugin_nes = {
    .name = "nes",
    .desc = "NES",
    .license = "BSD",
    .get_sdb = NULL,
    .load_buffer = &load_buffer,
    .destroy = &destroy,
    .check_buffer = &check_buffer,
    .baddr = NULL,
    .entries = NULL,
    .sections = NULL,
    .info = &info,
};

#ifdef R2_PLUGIN_INCORE
R_API RLibStruct radare_plugin = {
    .type = R_LIB_TYPE_BIN,
    .data = &r_bin_plugin_nes,
    .version = R2_VERSION
};
#endif

```

Some Examples

- XBE - <https://github.com/radareorg/radare2/pull/972>
- COFF - <https://github.com/radareorg/radare2/pull/645>
- TE - <https://github.com/radareorg/radare2/pull/61>
- Zimgz - <https://github.com/radareorg/radare2/commit/d1351cf836df3e2e63043a6dc728e880316f00eb>
- OMF - <https://github.com/radareorg/radare2/commit/44fd8b2555a0446ea759901a94c06f20566bbc40>

Charset

1. Create a file in `radare2/libr/util/d/yourfile.sdb.txt`. The extension `.sdb.txt` is important.

2. Edit the file `radare2/libr/util/charset.c`. -add `extern SdbGperf gperf_latin_1_ISO_8859_1_western_european;`. -then add your variable `&gperf_latin_1_ISO_8859_1_western_european`, in static `const SdbGperf *gperfs[]`
3. Update the Makefile: `radare2/libr/util/Makefile`: -Add `OBS+=d/latin_1_ISO_8859_1_western_european`
4. Update the Makefile `radare2/libr/util/d/Makefile` to add your file name with not `.sdb` and not `.txt` in `FILES=latin_1_ISO_8859_1_western_european`
5. Update the unit tests of `radare2/test/db/cmd/charset`

Congratulation! You can now type the command:

```
e cfg.charset=latin_1_ISO_8859_1_western_european;
```

If you have any issue with this tutorial you can check out the example at <https://github.com/radareorg/radare2/pull/19627/files>.

Implementing a new architecture

radare2 splits the logic of a CPU into several modules. You should write more than one plugin to get full support for a specific arch. Let's see which are those:

- `r_asm` : assembler and disassembler
- `r_anal` : code analysis (opcode,type,esil,...)
- `r_reg` : registers
- `r_syscall` : system calls
- `r_debug` : debugger

The most basic feature you usually want to support from a specific architecture is the disassembler. You first need to read into a human readable form the bytes in there.

Bear in mind that plugins can be compiled static or dynamically, this means that the arch will be embedded inside the core libraries or it will be distributed as a separated shared library.

To configure which plugins you want to compile use the `./configure-plugins` script which accepts the flags `--shared` and `--static` to specify them. You can also add it manually inside the `plugins.def.cfg` and then remove the `plugins.cfg` and run `./configure-plugins` again to update the `libr/config.mk` and `libr/config.h`.

You may find some examples of external plugins in radare2-extras repository.

Writing the `r_asm` plugin

The official way to make third-party plugins is to distribute them into a separate repository. This is a sample disasm plugin:

```
$ cd my-cpu
$ cat Makefile
```

```

NAME=mycpu
R2_PLUGIN_PATH=$(shell r2 -hh|grep R2_LIBR_PLUGINS|awk '{print $$2}')
CFLAGS=-g -fPIC $(shell pkg-config --cflags r_asm)
LDFLAGS=-shared $(shell pkg-config --libs r_asm)
OBJS=$(NAME).o
SO_EXT=$(shell uname|grep -q Darwin && echo dylib || echo so)
LIB=$(NAME).$(SO_EXT)

all: $(LIB)

clean:
    rm -f $(LIB) $(OBJS)

$(LIB): $(OBJS)
    $(CC) $(CFLAGS) $(LDFLAGS) $(OBJS) -o $(LIB)

install:
    cp -f $(NAME).$(SO_EXT) $(R2_PLUGIN_PATH)

uninstall:
    rm -f $(R2_PLUGIN_PATH)/$(NAME).$(SO_EXT)

$ cat mycpu.c
/* example r_asm plugin by pancake at 2014 */

#include <r_asm.h>
#include <r_lib.h>

#define OPS 17

static const char *ops[OPS*2] = {
    "nop", NULL,
    "if", "r",
    "ifnot", "r",
    "add", "rr",
    "addi", "ri",
    "sub", "ri",
    "neg", "ri",
    "xor", "ri",
    "mov", "ri",
    "cmp", "rr",
    "load", "ri",
    "store", "ri",
    "shl", "ri",
    "br", "r",
    "bl", "r",

```

```

    "ret", NULL,
    "sys", "i"
};

/* Main function for disassembly */
//b for byte, l for length
static int disassemble (RAsm *a, RAsmOp *op, const ut8 *b, int l) {
    char arg[32];
    int idx = (b[0]&0xf)*2;
    op->size = 2;
    if (idx>=(OPS*2)) {
        strcpy (op->buf_asm, "invalid");
        return -1;
    }
    strcpy (op->buf_asm, ops[idx]);
    if (ops[idx+1]) {
        const char *p = ops[idx+1];
        arg[0] = 0;
        if (!strcmp (p, "rr")) {
            sprintf (arg, "r%d, r%d", b[1]>>4, b[1]&0xf);
        } else
        if (!strcmp (p, "i")) {
            sprintf (arg, "%d", (char)b[1]);
        } else
        if (!strcmp (p, "r")) {
            sprintf (arg, "r%d, r%d", b[1]>>4, b[1]&0xf);
        } else
        if (!strcmp (p, "ri")) {
            sprintf (arg, "r%d, %d", b[1]>>4, (char)b[1]&0xf);
        }
        if (*arg) {
            strcat (op->buf_asm, " ");
            strcat (op->buf_asm, arg);
        }
    }
    return op->size;
}

/* Structure of exported functions and data */
RAsmPlugin r_asm_plugin_mycpu = {
    .name = "mycpu",
    .arch = "mycpu",
    .license = "LGPL3",
    .bits = 32,
    .desc = "My CPU disassembler",
    .disassemble = &disassemble,

```



```
};

#ifdef CORELIB
struct r_lib_struct_t radare_plugin = {
    .type = R_LIB_TYPE_ASM,
    .data = &r_asm_plugin_mycpu
};
#endif
```

To build and install this plugin just type this:

```
$ make
$ sudo make install
```

Write a debugger plugin

- Adding the debugger registers profile into the shlr/gdb/src/core.c
- Adding the registers profile and architecture support in the libr/debug/p/debug_native.c and libr/debug/p/debug_gdb.c
- Add the code to apply the profiles into the function `r_debug_gdb_attach(RDebug *dbg, int pid)`

If you want to add support for the gdb, you can see the register profile in the active gdb session using command `maint print registers`.

More to come..

- Related article: <http://radare.today/posts/extending-r2-with-new-plugins/>

Some commits related to “Implementing a new architecture”

- Extensa: <https://github.com/radareorg/radare2/commit/6f1655c49160fe9a287020537afe0fb8049085d7>
- Malbolge: <https://github.com/radareorg/radare2/pull/579>
- 6502: <https://github.com/radareorg/radare2/pull/656>
- h8300: <https://github.com/radareorg/radare2/pull/664>
- GBA: <https://github.com/radareorg/radare2/pull/702>
- CR16: <https://github.com/radareorg/radare2/pull/721/> && 726
- XCore: <https://github.com/radareorg/radare2/commit/bb16d1737ca5a471142f16ccfa7d444d2713a54d>
- SharpLH5801: <https://github.com/neuschaefer/radare2/commit/f4993cca634161ce6f82a64596fce45fe6b818>
- MSP430: <https://github.com/radareorg/radare2/pull/1426>
- HP-PA-RISC: <https://github.com/radareorg/radare2/commit/f8384feb6ba019b91229adb8fd6e0314b0656f7>
- V810: <https://github.com/radareorg/radare2/pull/2899>
- TMS320: <https://github.com/radareorg/radare2/pull/596>

Implementing a new pseudo architecture

This is an simple plugin for z80 that you may use as example:

<https://github.com/radareorg/radare2/commit/8ff6a92f65331cf8ad74cd0f44a60c258b137a06>

Plugins

radare2 is implemented on top of a bunch of libraries, almost every of those libraries support plugins to extend the capabilities of the library or add support for different targets.

This section aims to explain what are the plugins, how to write them and use them

Types of plugins

```
$ ls libr/*/p | grep : | awk -F / '{ print $2 }'
```

anal	# analysis plugins
asm	# assembler/disassembler plugins
bin	# binary format parsing plugins
bp	# breakpoint plugins
core	# core plugins (implement new commands)
crypto	# encrypt/decrypt/hash/...
debug	# debugger backends
egg	# shellcode encoders, etc
fs	# filesystems and partition tables
io	# io plugins
lang	# embedded scripting languages
parse	# disassembler parsing plugins
reg	# arch register logic

Listing plugins

Some r2 tools have the -L flag to list all the plugins associated to the functionality.

```
rasm2 -L    # list asm plugins
r2 -L       # list io plugins
rabin2 -L   # list bin plugins
rahash2 -L  # list hash/crypto/encoding plugins
```

There are more plugins in r2land, we can list them from inside r2, and this is done by using the L suffix.

Those are some of the commands:

L	# list core plugins
iL	# list bin plugins
dL	# list debug plugins
mL	# list fs plugins
ph	# print support hash algorithms

You can use the ? as value to get the possible values in the associated eval vars.

```
e asm.arch=? # list assembler/disassembler plugins
e anal.arch=? # list analysis plugins
```

Notes

Note there are some inconsistencies that most likely will be fixed in the future radare2 versions.

IO plugins

All access to files, network, debugger and all input/output in general is wrapped by an IO abstraction layer that allows radare to treat all data as if it were just a file.

IO plugins are the ones used to wrap the open, read, write and ‘system’ on virtual file systems. You can make radare understand anything as a plain file. E.g. a socket connection, a remote radare session, a file, a process, a device, a gdb session.

So, when radare reads a block of bytes, it is the task of an IO plugin to get these bytes from any place and put them into internal buffer. An IO plugin is chosen by a file’s URI to be opened. Some examples:

- Debugging URIs

```
$ r2 dbg:///bin/ls<br />
$ r2 pid:///1927
```

- Remote sessions

```
$ r2 rap://:1234<br />
$ r2 rap://<host>:1234//bin/ls
```

- Virtual buffers

```
$ r2 malloc:///512<br />
shortcut for
$ r2 -
```

You can get a list of the radare IO plugins by typing radare2 -L:

```
$ r2 -L
rw_ ar          Open ar/lib files [ar|lib]://[file//path] (LGPL3)
rw_ bfdbg       BrainFuck Debugger (bfdbg://path/to/file) (LGPL3)
rwd bochs       Attach to a BOCHS debugger (LGPL3)
r_d debug       Native debugger (dbg:///bin/ls dbg:///1388 pidof:// waitfor://) (LGPL3) v0.2.0
rw_ default     open local files using def_mmap:// (LGPL3)
rwd gdb         Attach to gdbserver, 'qemu -s', gdb://localhost:1234 (LGPL3)
rw_ gprobe      open gprobe connection using gprobe:// (LGPL3)
```

```

rw_ gzip      read/write gzipped files (LGPL3)
rw_ http      http get (http://rada.re/) (LGPL3)
rw_ ihex      Intel HEX file (ihex://eeproms.hex) (LGPL)
r__ mach      mach debug io (unsupported in this platform) (LGPL)
rw_ malloc    memory allocation (malloc://1024 hex://cd8090) (LGPL3)
rw_ mmap      open file using mmap:// (LGPL3)
rw_ null      null-plugin (null://23) (LGPL3)
rw_ procpid   /proc/pid/mem io (LGPL3)
rwd ptrace    ptrace and /proc/pid/mem (if available) io (LGPL3)
rwd qnx       Attach to QNX pdebug instance, qnx://host:1234 (LGPL3)
rw_ r2k       kernel access API io (r2k://) (LGPL3)
rw_ r2pipe    r2pipe io plugin (MIT)
rw_ r2web     r2web io client (r2web://cloud.rada.re/cmd/) (LGPL3)
rw_ rap       radare network protocol (rap://:port rap://host:port/file) (LGPL3)
rw_ rbuf      RBuffer IO plugin: rbuf:// (LGPL)
rw_ self      read memory from myself using 'self://' (LGPL3)
rw_ shm       shared memory resources (shm://key) (LGPL3)
rw_ sparse    sparse buffer allocation (sparse://1024 sparse://) (LGPL3)
rw_ tcp       load files via TCP (listen or connect) (LGPL3)
rwd windbg    Attach to a KD debugger (windbg://socket) (LGPL3)
rwd winedbg   Wine-dbg io and debug.io plugin for r2 (MIT)
rw_ zip       Open zip files [apk|ipa|zip|zipall]://[file//path] (BSD)

```

Python plugins

At first, to be able to write a plugins in Python for radare2 you need to install r2lang plugin: `r2pm -i lang-python`. Note - in the following examples there are missing functions of the actual decoding for the sake of readability!

For this you need to do this: 1. `import r2lang` and `from r2lang import R` (for constants) 2. Make a function with 2 subfunctions - `assemble` and `disassemble` and returning plugin structure - for RAsm plugin

```

def mycpu(a):
    def assemble(s):
        return [1, 2, 3, 4]

    def disassemble(memview, addr):
        try:
            opcode = get_opcode(memview) # https://docs.python.org/3/library/stdtypes.html#
            opstr = optbl[opcode][1]
            return [4, opstr]
        except:
            return [4, "unknown"]

```

3. This structure should contain a pointers to these 2 functions - `assemble`

and disassemble

```
return {
    "name" : "mycpu",
    "arch" : "mycpu",
    "bits" : 32,
    "endian" : R.R_SYS_ENDIAN_LITTLE,
    "license" : "GPL",
    "desc" : "MYCPU disasm",
    "assemble" : assemble,
    "disassemble" : disassemble,
}
```

4. Make a function with 2 subfunctions - `set_reg_profile` and `op` and returning plugin structure - for RAnal plugin

```
def mycpu_anal(a):
    def set_reg_profile():
        profile = "=PC pc\n" + \
            "=SP sp\n" + \
            "gpr r0 .32 0 0\n" + \
            "gpr r1 .32 4 0\n" + \
            "gpr r2 .32 8 0\n" + \
            "gpr r3 .32 12 0\n" + \
            "gpr r4 .32 16 0\n" + \
            "gpr r5 .32 20 0\n" + \
            "gpr sp .32 24 0\n" + \
            "gpr pc .32 28 0\n"
        return profile

    def op(memview, pc):
        analop = {
            "type" : R.R_ANAL_OP_TYPE_NULL,
            "cycles" : 0,
            "stackop" : 0,
            "stackptr" : 0,
            "ptr" : -1,
            "jump" : -1,
            "addr" : 0,
            "eob" : False,
            "esil" : "",
        }

        try:
            opcode = get_opcode(memview) # https://docs.python.org/3/library/stdtypes.html#
            esilstr = optbl[opcode][2]
            if optbl[opcode][0] == "J": # it's jump
                analop["type"] = R.R_ANAL_OP_TYPE_JMP
```

```

        analop["jump"] = decode_jump(opcode, j_mask)
        esilstr = jump_esil(esilstr, opcode, j_mask)

    except:
        result = analop
        # Don't forget to return proper instruction size!
    return [4, result]

```

5. This structure should contain a pointers to these 2 functions - set_reg_profile and op

```

return {
    "name" : "mycpu",
    "arch" : "mycpu",
    "bits" : 32,
    "license" : "GPL",
    "desc" : "MYCPU anal",
    "esil" : 1,
    "set_reg_profile" : set_reg_profile,
    "op" : op,
}

```

6. (Optional) To add extra information about op sizes and alignment, add a archinfo subfunction and point to it in the structure

```

def mycpu_anal(a):
    def set_reg_profile():
        [...]

    def archinfo(query):
        if query == R.R_ANAL_ARCHINFO_MIN_OP_SIZE:
            return 1
        if query == R.R_ANAL_ARCHINFO_MAX_OP_SIZE:
            return 8
        if query == R.R_ANAL_ARCHINFO_INV_OP_SIZE: # invalid op size
            return 2
        return 0

    def analop(memview, pc):
        [...]

    return {
        "name" : "mycpu",
        "arch" : "mycpu",
        "bits" : 32,
        "license" : "GPL",
        "desc" : "MYCPU anal",
        "esil" : 1,
        "set_reg_profile" : set_reg_profile,
        "archinfo": archinfo,
    }

```

```

        "op" : op,
    }

```

7. Register both plugins using `r2lang.plugin("asm")` and `r2lang.plugin("anal")` respectively

```

print("Registering MYCPU disasm plugin...")
print(r2lang.plugin("asm", mycpu))
print("Registering MYCPU analysis plugin...")
print(r2lang.plugin("anal", mycpu_anal))

```

You can combine everything in one file and load it using `-i` option:

```
r2 -I mycpu.py some_file.bin
```

Or you can load it from the r2 shell: `#!python mycpu.py`

See also:

- Python
- Javascript

Implementing new format plugin in Python

Note - in the following examples there are missing functions of the actual decoding for the sake of readability!

For this you need to do this: 1. `import r2lang` 2. Make a function with subfunctions: `- load - load_bytes - destroy - check_bytes - baddr - entries - sections - imports - relocs - binsym - info`

and returning plugin structure - for RAsm plugin

```

def le_format(a):
    def load(binf):
        return [0]

    def check_bytes(buf):
        try:
            if buf[0] == 77 and buf[1] == 90:
                lx_off, = struct.unpack("<I", buf[0x3c:0x40])
                if buf[lx_off] == 76 and buf[lx_off+1] == 88:
                    return [1]
            return [0]
        except:
            return [0]

```

and so on. Please be sure of the parameters for each function and format of returns. Note, that functions `entries`, `sections`, `imports`, `relocs` returns a list of special formed dictionaries - each with a different type. Other functions

return just a list of numerical values, even if single element one. There is a special function, which returns information about the file - `info`:

```
def info(binfmt):
    return [{
        "type" : "le",
        "bclass" : "le",
        "rclass" : "le",
        "os" : "OS/2",
        "subsystem" : "CLI",
        "machine" : "IBM",
        "arch" : "x86",
        "has_va" : 0,
        "bits" : 32,
        "big_endian" : 0,
        "dbg_info" : 0,
    }]
```

3. This structure should contain a pointers to the most important functions like `check_bytes`, `load` and `load_bytes`, `entries`, `relocs`, `imports`.

```
return {
    "name" : "le",
    "desc" : "OS/2 LE/LX format",
    "license" : "GPL",
    "load" : load,
    "load_bytes" : load_bytes,
    "destroy" : destroy,
    "check_bytes" : check_bytes,
    "baddr" : baddr,
    "entries" : entries,
    "sections" : sections,
    "imports" : imports,
    "symbols" : symbols,
    "relocs" : relocs,
    "binsym" : binsym,
    "info" : info,
}
```

4. Then you need to register it as a file format plugin:

```
print("Registering OS/2 LE/LX plugin...")
print(r2lang.plugin("bin", le_format))
```


Creating an r2pm package of the plugin

As you remember radare2 has its own packaging manager and we can easily add newly written plugin for everyone to access.

All packages are located in radare2-pm repository, and have very simple text format.

```
R2PM_BEGIN

R2PM_GIT "https://github.com/user/mycpu"
R2PM_DESC "[r2-arch] MYCPU disassembler and analyzer plugins"

R2PM_INSTALL() {
    ${MAKE} clean
    ${MAKE} all || exit 1
    ${MAKE} install R2PM_PLUGDIR="${R2PM_PLUGDIR}"
}

R2PM_UNINSTALL() {
    rm -f "${R2PM_PLUGDIR}/asm_mycpu.*"
    rm -f "${R2PM_PLUGDIR}/anal_mycpu.*"
}

R2PM_END
```

Then add it in the /db directory of radare2-pm repository and send a pull request to the mainline.

Testing the plugin

This plugin is used by rasm2 and r2. You can verify that the plugin is properly loaded with this command:

```
$ rasm2 -L | grep mycpu
_d mycpu          My CPU disassembler  (LGPL3)
```

Let's open an empty file using the 'mycpu' arch and write some random code there.

```
$ r2 -
-- I endians swap
[0x00000000]> e asm.arch=mycpu
[0x00000000]> woR
[0x00000000]> pd 10
           0x00000000    888e    mov r8, 14
           0x00000002    b2a5    ifnot r10, r5
           0x00000004    3f67    ret
```

0x00000006	7ef6	bl r15, r6
0x00000008	2701	xor r0, 1
0x0000000a	9826	mov r2, 6
0x0000000c	478d	xor r8, 13
0x0000000e	6b6b	store r6, 11
0x00000010	1382	add r8, r2
0x00000012	7f15	ret

Yay! it works.. and the mandatory oneliner too!

```
r2 -nqamycpu -cwoR -cpd' 10' -
```

Radare2 Reference Card

This chapter is based on the Radare 2 reference card by Thanat0s, which is under the GNU GPL. Original license is as follows:

This card may be freely distributed under the terms of the GNU general public licence - Copyright by Thanat0s - v0.1 -

Survival Guide

Those are the basic commands you will want to know and use for moving around a binary and getting information about it.

Command	Description
s (tab)	Seek to a different place
x [nbytes]	Hexdump of nbytes, \$b by default
aa	Auto analyze
pdf@ funcname	Disassemble function (main, fcn, etc.)
f fcn(Tab)	List functions
f str(Tab)	List strings
fr [flagname] [newname]	Rename flag
psz [offset]~grep	Print strings and grep for one
axF [flag]	Find cross reference for a flag

Flags

Flags are like bookmarks, but they carry some extra information like size, tags or associated flagospace. Use the **f** command to list, set, get them.

Command	Description
f	List flags
fd \$\$	Describe an offset
fj	Display flags in JSON

Command	Description
fl	Show flag length
fx [flagname]	Show hexdump of flag
fC [name] [comment]	Set flag comment

Flagspaces

Flags are created into a flagospace, by default none is selected, and listing flags will list them all. To display a subset of flags you can use the **fs** command to restrict it.

Command	Description
fs	Display flagspaces
fs *	Select all flagspaces
fs [space]	Select one flagospace

Information

Binary files have information stored inside the headers. The **i** command uses the RBin api and allows us to do the same things rabin2 do. Those are the most common ones.

Command	Description
ii	Information on imports
iI	Info on binary
ie	Display entrypoint
iS	Display sections
ir	Display relocations
iz	List strings (izz, izzz)

Print string

There are different ways to represent a string in memory. The **ps** command allows us to print it in utf-16, pascal, zero terminated, .. formats.

Command	Description
psz [offset]	Print zero terminated string
psb [offset]	Print strings in current block
psx [offset]	Show string with escaped chars
psp [offset]	Print pascal string
psw [offset]	Print wide string

Visual mode

The visual mode is the standard interactive interface of radare2.

To enter in visual mode use the `v` or `V` command, and then you'll only have to press keys to get the actions happen instead of commands.

Command	Description
V	Enter visual mode
p/P	Rotate modes (hex, disasm, debug, words, buf)
c	Toggle (c)ursor
q	Back to Radare shell
h/jkl	Move around (or HJKL) (left-down-up-right)
Enter	Follow address of jump/call
sS	Step/step over
o	Toggle asm.pseudo and asm.esil
.	Seek to program counter
/	In cursor mode, search in current block
:cmd	Run radare command
;-]cmt	Add/remove comment
/*+~	Change block size, ~ = resize hex.cols
<, >	Seek aligned to block size
i/a/A	(i)nsert hex, (a)ssemble code, visual (A)ssembler
b	Toggle breakpoint
B	Browse evals, symbols, flags, classes, ...
d[f?]	Define function, data, code, ..
D	Enter visual diff mode (set diff.from/to)
e	Edit eval configuration variables
f/F	Set/unset flag
gG	Go seek to begin and end of file (0-\$s)
mK/'K	Mark/go to Key (any key)
M	Walk the mounted filesystems
n/N	Seek next/prev function/flag/hit (scr.nkey)
C	Toggle (C)olors
R	Randomize color palette (ecr)
tT	Tab related. see also tab
v	Visual code analysis menu
V	(V)iew graph (agv?)
wW	Seek cursor to next/prev word
uU	Undo/redo seek
x	Show xrefs of current func from/to data/code
yY	Copy and paste selection
z	fold/unfold comments in diassembly

Searching

There are many situations where we need to find a value inside a binary or in some specific regions. Use the `e search.in=?` command to choose where the `/` command may search for the given value.

Command	Description
<code>/ foo\00</code>	Search for string 'foo\0'
<code>/b</code>	Search backwards
<code>//</code>	Repeat last search
<code>/w foo</code>	Search for wide string 'f\0o\0o\0'
<code>/wi foo</code>	Search for wide string ignoring case
<code>/! ff</code>	Search for first occurrence not matching
<code>/i foo</code>	Search for string 'foo' ignoring case
<code>/e /E.F/i</code>	Match regular expression
<code>/x a1b2c3</code>	Search for bytes; spaces and uppercase nibbles are allowed, same as <code>/x A1 B2 C3</code>
<code>/x a1..c3</code>	Search for bytes ignoring some nibbles (auto-generates mask, in this example: <code>ff00ff</code>)
<code>/x a1b2:fff3</code>	Search for bytes with mask (specify individual bits)
<code>/d 101112</code>	Search for a deltified sequence of bytes
<code>/!x 00</code>	Inverse hexa search (find first byte != 0x00)
<code>/c jmp [esp]</code>	Search for asm code (see <code>search.asmstr</code>)
<code>/a jmp eax</code>	Assemble opcode and search its bytes
<code>/A</code>	Search for AES expanded keys
<code>/r sym.printf</code>	Analyze opcode reference an offset
<code>/R</code>	Search for ROP gadgets
<code>/P</code>	Show offset of previous instruction
<code>/m magicfile</code>	Search for matching magic file
<code>/p patternsize</code>	Search for pattern of given size
<code>/z min max</code>	Search for strings of given size
<code>/v[?248] num</code>	Look for a asm.bigendian 32bit value

Saving (Broken)

This feature has broken and not been resolved at the time of writing these words (Nov.16th 2020). check [#Issue 6945: META - Project files](#) and [#Issue 17034](#) for more details.

To save your analysis for now, write your own script which records the function name, variable name, etc. for example:

```
vim sample_A.r2

e scr.utf8 = false
s 0x000403ce0
```

```

aaa
s fcn.00403130
afn return_delta_to_heapaddr
afvn iter var_04h
...

```

Usable variables in expression

The `??` command will display the variables that can be used in any math operation inside the r2 shell. For example, using the `? $$` command to evaluate a number or `?v` to just the value in one format.

All commands in r2 that accept a number supports the use of those variables.

Command	Description
<code> here(currentvirtualseek)</code>	non-temporary virtual seek
<code> \$? </code>	last comparison value
<code>\$alias=value</code>	alias commands (simple macros)
<code>\$b</code>	block size
<code>\$B</code>	base address (aligned lowest map address)
<code>\$f</code>	jump fail address (e.g. <code>jz 0x10 => next instruction</code>)
<code>\$fl</code>	flag length (size) at current address (fla; <code>pD \$l @ entry0</code>)
<code>\$F</code>	current function size
<code>\$FB</code>	begin of function
<code>\$Fb</code>	address of the current basic block
<code>\$Fs</code>	size of the current basic block
<code>\$FE</code>	end of function
<code>\$FS</code>	function size
<code>\$Fj</code>	function jump destination
<code>\$Ff</code>	function false destination
<code>\$FI</code>	function instructions
<code>c,r</code>	get width and height of terminal
<code>\$Cn</code>	get nth call of function
<code>\$Dn</code>	get nth data reference in function
<code>\$D</code>	current debug map base address <code>?v \$D @ rsp</code>
<code>\$DD</code>	current debug map size
<code>\$e</code>	1 if end of block, else 0
<code>\$j</code>	jump address (e.g. <code>jmp 0x10, jz 0x10 => 0x10</code>)
<code>\$Ja</code>	get nth jump of function
<code>\$Xn</code>	get nth xref of function
<code>\$l</code>	opcode length
<code>\$m</code>	opcode memory reference (e.g. <code>mov eax,[0x10] => 0x10</code>)
<code>\$M</code>	map address (lowest map address)
<code>\$o</code>	here (current disk io offset)
<code>\$p</code>	<code>getpid()</code>
<code>\$P</code>	pid of children (only in debug)

Command	Description
\$s	file size
\$S	section offset
\$\$S	section size
\$v	opcode immediate value (e.g. lui a0,0x8010 => 0x8010)
\$w	get word size, 4 if asm.bits=32, 8 if 64, ...
\${ev}	get value of eval config variable
\$r{reg}	get value of named register
\$k{kv}	get value of an sdb query value
\$s{flag}	get size of flag
RNum	\$variables usable in math expressions

Scripting

Radare2 provides a wide set of a features to automate boring work. It ranges from the simple sequencing of the commands to the calling scripts/another programs via IPC (Inter-Process Communication), called r2pipe.

As mentioned a few times before there is an ability to sequence commands using ; semicolon operator.

```
[0x00404800]> pd 1 ; ao 1
               0x00404800      b827e66100      mov eax, 0x61e627      ; "tab"
address: 0x404800
opcode: mov eax, 0x61e627
prefix: 0
bytes: b827e66100
ptr: 0x0061e627
refptr: 0
size: 5
type: mov
esil: 6415911, rax, =
stack: null
family: cpu
[0x00404800]>
```

It simply runs the second command after finishing the first one, like in a shell.

The second important way to sequence the commands is with a simple pipe |

```
ao|grep address
```

Note, the | pipe only can pipe output of r2 commands to external (shell) commands, like system programs or builtin shell commands. There is a similar way to sequence r2 commands, using the backtick operator ``command``. The quoted part will undergo command substitution and the output will be used as an argument of the command line.

For example, we want to see a few bytes of the memory at the address referred to by the 'mov eax, addr' instruction. We can do that without jumping to it, using a sequence of commands:

```
[0x00404800]> pd 1
               0x00404800      b827e66100      mov eax, 0x61e627      ; "tab"
[0x00404800]> ao
address: 0x404800
opcode: mov eax, 0x61e627
prefix: 0
bytes: b827e66100
ptr: 0x0061e627
refptr: 0
size: 5
type: mov
esil: 6415911, rax, =
stack: null
family: cpu
[0x00404800]> ao~ptr[1]
0x0061e627
0
[0x00404800]> px 10 @ `ao~ptr[1]`
- offset -   0 1  2 3  4 5  6 7  8 9  A B  C D  E F  0123456789ABCDEF
0x0061e627  7461 6200 2e69 6e74 6572                                tab..inter
[0x00404800]>
```

And of course it's possible to redirect the output of an r2 command into a file, using the > and >> commands

```
[0x00404800]> px 10 @ `ao~ptr[1]` > example.txt
[0x00404800]> px 10 @ `ao~ptr[1]` >> example.txt
```

Radare2 also provides quite a few Unix type file processing commands like head, tail, cat, grep and many more. One such command is Uniq, which can be used to filter a file to display only non-duplicate content. So to make a new file with only unique strings, you can do:

```
[0x00404800]> uniq file > uniq_file
```

The head command can be used to see the first N number of lines in the file, similarly tail command allows the last N number of lines to be seen.

```
[0x00404800]> head 3 foodtypes.txt
1 Protein
2 Carbohydrate
3 Fat
[0x00404800]> tail 2 foodtypes.txt
3 Shake
4 Milk
```


The join command could be used to merge two different files with common first field.

```
[0x00404800]> cat foodtypes.txt
1 Protein
2 Carbohydrate
3 Fat
[0x00404800]> cat foods.txt
1 Cheese
2 Potato
3 Butter
[0x00404800]> join foodtypes foods.txt
1 Protein Cheese
2 Carbohydrate Potato
3 Fat Butter
```

Similarly, sorting the content is also possible with the sort command. A typical example could be:

```
[0x00404800]> sort file
eleven
five
five
great
one
one
radare
```

The `??` command describes several helpful variables you can use to do similar actions even more easily, like the `$v` “immediate value” variable, or the `$m` opcode memory reference variable.

Loops

One of the most common task in automation is looping through something, there are multiple ways to do this in radare2.

We can loop over flags:

```
@@ flagname-regex
```

For example, we want to see function information with `afi` command:

```
[0x004047d6]> afi
#
offset: 0x004047d0
name: entry0
size: 42
realsz: 42
```

```

stackframe: 0
call-convention: amd64
cyclomatic-complexity: 1
bits: 64
type: fcn [NEW]
num-bbs: 1
edges: 0
end-bbs: 1
call-refs: 0x00402450 C
data-refs: 0x004136c0 0x00413660 0x004027e0
code-xrefs:
data-xrefs:
locals:0
args: 0
diff: type: new
[0x004047d6]>

```

Now let's say, for example, that we'd like see a particular field from this output for all functions found by analysis. We can do that with a loop over all function flags (whose names begin with `fcn.`):

```

[0x004047d6]> fs functions
[0x004047d6]> afi @@ fcn.* ~name

```

This command will extract the `name` field from the `afi` output of every flag with a name matching the regexp `fcn.*`. There are also a predefined loop called `@@f`, which runs your command on every functions found by `r2`:

```

[0x004047d6]> afi @@f ~name

```

We can also loop over a list of offsets, using the following syntax:

```
@@=1 2 3 ... N
```

For example, say we want to see the opcode information for 2 offsets: the current one, and at current + 2:

```

[0x004047d6]> ao @@=$$ $$+2
address: 0x4047d6
opcode: mov rdx, rsp
prefix: 0
bytes: 4889e2
refptr: 0
size: 3
type: mov
esil: rsp,rdx,=
stack: null
family: cpu
address: 0x4047d8
opcode: loop 0x404822

```

```

prefix: 0
bytes: e248
refptr: 0
size: 2
type: cjmp
esil: 1,rcx,-=,rcx,{,4212770,rip,=,}
jump: 0x00404822
fail: 0x004047da
stack: null
cond: al
family: cpu
[0x004047d6]>

```

Note we're using the \$\$ variable which evaluates to the current offset. Also note that \$\$+2 is evaluated before looping, so we can use the simple arithmetic expressions.

A third way to loop is by having the offsets be loaded from a file. This file should contain one offset per line.

```

[0x004047d0]> ?v $$ > offsets.txt
[0x004047d0]> ?v $$+2 >> offsets.txt
[0x004047d0]> !cat offsets.txt
4047d0
4047d2
[0x004047d0]> pi 1 @@.offsets.txt
xor ebp, ebp
mov r9, rdx

```

radare2 also offers various `foreach` constructs for looping. One of the most useful is for looping through all the instructions of a function:

```

[0x004047d0]> pdf
/ (fcn) entry0 42
|; UNKNOWN XREF from 0x00400018 (unk)
|; DATA XREF from 0x004064bf (sub.strlen_460)
|; DATA XREF from 0x00406511 (sub.strlen_460)
|; DATA XREF from 0x0040b080 (unk)
|; DATA XREF from 0x0040b0ef (unk)
|0x004047d0 xor ebp, ebp
|0x004047d2 mov r9, rdx
|0x004047d5 pop rsi
|0x004047d6 mov rdx, rsp
|0x004047d9 and rsp, 0xfffffffffffffff0
|0x004047dd push rax
|0x004047de push rsp
|0x004047df mov r8, 0x4136c0
|0x004047e6 mov rcx, 0x413660 ; "AWA..AVI..AUI..ATL.%.. "

```

```

0A..AVI..AUI.
|0x004047ed  mov rdi, main          ; "AWAVAUATUH..S..H...." @
0
|0x004047f4  call sym.imp.__libc_start_main
\0x004047f9  hlt
[0x004047d0]> pi 1 @@i
mov r9, rdx
pop rsi
mov rdx, rsp
and rsp, 0xfffffffffffffff0
push rax
push rsp
mov r8, 0x4136c0
mov rcx, 0x413660
mov rdi, main
call sym.imp.__libc_start_main
hlt

```

In this example the command `pi 1` runs over all the instructions in the current function (entry0). There are other options too (not complete list, check `@@?` for more information): - `@@k sdbquery` - iterate over all offsets returned by that `sdbquery` - `@@t` - iterate over on all threads (see `dp`) - `@@b` - iterate over all basic blocks of current function (see `afb`) - `@@f` - iterate over all functions (see `afq`)

The last kind of looping lets you loop through predefined iterator types:

- symbols
- imports
- registers
- threads
- comments
- functions
- flags

This is done using the `@@@` command. The previous example of listing information about functions can also be done using the `@@@` command:

```
[0x004047d6]> afi @@@ functions ~name
```

This will extract `name` field from `afi` output and will output a huge list of function names. We can choose only the second column, to remove the redundant `name:` on every line:

```
[0x004047d6]> afi @@@ functions ~name[1]
```

Beware, `@@@` is not compatible with JSON commands.

Macros

Apart from simple sequencing and looping, radare2 allows to write simple macros, using this construction:

```
[0x00404800]> (qwe; pd 4; ao)
```

This will define a macro called 'qwe' which runs sequentially first 'pd 4' then 'ao'. Calling the macro using syntax `.(macro)` is simple:

```
[0x00404800]> (qwe; pd 4; ao)
[0x00404800]> .(qwe)
0x00404800  mov eax, 0x61e627      ; "tab"
0x00404805  push rbp
0x00404806  sub rax, section_end.LOAD1
0x0040480c  mov rbp, rsp
```

```
address: 0x404800
opcode: mov eax, 0x61e627
prefix: 0
bytes: b827e66100
ptr: 0x0061e627
refptr: 0
size: 5
type: mov
esil: 6415911,rax,=
stack: null
family: cpu
[0x00404800]>
```

To list available macroses simply call `(*)`:

```
[0x00404800]> (*)
(qwe ; pd 4; ao)
```

And if want to remove some macro, just add '-' before the name:

```
[0x00404800]> (-qwe)
Macro 'qwe' removed.
[0x00404800]>
```

Moreover, it's possible to create a macro that takes arguments, which comes in handy in some simple scripting situations. To create a macro that takes arguments you simply add them to macro definition.

```
[0x00404800]
[0x004047d0]> (foo x y; pd $0; s +$1)
[0x004047d0]> .(foo 5 6)
;-- entry0:
0x004047d0      xor ebp, ebp
```

```

0x004047d2      mov r9, rdx
0x004047d5      pop rsi
0x004047d6      mov rdx, rsp
0x004047d9      and rsp, 0xfffffffffffffff0
[0x004047d6]>

```

As you can see, the arguments are named by index, starting from 0: \$0, \$1, ...

Aliases

radare2 also offers aliases which might help you save time by quickly executing your most used commands. They are under \$?

The general usage of the feature is: \$alias=cmd

```
[0x00404800]> $disas=pdf
```

The above command will create an alias **disas** for **pdf**. The following command prints the disassembly of the main function.

```
[0x00404800]> $disas @ main
```

Apart from commands, you can also alias a text to be printed, when called.

```

[0x00404800]> $my_alias=$test input
[0x00404800]> $my_alias
test input

```

To undefine alias, use \$alias=:

```

[0x00404800]> $pmore='b 300;px'
[0x00404800]> $
$pmore
[0x00404800]> $pmore=
[0x00404800]> $

```

A single \$ in the above will list all defined aliases. It's also possible check the aliased command of an alias:

```

[0x00404800]> $pmore?
b 200; px

```

Can we create an alias contains alias ? The answer is yes:

```

[0x00404800]> $pStart='s 0x0;$pmore'
[0x00404800]> $pStart
- offset -   0 1  2 3  4 5  6 7  8 9  A B  C D  E F  0123456789ABCDEF
0x00000000  7f45 4c46 0201 0100 0000 0000 0000 0000  .ELF.....
0x00000010  0300 3e00 0100 0000 1014 0000 0000 0000  ..>.....
0x00000020  4000 0000 0000 0000 5031 0000 0000 0000  @.....P1.....
0x00000030  0000 0000 4000 3800 0d00 4000 1e00 1d00  ....@.8...@.....

```

```

0x00000040 0600 0000 0400 0000 4000 0000 0000 0000 .....@.....
0x00000050 4000 0000 0000 0000 4000 0000 0000 0000 @.....@.....
0x00000060 d802 0000 0000 0000 d802 0000 0000 0000 .....
0x00000070 0800 0000 0000 0000 0300 0000 0400 0000 .....
0x00000080 1803 0000 0000 0000 1803 0000 0000 0000 .....
0x00000090 1803 0000 0000 0000 1c00 0000 0000 0000 .....
0x000000a0 1c00 0000 0000 0000 0100 0000 0000 0000 .....
0x000000b0 0100 0000 0400 0000 0000 0000 0000 0000 .....
0x000000c0 0000 0000 0000 0000 .....
[0x00000000]>

```

R2pipe

The r2pipe api was initially designed for NodeJS in order to support reusing the web's r2.js API from the commandline. The r2pipe module permits interacting with r2 instances in different methods:

- spawn pipes (r2 -0)
- http queries (cloud friendly)
- tcp socket (r2 -c)

	pipe	spawn	async	http	tcp	rap	json
nodejs	x	x	x	x	x	-	x
python	x	x	-	x	x	x	x
swift	x	x	x	x	-	-	x
dotnet	x	x	x	x	-	-	-
haskell	x	x	-	x	-	-	x
java	-	x	-	x	-	-	-
golang	x	x	-	-	-	-	x
ruby	x	x	-	-	-	-	x
rust	x	x	-	-	-	-	x
vala	-	x	x	-	-	-	-
erlang	x	x	-	-	-	-	-
newlisp	x	-	-	-	-	-	-
dlang	x	-	-	-	-	-	x
perl	x	-	-	-	-	-	-

Examples

Python

```

$ pip install r2pipe

import r2pipe

r2 = r2pipe.open("/bin/ls")

```

```

r2.cmd('aa')
print(r2.cmd("afl"))
print(r2.cmdj("aflj"))  # evaluates JSONs and returns an object

```

NodeJS

Use this command to install the r2pipe bindings

```
$ npm install r2pipe
```

Here's a sample hello world

```

const r2pipe = require('r2pipe');
r2pipe.open('/bin/ls', (err, res) => {
  if (err) {
    throw err;
  }
  r2.cmd ('af @ entry0', function (o) {
    r2.cmd ("pdf @ entry0", function (o) {
      console.log (o);
      r.quit ()
    });
  });
});

```

Checkout the GIT repository for more examples and details.

<https://github.com/radareorg/radare2-r2pipe/blob/master/nodejs/r2pipe/README.md>

Go

```
$ r2pm -i r2pipe-go
```

<https://github.com/radare/r2pipe-go>

```
package main
```

```

import (
    "fmt"
    "github.com/radare/r2pipe-go"
)

func main() {
    r2p, err := r2pipe.NewPipe("/bin/ls")
    if err != nil {
        panic(err)
    }
    defer r2p.Close()
    buf1, err := r2p.Cmd("?E Hello World")
    if err != nil {

```



```

        panic(err)
    }
    fmt.Println(buf1)
}

```

Rust

```

$ cat Cargo.toml
...
[dependencies]
r2pipe = "*"

#[macro_use]
extern crate r2pipe;
use r2pipe::R2Pipe;
fn main() {
    let mut r2p = open_pipe!(Some("/bin/ls")).unwrap();
    println!("{:?}", r2p.cmd("?e Hello World"));
    let json = r2p.cmdj("ij").unwrap();
    println!("{}", serde_json::to_string_pretty(&json).unwrap());
    println!("ARCH {}", json["bin"]["arch"]);
    r2p.close();
}

```

Ruby

```

$ gem install r2pipe

require 'r2pipe'
puts 'r2pipe ruby api demo'
puts '===== '
r2p = R2Pipe.new '/bin/ls'
puts r2p.cmd 'pi 5'
puts r2p.cmd 'pij 1'
puts r2p.json(r2p.cmd 'pij 1')
puts r2p.cmd 'px 64'
r2p.quit

```

Perl

```

#!/usr/bin/perl

use R2::Pipe;
use strict;

my $r = R2::Pipe->new ("/bin/ls");
print $r->cmd ("pd 5")."\n";

```

```
print $r->cmd ("px 64")."\n";
$r->quit ();
```

Erlang

```
#!/usr/bin/env escript
%% -*- erlang -*-
%%! -smp enable

%% -sname hr
-mode(compile).

-export([main/1]).

main(_Args) ->
    %% adding r2pipe to modulepath, set it to your r2pipe_erl location
    R2pipePATH = filename:dirname(escript:script_name()) ++ "/ebin",
    true = code:add_pathz(R2pipePATH),

    %% initializing the link with r2
    H = r2pipe:init(lpipe),

    %% all work goes here
    io:format("~s", [r2pipe:cmd(H, "i")]).
```

Haskell

```
import R2pipe
import qualified Data.ByteString.Lazy as L

showMainFunction ctx = do
    cmd ctx "s main"
    L.putStr =<< cmd ctx "pD `fl $$`"

main = do
    -- Run r2 locally
    open "/bin/ls" >=> showMainFunction
    -- Connect to r2 via HTTP (e.g. if "r2 -qc=h /bin/ls" is running)
    open "http://127.0.0.1:9090" >=> showMainFunction
```

Dotnet

```
using System;
using System.Collections.Generic;
using System.Diagnostics;
using System.Linq;
```

```

using System.Text;
using System.Threading.Tasks;
using r2pipe;

namespace LocalExample {
    class Program {
        static void Main(string[] args) {
#if __MonoCS__
            using (IR2Pipe pipe = new R2Pipe("/bin/ls")) {
#else
            using (IR2Pipe pipe = new R2Pipe(@"C:\Windows\notepad.exe",
                @"C:\radare2\radare2.exe")) {
#endif
                Console.WriteLine("Hello r2! " + pipe.RunCommand("?V"));
                Task<string> async = pipe.RunCommandAsync("?V");
                Console.WriteLine("Hello async r2!" + async.Result);
                QueuedR2Pipe qr2 = new QueuedR2Pipe(pipe);
                qr2.Enqueue(new R2Command("x", (string result) => {
                    Console.WriteLine("Result of x:\n {0}", result); }));
                qr2.Enqueue(new R2Command("pi 10", (string result) => {
                    Console.WriteLine("Result of pi 10:\n {0}", result); }));
                qr2.ExecuteCommands();
            }
        }
    }
}

```

Java

```

import org.radare.r2pipe.R2Pipe;

public class Test {
    public static void main (String[] args) {
        try {
            R2Pipe r2p = new R2Pipe ("/bin/ls");
            // new R2Pipe ("http://cloud.rada.re/cmd/", true);
            System.out.println (r2p.cmd ("pd 10"));
            System.out.println (r2p.cmd ("px 32"));
            r2p.quit();
        } catch (Exception e) {
            System.err.println (e);
        }
    }
}

```

Swift

```
if let r2p = R2Pipe(url:nil) {
    r2p.cmd ("?V", closure:{
        (str:String?) in
        if let s = str {
            print ("Version: \(s)");
            exit (0);
        } else {
            debugPrint ("R2PIPE. Error");
            exit (1);
        }
    });
    NSRunLoop.currentRunLoop().run();
} else {
    print ("Needs to run from r2")
}
```

Vala

```
public static int main (string[] args) {
    MainLoop loop = new MainLoop ();
    var r2p = new R2Pipe ("/bin/ls");
    r2p.cmd ("pi 4", (x) => {
        stdout.printf ("Disassembly:\n%s\n", x);
        r2p.cmd ("ie", (x) => {
            stdout.printf ("Entrypoint:\n%s\n", x);
            r2p.cmd ("q");
        });
    });
    ChildWatch.add (r2p.child_pid, (pid, status) => {
        Process.close_pid (pid);
        loop.quit ();
    });
    loop.run ();
    return 0;
}
```

NewLisp

```
(load "r2pipe.lsp")
(println "pd 3:\n" (r2pipe:cmd "pd 3"))
(exit)
```

Dlang

```
import std.stdio;
import r2pipe;

void main() {
    auto r2 = r2pipe.open ();
    writeln ("Hello "~ r2.cmd("?e World"));
    writeln ("Hello "~ r2.cmd("?e Works"));

    string uri = r2.cmdj("ij")["core"]["uri"].str;
    writeln ("Uri: ",uri);
}
```

Search Automation

The `cmd.hit` configuration variable is used to define a radare2 command to be executed when a matching entry is found by the search engine. If you want to run several commands, separate them with `;`. Alternatively, you can arrange them in a separate script, and then invoke it as a whole with `. script-file-name` command. For example:

```
[0x00404888]> e cmd.hit = p8 8
[0x00404888]> / lib
Searching 3 bytes from 0x00400000 to 0x0041ae08: 6c 69 62
hits: 9
0x00400239 hit4_0 "lib64/ld-linux-x86-64.so.2"
31ed4989d15e4889
0x00400f19 hit4_1 "libselinux.so.1"
31ed4989d15e4889
0x00400fae hit4_2 "librt.so.1"
31ed4989d15e4889
0x00400fc7 hit4_3 "libacl.so.1"
31ed4989d15e4889
0x00401004 hit4_4 "libc.so.6"
31ed4989d15e4889
0x004013ce hit4_5 "libc_start_main"
31ed4989d15e4889
0x00416542 hit4_6 "libs/"
31ed4989d15e4889
0x00417160 hit4_7 "lib/xstrtol.c"
31ed4989d15e4889
0x00417578 hit4_8 "lib"
31ed4989d15e4889
```

Searching Backwards

Sometimes you want to find a keyword backwards. This is, before the current offset, to do this you can seek back and search forward by adding some search.from/to restrictions, or use the /b command.

```
[0x100001200]> / nop
0x100004b15 hit0_0 .STUWabcdefghijklmnopqrstuvwxyzbin/ls.
0x100004f50 hit0_1 .STUWabcdefghijklmnopqrstuvwxyz1] [file .
[0x100001200]> /b nop
[0x100001200]> s 0x100004f50p
[0x100004f50]> /b nop
0x100004b15 hit2_0 .STUWabcdefghijklmnopqrstuvwxyzbin/ls.
[0x100004f50]>
```

Note that /b is doing the same as /, but backward, so what if we want to use /x backward? We can use /bx, and the same goes for other search subcommands:

```
[0x100001200]> /x 90
0x100001a23 hit1_0 90
0x10000248f hit1_1 90
0x1000027b2 hit1_2 90
0x100002b2e hit1_3 90
0x1000032b8 hit1_4 90
0x100003454 hit1_5 90
0x100003468 hit1_6 90
0x10000355b hit1_7 90
0x100003647 hit1_8 90
0x1000037ac hit1_9 90
0x10000389c hit1_10 90
0x100003c5c hit1_11 90

[0x100001200]> /bx 90
[0x100001200]> s 0x10000355b
[0x10000355b]> /bx 90
0x100003468 hit3_0 90
0x100003454 hit3_1 90
0x1000032b8 hit3_2 90
0x100002b2e hit3_3 90
0x1000027b2 hit3_4 90
0x10000248f hit3_5 90
0x100001a23 hit3_6 90
[0x10000355b]>
```

Basic Search

A basic search for a plain text string in a file would be something like:

```
$ r2 -q -c "/ lib" /bin/ls
Searching 3 bytes from 0x00400000 to 0x0041ae08: 6c 69 62
hits: 9
0x00400239 hit0_0 "lib64/ld-linux-x86-64.so.2"
0x00400f19 hit0_1 "libselinux.so.1"
0x00400fae hit0_2 "librt.so.1"
0x00400fc7 hit0_3 "libacl.so.1"
0x00401004 hit0_4 "libc.so.6"
0x004013ce hit0_5 "libc_start_main"
0x00416542 hit0_6 "libs/"
0x00417160 hit0_7 "lib/xstrtol.c"
0x00417578 hit0_8 "lib"
```

As can be seen from the output above, radare2 generates a “hit” flag for every entry found. You can then use the `ps` command to see the strings stored at the offsets marked by the flags in this group, and they will have names of the form `hit0_<index>`:

```
[0x00404888]> / ls
...
[0x00404888]> ps @ hit0_0
lseek
```

You can search for wide-char strings (e.g., unicode letters) using the `/w` command:

```
[0x00000000]> /w Hello
0 results found.
```

To perform a case-insensitive search for strings use `/i`:

```
[0x0040488f]> /i Stallman
Searching 8 bytes from 0x00400238 to 0x0040488f: 53 74 61 6c 6c 6d 61 6e
[# ]hits: 004138 < 0x0040488f hits = 0
```

It is possible to specify hexadecimal escape sequences in the search string by prepending them with `\x`:

```
[0x00000000]> / \x7FELF
```

if, instead, you are searching for a string of hexadecimal values, you’re probably better off using the `/x` command:

```
[0x00000000]> /x 7F454C46
```

If you want to mask some nibble during the search you can use the symbol `.` to allow any nibble value to match:

```
[0x00407354]> /x 80..80
0x0040d4b6 hit3_0 800080
0x0040d4c8 hit3_1 808080
0x004058a6 hit3_2 80fb80
```

You may not know some bit values of your hexadecimal pattern. Thus you may use a bit mask on your pattern. Each bit set to one in the mask indicates to search the bit value in the pattern. A bit set to zero in the mask indicates that the value of a matching value can be 0 or 1:

```
[0x00407354]> /x 808080:ff80ff
0x0040d4c8 hit4_0 808080
0x0040d7b0 hit4_1 808080
0x004058a6 hit4_2 80fb80
```

You can notice that the command `/x 808080:ff00ff` is equivalent to the command `/x 80..80`.

Once the search is done, the results are stored in the `searches` flag space.

```
[0x00000000]> fs
0    0 . strings
1    0 . symbols
2    6 . searches
```

```
[0x00000000]> f
0x00000135 512 hit0_0
0x00000b71 512 hit0_1
0x00000bad 512 hit0_2
0x00000bdd 512 hit0_3
0x00000bfb 512 hit0_4
0x00000f2a 512 hit0_5
```

To remove “hit” flags after you do not need them anymore, use the `f- hit*` command.

Often, during long search sessions, you will need to launch the latest search more than once. You can use the `//` command to repeat the last search.

```
[0x00000f2a]> //      ; repeat last search
```

Configuring Search Options

The radare2 search engine can be configured through several configuration variables, modifiable with the `e` command.

```
e cmd.hit = x          ; radare2 command to execute on every search hit
e search.distance = 0 ; search string distance
e search.in = [foo]    ; specify search boundaries. Supported values are listed under e search
e search.align = 4      ; only show search results aligned by specified boundary.
e search.from = 0       ; start address
e search.to = 0         ; end address
e search.asmstr = 0     ; search for string instead of assembly
e search.flags = true ; if enabled, create flags on hits
```


The `search.align` variable is used to limit valid search hits to certain alignment. For example, with `e search.align=4` you will see only hits found at 4-bytes aligned offsets.

The `search.flags` boolean variable instructs the search engine to flag hits so that they can be referenced later. If a currently running search is interrupted with `Ctrl-C` keyboard sequence, current search position is flagged with `search_stop`.

Searching for Bytes

The radare2 search engine is based on work done by esteve, plus multiple features implemented on top of it. It supports multiple keyword searches, binary masks, and hexadecimal values. It automatically creates flags for search hit locations ease future referencing.

Search is initiated by `/` command.

```
[0x00000000]> /?
|Usage: /[!bf] [arg]Search stuff (see 'e??search' for options)
|Use io.va for searching in non virtual addressing spaces
| / foo\x00          search for string 'foo\0'
| /j foo\x00         search for string 'foo\0' (json output)
| /! ff             search for first occurrence not matching, command modifier
| /!x 00            inverse hexa search (find first byte != 0x00)
| /+ /bin/sh        construct the string with chunks
| //               repeat last search
| /a jmp eax        assemble opcode and search its bytes
| /A jmp           find analyzed instructions of this type (/A? for help)
| /b               search backwards, command modifier, followed by other command
| /B               search recognized RBin headers
| /c jmp [esp]      search for asm code matching the given string
| /ce rsp,rbp       search for esil expressions matching
| /C[ar]           search for crypto materials
| /d 101112         search for a deltified sequence of bytes
| /e /E.F/i        match regular expression
| /E esil-expr      offset matching given esil expressions %%= here
| /f               search forwards, command modifier, followed by other command
| /F file [off] [sz] search contents of file with offset and size
| /g[g] [from]     find all graph paths A to B (/gg follow jumps, see search.count and
anal.depth)
| /h[t] [hash] [len] find block matching this hash. See ph
| /i foo           search for string 'foo' ignoring case
| /m magicfile     search for matching magic file (use blocksize)
| /M               search for known filesystems and mount them automatically
| /o [n]           show offset of n instructions backward
| /O [n]           same as /o, but with a different fallback if anal cannot be used
```

/p patternsize	search for pattern of given size
/P patternsize	search similar blocks
/r[erwx][?] sym.printf	analyze opcode reference an offset (/re for esil)
/R [grepopcode]	search for matching ROP gadgets, semicolon-separated
/s	search for all syscalls in a region (EXPERIMENTAL)
/v[1248] value	look for an `cfg.bigendian` 32bit value
/V[1248] min max	look for an `cfg.bigendian` 32bit value in range
/w foo	search for wide string 'f\0o\0o\0'
/wi foo	search for wide string ignoring case 'f\0o\0o\0'
/x ff..33	search for hex string ignoring some nibbles
/x ff0033	search for hex string
/x ff43:ffd0	search for hexpair with mask
/z min max	search for strings of given size

Because everything is treated as a file in radare2, it does not matter whether you search in a socket, a remote device, in process memory, or a file.

note that `'/'` starts *multiline comment*. It's not for searching. type `'''` to end comment.

Pattern Matching Search

The `/p` command allows you to apply repeated pattern searches on IO backend storage. It is possible to identify repeated byte sequences without explicitly specifying them. The only command's parameter sets minimum detectable pattern length. Here is an example:

```
[0x00000000]> /p 10
```

This command output will show different patterns found and how many times each of them is encountered.

It is possible to search patterns with a known difference between consecutive bytes with `/d` command. For example, the command to search all the patterns with the first and second bytes having the first bit which differs and the second and third bytes with the second bit which differs is:

```
[0x00000000]> /d 0102
Searching 2 bytes in [0x0-0x400]
hits: 2
0x00000118 hit2_0 9a9b9d
0x00000202 hit2_1 a4a5a7
```

Assembler Search

If you want to search for a certain assembler opcodes, you can use `/a` commands.

The command `/ad/ jmp [esp]` searches for the specified category of assembly mnemonic:

```
[0x00404888]> /ad/ jmp qword [rdx]
f hit_0 @ 0x0040e50d # 2: jmp qword [rdx]
f hit_1 @ 0x00418dbb # 2: jmp qword [rdx]
f hit_2 @ 0x00418fcb # 3: jmp qword [rdx]
f hit_3 @ 0x004196ab # 6: jmp qword [rdx]
f hit_4 @ 0x00419bf3 # 3: jmp qword [rdx]
f hit_5 @ 0x00419c1b # 3: jmp qword [rdx]
f hit_6 @ 0x00419c43 # 3: jmp qword [rdx]
```

The command `/a jmp eax` assembles a string to machine code, and then searches for the resulting bytes:

```
[0x00404888]> /a jmp eax
hits: 1
0x004048e7 hit3_0 ffe00f1f8000000000b8
```

Searching for Cryptography materials

Searching AES keys

radare2 is capable of finding **expanded AES** keys with `/ca` command. It searches from current seek position up to the `search.distance` limit, or until end of file is reached. You can interrupt current search by pressing **Ctrl-C**. For example, to look for AES keys in a memory dump:

```
0x00000000]> /ca
Searching 40 bytes in [0x0-0x1ab]
hits: 1
0x000000fb hit0_0 6920e299a5202a6d656e636869746f2a
```

The output length gives you the size of the AES key used: 128, 192 or 256 bits. If you are simply looking for plaintext AES keys in your binary, `/ca` will not find them they must have been expanded by the key expansion algorithm.

Searching private keys and certificates

`/cr` command implements the search of private keys (RSA and ECC). `/cd` command implements a similar feature to search certificates.

```
[0x00000000]> /cr
Searching 11 bytes in [0x0-0x15a]
hits: 2
0x000000fa hit1_0 302e020100300506032b657004220420fb3d588296fed5694ff7049eafb74490bf4bc6467e
```

Entropy analysis

`p=e` might give some hints if high entropy sections are found trying to cover up a hardcoded secret.

There is the possibility to delimit entropy sections for later use with `\s` command:

```

[0x00000000]> b
0x100
[0x00000000]> b 4096
[0x00000000]> /s
0x00100000 - 0x00101000 ~ 5.556094
0x014e2c88 - 0x014e3c88 ~ 0.000000
0x01434374 - 0x01435374 ~ 6.332087
0x01435374 - 0x0144c374 ~ 3.664636
0x0144c374 - 0x0144d374 ~ 1.664368
0x0144d374 - 0x0144f374 ~ 4.229199
0x0144f374 - 0x01451374 ~ 2.000000
(...)
[0x00000000]> /s*
f entropy_section_0 0x00001000 0x00100000
f entropy_section_1 0x00001000 0x014e2c88
f entropy_section_2 0x00001000 0x01434374
f entropy_section_3 0x00017000 0x01435374
f entropy_section_4 0x00001000 0x0144c374
f entropy_section_5 0x00002000 0x0144d374
f entropy_section_6 0x00002000 0x0144f374

```

The blocksize is increased to 4096 bytes from the default 100 bytes so that the entropy search `/s` can work on reasonably sized chunks for entropy analysis. The sections flags can be applied with the dot operator, `./s*` and then looped through `px 32 @@ entropy*`.

Signatures

Radare2 has its own format of the signatures, allowing to both load/apply and create them on the fly. They are available under the `z` command namespace:

```

[0x00000000]> z?
Usage: z[*j-aof/cs] [args] # Manage zignatures
| z          show zignatures
| z.         find matching zignatures in current offset
| zb[?][n=5] search for best match
| z*         show zignatures in radare format
| zq         show zignatures in quiet mode
| zj         show zignatures in json format
| zk         show zignatures in sdb format
| z-zignature delete zignature
| z-*        delete all zignatures
| za[?]      add zignature
| zg         generate zignatures (alias for zaF)
| zo[?]      manage zignature files
| zf[?]      manage FLIRT signatures

```

```
| z/[?]      search signatures
| zc[?]      compare current zignspace signatures with another one
| zs[?]      manage zignspaces
| zi         show zignatures matching information
```

To load the created signature file you need to load it from SDB file using `zo` command or from the compressed SDB file using `zoz` command.

To create signature you need to make function first, then you can create it from the function:

```
r2 /bin/ls
[0x000051c0]> aaa # this creates functions, including 'entry0'
[0x000051c0]> zaf entry0 entry
[0x000051c0]> z
entry:
  bytes: 31ed4989d15e4889e24883e4f050544c.....48.....48.....ff.....
  graph: cc=1 nbbs=1 edges=0 ebbs=1
  offset: 0x000051c0
[0x000051c0]>
```

As you can see it made a new signature with a name `entry` from a function `entry0`. You can show it in JSON format too, which can be useful for scripting:

```
[0x000051c0]> zj~{}
[
  {
    "name": "entry",
    "bytes": "31ed4989d15e4889e24883e4f050544c.....48.....48.....ff.....",
    "graph": {
      "cc": "1",
      "nbbs": "1",
      "edges": "0",
      "ebbs": "1"
    },
    "offset": 20928,
    "refs": [
    ]
  }
]
[0x000051c0]>
```

To remove it just run `z-entry`.

If you want, instead, to save all created signatures, you need to save it into the SDB file using command `zos myentry`.

Then we can apply them. Lets open a file again:

```
r2 /bin/ls
```

```
-- Log On. Hack In. Go Anywhere. Get Everything.
```

```
[0x000051c0]> zo myentry
```

```
[0x000051c0]> z
```

```
entry:
```

```
bytes: 31ed4989d15e4889e24883e4f050544c.....48.....48.....ff.....
```

```
graph: cc=1 nbbs=1 edges=0 ebbs=1
```

```
offset: 0x000051c0
```

```
[0x000051c0]>
```

This means that the signatures were successfully loaded from the file `myentry` and now we can search matching functions:

```
[0x000051c0]> z.
```

```
[+] searching 0x000051c0 - 0x000052c0
```

```
[+] searching function metrics
```

```
hits: 1
```

```
[0x000051c0]>
```

Note that `z.` command just checks the signatures against the current address. To search signatures across the all file we need to do a bit different thing. There is an important moment though, if we just run it “as is” - it wont find anything:

```
[0x000051c0]> z/
```

```
[+] searching 0x0021dfd0 - 0x002203e8
```

```
[+] searching function metrics
```

```
hits: 0
```

```
[0x000051c0]>
```

Note the searching address - this is because we need to adjust the searching range first:

```
[0x000051c0]> e search.in=io.section
```

```
[0x000051c0]> z/
```

```
[+] searching 0x000038b0 - 0x00015898
```

```
[+] searching function metrics
```

```
hits: 1
```

```
[0x000051c0]>
```

We are setting the search mode to `io.section` (it was `file` by default) to search in the current section (assuming we are currently in the `.text` section of course). Now we can check, what radare2 found for us:

```
[0x000051c0]> pd 5
```

```
-- entry0:
```

```
-- sign.bytes.entry_0:
```

0x000051c0	31ed	xor ebp, ebp
0x000051c2	4989d1	mov r9, rdx
0x000051c5	5e	pop rsi
0x000051c6	4889e2	mov rdx, rsp
0x000051c9	4883e4f0	and rsp, 0xfffffffffffffff0

```
[0x000051c0]>
```

Here we can see the comment of `entry0`, which is taken from the ELF parsing, but also the `sign.bytes.entry_0`, which is exactly the result of matching signature.

Signatures configuration stored in the `sign. config vars'` namespace:

```
[0x000051c0]> e? sign.
    sign.autoload: Autoload all signatures located in ~/.local/share/radare2/zigns
    sign.bytes: Use bytes patterns for matching
    sign.diff.bthresh: Threshold for diffing sign bytes [0, 1] (see zc?)
    sign.diff.gthresh: Threshold for diffing sign graphs [0, 1] (see zc?)
    sign.graph: Use graph metrics for matching
    sign.hash: Use Hash for matching
    sign.maxsz: Maximum signature length
    sign.mincc: Minimum cyclomatic complexity for matching
    sign.minsz: Minimum signature length for matching
    sign.offset: Use original offset for matching
    sign.prefix: Default prefix for signatures matches
    sign.refs: Use references for matching
    sign.threshold: Minimum similarity required for inclusion in zb output
    sign.types: Use types for matching
[0x000051c0]>
```

Finding Best Matches `zb`

Often you know the signature should exist somewhere in a binary but `z/` and `z.` still fail. This is often due to very minor differences between the signature and the function. Maybe the compiler switched two instructions, or your signature is not for the correct function version. In these situations the `zb` commands can still help point you in the right direction by listing near matches.

```
[0x000040a0]> zb?
Usage: zb[r?] [args] # search for closest matching signatures
| zb [n]             find n closest matching signatures to function at current offset
| zbr ziname [n]     search for n most similar functions to ziname
```

The `zb` (zign best) command will show the top 5 closest signatures to a function. Each will contain a score between 1.0 and 0.0.

```
[0x0041e390]> s sym.fclose
[0x0040fc10]> zb
0.96032 0.92400 B 0.99664 G sym.fclose
0.65971 0.35600 B 0.96342 G sym._nl_expand_alias
0.65770 0.37800 B 0.93740 G sym.fdopen
0.65112 0.35000 B 0.95225 G sym.__run_exit_handlers
0.62532 0.34800 B 0.90264 G sym.__cxa_finalize
```

In the above example, **zb** correctly associated the **sym.fclose** signature to the current function. The **z/** and **z.** command would have failed to match here since both the Byte and Graph scores are less than 1.0. A 30% separation between the first and second place results is also a good indication of a correct match.

The **zbr** (zign best reverse) accepts a signature name and attempts to find the closest matching functions. Use an analysis command, like **aa** to find functions first.

```
[0x00401b20]> aa
[x] Analyze all flags starting with sym. and entry0 (aa)
[0x00401b20]> zo ./libc.sdb
[0x00401b20]> zbr sym.__libc_malloc 10
0.94873 0.89800 B 0.99946 G sym.malloc
0.65245 0.40600 B 0.89891 G sym._mid_memalign
0.59470 0.38600 B 0.80341 G sym._IO_flush_all_lockp
0.59200 0.28200 B 0.90201 G sym._IO_file_underflow
0.57802 0.30400 B 0.85204 G sym.__libc_realloc
0.57094 0.35200 B 0.78988 G sym.__calloc
0.56785 0.34000 B 0.79570 G sym._IO_un_link.part.0
0.56358 0.36200 B 0.76516 G sym._IO_cleanup
0.56064 0.26000 B 0.86127 G sym.intel_check_word.constprop.0
0.55726 0.28400 B 0.83051 G sym.linear_search_fdes
```

Tools

Radare2 is not just the only tool provided by the radare2 project. The rest of the chapters in this book are focused on explaining the use of the radare2 tool, this chapter will focus on explaining all the other companion tools that are shipped inside the radare2 project.

All the functionalities provided by the different APIs and plugins have also different tools to allow to use them from the commandline and integrate them with shells scripts easily.

Thanks to the orthogonal design of the framework it is possible to do all the things that r2 is able from different places:

- these companion tools
- native library apis
- scripting with r2pipe
- the r2 shell

Visual Mode

The visual mode is a more user-friendly interface alternative to radare2's command-line prompt. It allows easy navigation, has a cursor mode for se-

lecting bytes, and offers numerous key bindings to simplify debugger use. To enter visual mode, use **V** command. To exit from it back to command line, press **q**.

Navigation

Navigation can be done using HJKL or arrow keys and PgUp/PgDown keys. It also understands usual Home/End keys. Like in Vim the movements can be repeated by preceding the navigation key with the number, for example 5j will move down for 5 lines, or 2l will move 2 characters right.

Visual Mode

Print Modes, a.k.a.: Panels

The Visual mode uses “print modes” which are basically different panels that you can rotate. By default those are:

Hexdump panel -> Disassembly panel -> Debugger panel -> Hexadecimal words dump panel -> Hex-less hexdump panel -> Op analysis color map panel -> Annotated hexdump panel -> Hexdump panel -> [...]

Notice that the top of the panel contains the command which is used, for example for the disassembly panel:

```
[0x00404890 16% 120 /bin/ls]> pd $r @ entry0
```

Getting Help

To see help on all key bindings defined for visual mode, press **?**:

Visual mode help:

```
?      show this help
??     show the user-friendly hud
%      in cursor mode finds matching pair, or toggle autoblocksz
@      redraw screen every 1s (multi-user view)
^      seek to the beginning of the function
!      enter into the visual panels mode
_      enter the flag/comment/functions/.. hud (same as VF_)
=      set cmd.vprompt (top row)
|      set cmd.cprompt (right column)
.      seek to program counter
\      toggle visual split mode
"      toggle the column mode (uses pC..)
/      in cursor mode search in current block
:cmd   run radare command
;[-]cmt add/remove comment
```

0	seek to beginning of current function
[1-9]	follow jmp/call identified by shortcut (like ;[1])
,file	add a link to the text file
/*+-[]	change block size, [] = resize hex.cols
</>	seek aligned to block size (seek cursor in cursor mode)
a/A	(a)ssemble code, visual (A)ssembler
b	browse symbols, flags, configurations, classes, ...
B	toggle breakpoint
c/C	toggle (c)ursor and (C)olors
d[f?]	define function, data, code, ..
D	enter visual diff mode (set diff.from/to
e	edit eval configuration variables
f/F	set/unset or browse flags. f- to unset, F to browse, ..
gG	go seek to begin and end of file (0-\$s)
hjkl	move around (or HJKL) (left-down-up-right)
i	insert hex or string (in hexdump) use tab to toggle
mK/'K	mark/go to Key (any key)
M	walk the mounted filesystems
n/N	seek next/prev function/flag/hit (scr.nkey)
g	go/seek to given offset
O	toggle asm.pseudo and asm.esil
p/P	rotate print modes (hex, disasm, debug, words, buf)
q	back to radare shell
r	refresh screen / in cursor mode browse comments
R	randomize color palette (ecr)
sS	step / step over
t	browse types
T	enter textlog chat console (TT)
uU	undo/redo seek
v	visual function/vars code analysis menu
V	(V)iew graph using cmd.graph (agv?)
wW	seek cursor to next/prev word
xX	show xrefs/refs of current function from/to data/code
yY	copy and paste selection
z	fold/unfold comments in disassembly
Z	toggle zoom mode
Enter	follow address of jump/call

Function Keys: (See 'e key.'), defaults to:

F2	toggle breakpoint
F4	run to cursor
F7	single step
F8	step over
F9	continue

Visual Assembler

You can use Visual Mode to assemble code using **A**. For example let's replace the **push** by a **jmp**:

Before

Notice the preview of the disassembly and arrows:

After

You need to open the file in writing mode (**r2 -w** or **oo+**) in order to patch the file. You can also use the cache mode: **e io.cache = true** and **wc?**.

Remember that patching files in debug mode only patch the memory not the file.

Visual Configuration Editor

Ve or **e** in visual mode allows you to edit radare2 configuration visually. For example, if you want to change the assembly display just select **asm** in the list and choose your assembly display flavor.

First Select asm

Example switch to pseudo disassembly:

Pseudo disassembly disabled

Pseudo disassembly enabled

Visual Disassembly

Navigation

Move within the Disassembly using arrow keys or **h j k l**. Use **g** to seek directly to a flag or an offset, type it when requested by the prompt: **[offset]>**. Follow a jump or a call using the **number** of your keyboard **[0-9]** and the number on the right in disassembly to follow a call or a jump. In this example typing **1** on the keyboard would follow the call to **sym.imp.__libc_start_main** and therefore, seek at the offset of this symbol.

```
0x00404894      e857dcffff      call sym.imp.__libc_start_main ;[1]
```

Seek back to the previous location using **u**, **U** will allow you to redo the seek.

d as define

d can be used to change the type of data of the current block, several basic types/structures are available as well as more advanced one using **pf** template:

```

d → ...
0x004048f7      48c1e83f      shr rax, 0x3f
d → b
0x004048f7 .byte 0x48
d → B
0x004048f7 .word 0xc148
d → d
0x004048f7 hex length=165 delta=0
0x004048f7  48c1 e83f 4801 c648 d1fe 7415 b800 0000
...

```

To improve code readability you can change how radare2 presents numerical values in disassembly, by default most of disassembly display numerical value as hexadecimal. Sometimes you would like to view it as a decimal, binary or even custom defined constant. To change value format you can use **d** following by **i** then choose what base to work in, this is the equivalent to **ahi**:

```

d → i → ...
0x004048f7      48c1e83f      shr rax, 0x3f
d → i → 10
0x004048f7      48c1e83f      shr rax, 63
d → i → 2
0x004048f7      48c1e83f      shr rax, '?'

```

Usage of the Cursor for Inserting/Patching...

Remember that, to be able to actually edit files loaded in radare2, you have to start it with the **-w** option. Otherwise a file is opened in read-only mode.

Pressing lowercase **c** toggles the cursor mode. When this mode is active, the currently selected byte (or byte range) is highlighted.

Cursor at 0x00404896

The cursor is used to select a range of bytes or simply to point to a byte. You can use the cursor to create a named flag at specific location. To do so, seek to the required position, then press **f** and enter a name for a flag. If the file was opened in write mode using the **-w** flag or the **o+** command, you can also use the cursor to overwrite a selected range with new values. To do so, select a range of bytes (with **HJKL** and **SHIFT** key pressed), then press **i** and enter the hexpair values for the new data. The data will be repeated as needed to fill the range selected. For example:

```

<select 10 bytes in visual mode using SHIFT+HJKL>
<press 'i' and then enter '12 34'>

```

The 10 bytes you have selected will be changed to “12 34 12 34 12 ...”.

The Visual Assembler is a feature that provides a live-preview while you type in new instructions to patch into the disassembly. To use it, seek or place the cursor

at the wanted location and hit the ‘A’ key. To provide multiple instructions, separate them with semicolons, ;.

XREF

When radare2 has discovered a XREF during the analysis, it will show you the information in the Visual Disassembly using **XREF** tag:

```
; DATA XREF from 0x00402e0e (unk)
str.David_MacKenzie:
```

To see where this string is called, press **x**, if you want to jump to the location where the data is used then press the corresponding number [0-9] on your keyboard. (This functionality is similar to **axt**)

X corresponds to the reverse operation aka **axf**.

Function Argument display

To enable this view use this config var **e dbg.funcarg = true**
funcarg

Add a comment

To add a comment press ;.

Type other commands

Quickly type commands using :.

Search

/: allows highlighting of strings in the current display. **:cmd** allows you to use one of the “/?” commands that perform more specialized searches.

The HUDS

The “UserFriendly HUD”

The “UserFriendly HUD” can be accessed using the **??** key-combination. This HUD acts as an interactive Cheat Sheet that one can use to more easily find and execute commands. This HUD is particularly useful for new-comers. For experienced users, the other HUDS which are more activity-specific may be more useful.

The “flag/comment/functions/.. HUD”

This HUD can be displayed using the `_` key, it shows a list of all the flags defined and lets you jump to them. Using the keyboard you can quickly filter the list down to a flag that contains a specific pattern.

Hud input mode can be closed using `^C`. It will also exit when backspace is pressed when the user input string is empty.

Tweaking the Disassembly

The disassembly’s look-and-feel is controlled using the “asm.* configuration keys, which can be changed using the `e` command. All configuration keys can also be edited through the Visual Configuration Editor.

Visual Configuration Editor

This HUD can be accessed using the `e` key in visual mode. The editor allows you to easily examine and change radare2’s configuration. For example, if you want to change something about the disassembly display, select `asm` from the list, navigate to the item you wish to modify it, then select it by hitting `Enter`. If the item is a boolean variable, it will toggle, otherwise you will be prompted to provide a new value.

First Select `asm`

Example switch to pseudo disassembly:

Pseudo disassembly disabled

Pseudo disassembly enabled

Following are some example of eval variable related to disassembly.

Examples

asm.arch: Change Architecture && asm.bits: Word size in bits at assembler You can view the list of all arch using `e asm.arch=?`

```
e asm.arch = dalvik
0x00404870      31ed4989      cmp-long v237, v73, v137
0x00404874      d15e4889      rsub-int v14, v5, 0x8948
0x00404878      e24883e4      ushr-int/lit8 v72, v131, 0xe4
0x0040487c      f0505449c7c0 +invoke-object-init-range {}, method+18772 ;[0]
0x00404882      90244100      add-int v36, v65, v0

e asm.bits = 16
0000:4870      31ed          xor bp, bp
0000:4872      49            dec cx
0000:4873      89d1          mov cx, dx
```

```

0000:4875      5e          pop si
0000:4876      48          dec ax
0000:4877      89e2         mov dx, sp

```

This latest operation can also be done using `&` in Visual mode.

asm.pseudo: Enable pseudo syntax

```

e asm.pseudo = true
0x00404870      31ed          ebp = 0
0x00404872      4989d1        r9 = rdx
0x00404875      5e          pop rsi
0x00404876      4889e2        rdx = rsp
0x00404879      4883e4f0      rsp &= 0xfffffffffffffff0

```

asm.syntax: Select assembly syntax (intel, att, masm...)

```

e asm.syntax = att
0x00404870      31ed          xor %ebp, %ebp
0x00404872      4989d1        mov %rdx, %r9
0x00404875      5e          pop %rsi
0x00404876      4889e2        mov %rsp, %rdx
0x00404879      4883e4f0      and $0xfffffffffffffff0, %rsp

```

asm.describe: Show opcode description

```

e asm.describe = true
0x00404870  xor ebp, ebp    ; logical exclusive or
0x00404872  mov r9, rdx     ; moves data from src to dst
0x00404875  pop rsi        ; pops last element of stack and stores the result in argument
0x00404876  mov rdx, rsp     ; moves data from src to dst
0x00404879  and rsp, -0xf    ; binary and operation between src and dst, stores result on dst

```

Visual Panels

Concept

Visual Panels is characterized by the following core functionalities:

1. Split Screen
2. Display multiple screens such as Symbols, Registers, Stack, as well as custom panels
3. Menu will cover all those commonly used commands for you so that you don't have to memorize any of them

CUI met some useful GUI as the menu, that is Visual Panels.

Panels can be accessed by using `v` or by using `!` from the visual mode.

Overview

Panels Overview

Commands

```
|Visual Ascii Art Panels:
| |      split the current panel vertically
| -      split the current panel horizontally
| :      run r2 command in prompt
| ;      add/remove comment
| _      start the hud input mode
| \      show the user-friendly hud
| ?      show this help
| !      run r2048 game
| .      seek to PC or entrypoint
| *      show decompiler in the current panel
| "      create a panel from the list and replace the current one
| /      highlight the keyword
| (      toggle snow
| &      toggle cache
| [1-9]  follow jmp/call identified by shortcut (like ;[1])
| ' '    (space) toggle graph / panels
| tab    go to the next panel
| Enter  start Zoom mode
| a      toggle auto update for decompiler
| b      browse symbols, flags, configurations, classes, ...
| c      toggle cursor
| C      toggle color
| d      define in the current address. Same as Vd
| D      show disassembly in the current panel
| e      change title and command of current panel
| f      set/add filter keywords
| F      remove all the filters
| g      go/seek to given offset
| G      go/seek to highlight
| i      insert hex
| hjkl   move around (left-down-up-right)
| HJKL   move around (left-down-up-right) by page
| m      select the menu panel
| M      open new custom frame
| n/N    seek next/prev function/flag/hit (scr.nkey)
| p/P    rotate panel layout
| q      quit, or close a tab
| Q      close all the tabs and quit
| r      toggle callhints/jmphints/leahints
```



```
| R      randomize color palette (ecr)
| s/S    step in / step over
| t/T    tab prompt / close a tab
| u/U    undo / redo seek
| w      start Window mode
| V      go to the graph mode
| xX     show xrefs/refs of current function from/to data/code
| z      swap current panel with the first one
```

Basic Usage

Use **tab** to move around the panels until you get to the targeted panel. Then, use **h j k l**, just like in vim, to scroll the panel you are currently on. Use **S** and **s** to step over/in, and all the panels should be updated dynamically while you are debugging. Either in the Registers or Stack panels, you can edit the values by inserting hex. This will be explained later. While hitting **tab** can help you moving between panels, it is highly recommended to use **m** to open the menu. As usual, you can use **h j k l** to move around the menu and will find tons of useful stuff there. You can also press **"** to quickly browse through the different options View offers and change the contents of the selected panel.

Split Screen

| is for the vertical and - is for the horizontal split. You can delete any panel by pressing **X**.

Split panels can be resized from Window Mode, which is accessed with **w**.

Window Mode Commands

```
|Panels Window mode help:
| ?      show this help
| ??     show the user-friendly hud
| Enter  start Zoom mode
| c      toggle cursor
| h j k l move around (left-down-up-right)
| JK     resize panels vertically
| HL     resize panels horizontally
| q      quit Window mode
```

Edit Values

Either in the Register or Stack panel, you can edit the values. Use **c** to activate cursor mode and you can move the cursor by pressing **h j k l**, as usual. Then, hit **i**, just like the insert mode of vim, to insert a value.

Tabs

Visual Panels also offer tabs to quickly access multiple forms of information easily. Press **t** to enter Tab Mode. All the tabs numbers will be visible in the top right corner.

By default you will have one tab and you can press **t** to create a new tab with the same panels and **T** to create a new panel from scratch.

For traversing through the tabs, you can type in the tab number while in Tab Mode.

And pressing **-** deletes the tab you are in.

Saving layouts

You can save your custom layout of your visual panels either by picking the option ‘Save Layout’ from the File menu of the menu bar or by running:

```
v= test
```

Where **test** is the name with which you’d like to save it.

You can open a saved layout by passing the name as the parameter to **v**:

```
v test
```

More about that can be found under **v?**.